


Added – TCP/IP Security

- Internet Protocol Review
- Security Problems in Internet Protocols

Internet protocol stack

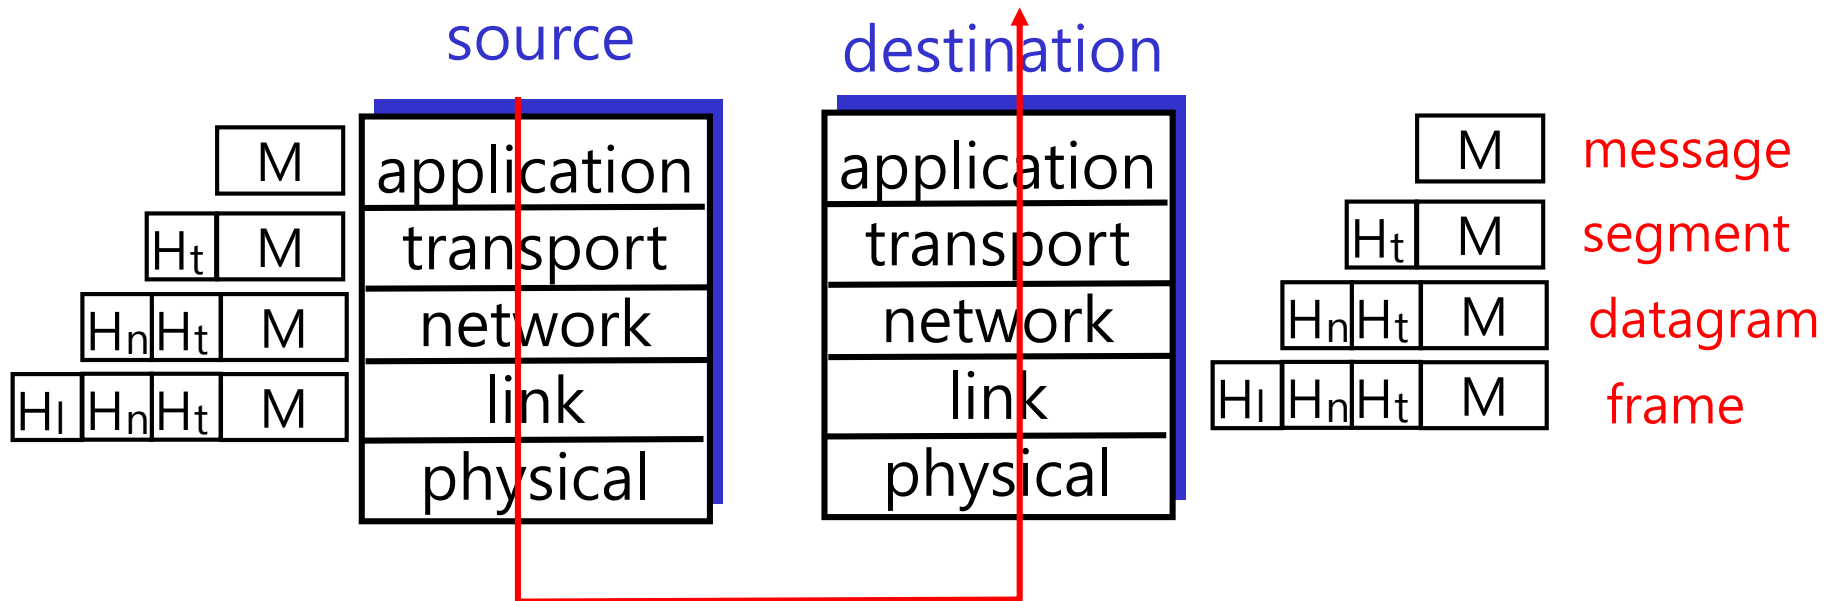
- application: supporting network applications
 - ftp, smtp, http
- transport: host-to-host data transfer services
 - tcp, udp
- network: routing of datagrams from source to destination
 - ip, routing protocols (rip, ospf, ...)
- link: data transfer between neighboring network elements
 - ppp, ethernet
- physical: bits “on the wire”



Internet protocol stack


Each layer takes data from above

- adds header information to create new data unit
- passes new data unit to layer below

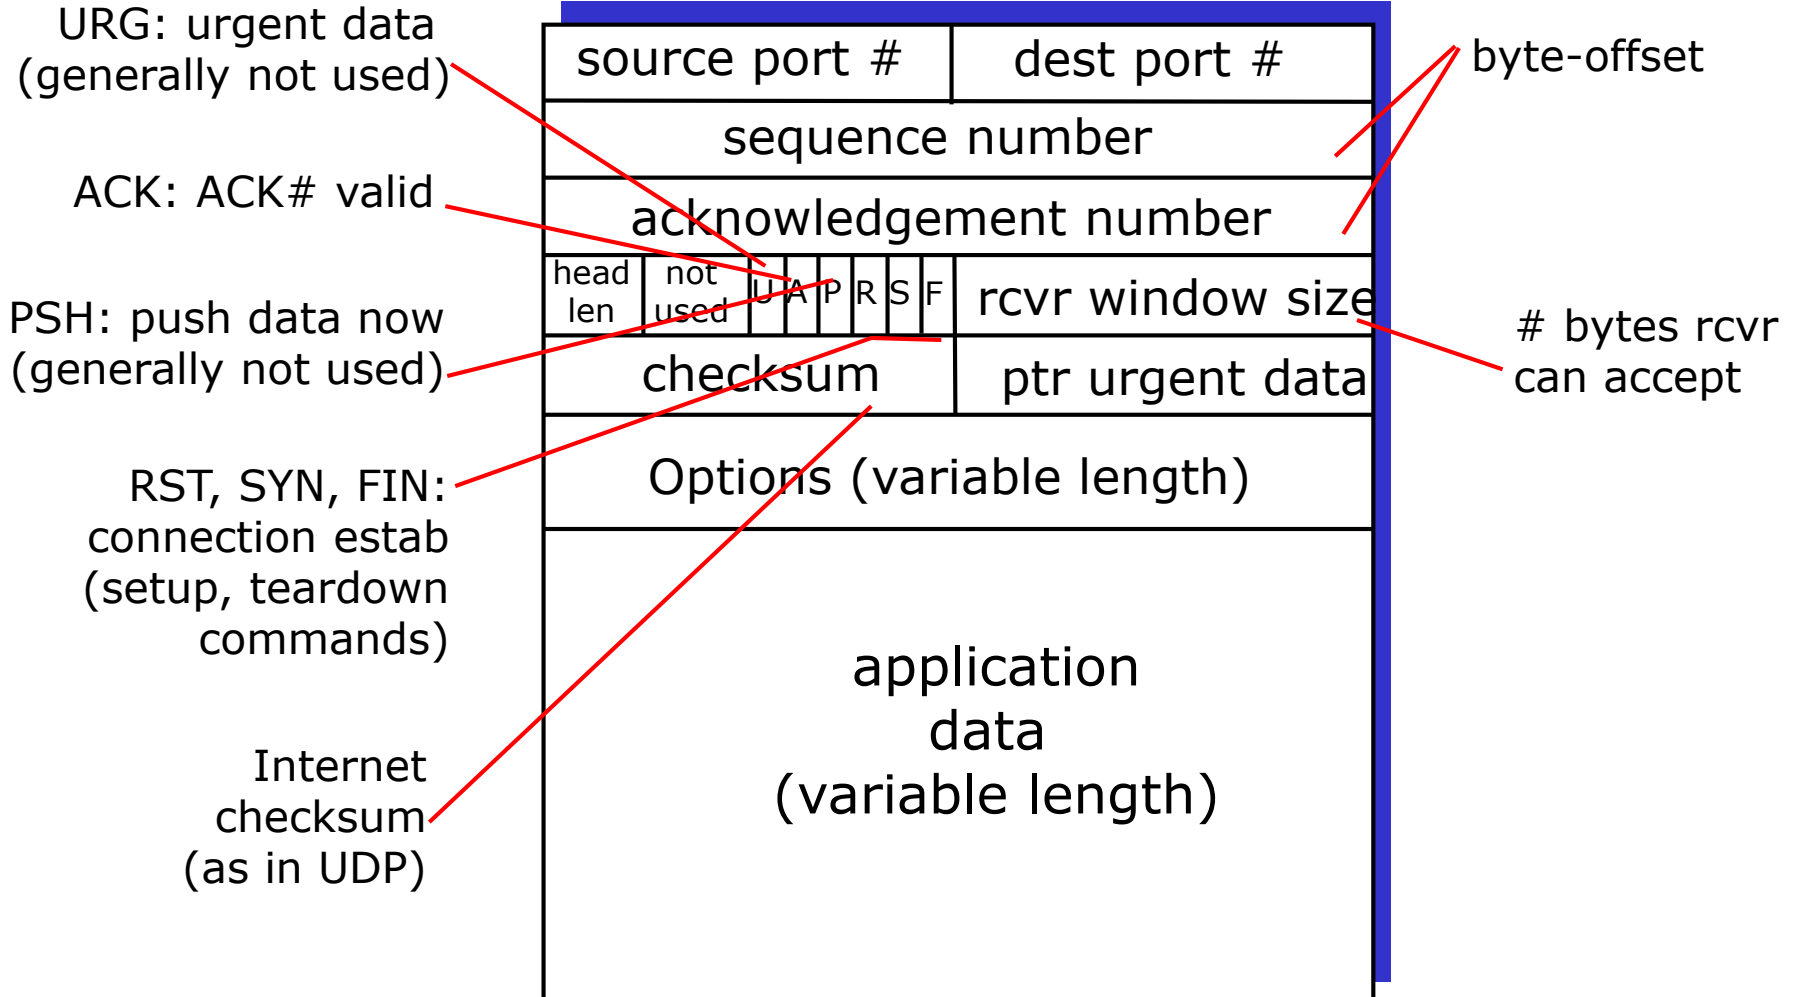


TCP Protocol Review

- point-to-point
- reliable, in-order *byte stream*:
 - no “message boundaries”
- pipelined:
 - TCP congestion and flow control set window size
- Buffering:
 - *send & receive buffers*
- full duplex data:
 - bi-directional data flow in same connection
- connection-oriented:
 - Handshaking: initialize sender, receiver state before data exchange
- flow controlled:
 - sender will not overwhelm receiver



TCP Protocol Review



ICMP: Internet Control Message Protocol

- used by hosts, routers, to get network-level information
 - error reporting: unreachable host, network, port, protocol
 - echo request/reply (used by ping)
- ICMP msgs carried in IP datagrams
- **ICMP message:** type, code plus first 8 bytes of IP datagram causing error

<u>Type</u>	<u>Code</u>	<u>description</u>
0	0	echo reply (ping)
3	0	dest. network unreachable
3	1	dest host unreachable
3	2	dest protocol unreachable
3	3	dest port unreachable
3	6	dest network unknown
3	7	dest host unknown
4	0	source quench(congestion control - not used)
5	0~3	icmp redirect
8	0	echo request (ping)
9	0	route advertisement
10	0	router discovery
11	0	TTL expired
12	0	bad IP header

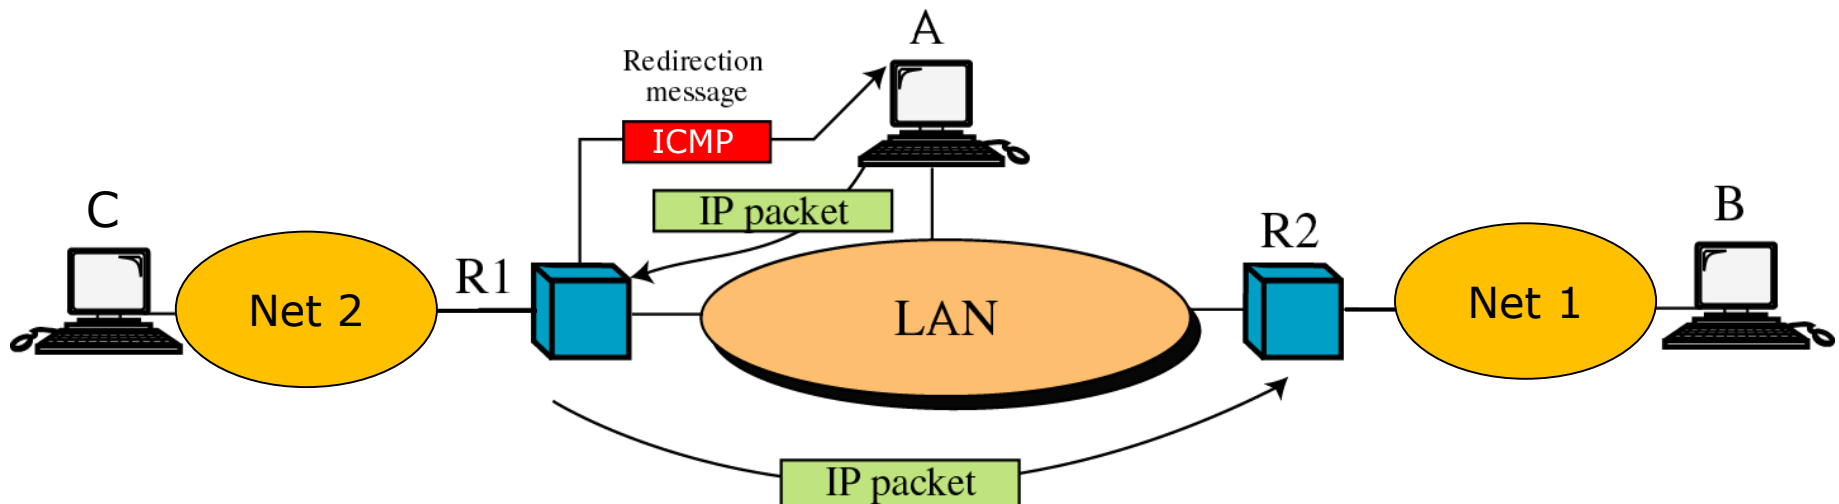
ICMP Redirects

- ICMP is used for routing error messages
 - TTL expired (traceroute)
 - Host unreachable
 - Echo request (ping program)
- Also used by default routers to redirect along quicker path

ICMP Redirect Routing


□ Redirection

- A host usually starts with a small routing table that is gradually augmented and updated
- A **ICMP redirection message** is sent from a router to a host on the same local network



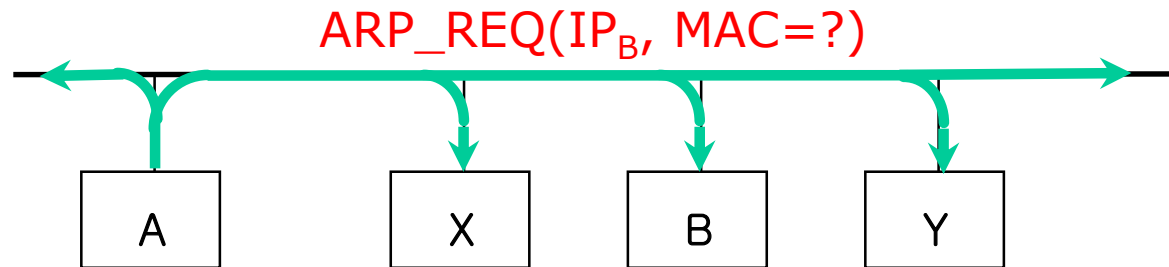
ARP: On-the-same-LAN routing

1. Route lookup determines its on the same subnet
2. Use ARP to determine what link layer address to send it to
3. Give it to Link layer




Address Resolution Protocol (ARP)

- Used in the broadcast LANs



a) host A broadcasts an ARP request containing I_B



b) host B responds with an ARP reply containing the pair (I_B, P_B)


- IP address set-up using **gratuitous ARP**

Inter-AS routing

□ AS: autonomous system

- Keeps the routing information within its AS
- Exports to the other AS router (exterior gateway)

□ Interior routing vs. exterior routing protocols



Inter-AS routing

- BGP (Border Gateway Protocol): the de facto standard
- Path Vector algorithm: and extension of Distance Vector
- Each Border Gateway broadcast to neighbors (peers) the entire path (ie, sequence of ASs) to destination
- For example, Gateway X may store the following path to destination Z:

Path (X,Z) = X,Y1,Y2,Y3,...,Z

Inter-AS routing

- Now, suppose GW X send its path to peer GW W
- GW W may or may not select the path offered by GW X, because of **policy** (cost, or loop prevention reasons)
- If GW W selects the path advertised by GW X, then:

Path (W,Z) = W, X,Y1,Y2,Y3,...,Z

Security Problems in Internet Protocols

- Reference: Security Problems in the TCP/IP Protocol Suite by Steve Bellovin
- R-services
- Source-routing
- ARP attacks
- Session hijacking
- TCP session stealing

Security problems in r-services


- `rsh` and `rcp` use the `.rhosts` file in your directory, which lists hosts and accounts to allow access from without a password
- r-services: allowed by `/etc/inetd`
- Example `.rhosts` (or `/etc/rhosts`) file:

```
www.ulsan.ac.kr mkkim  
*.ulsan.ac.kr mkkim  
* *
```

Making a Connection to rsh Server

□ 3-way handshaking


- C->S: $\text{SYN}(\text{ISN}_C)$
- S->C: $\text{SYN}(\text{ISN}_S), \text{ACK}(\text{ISN}_C+1)$
- C->S: $\text{ACK}(\text{ISN}_S + 1)$
- Client and Server exchange data (rsh command)



Making a Connection to rsh Server

□ 3-way handshaking

- X → S (spoof): $\text{SYN}(\text{ISN}_C)$
- S → C: $\text{SYN}(\text{ISN}_S), \text{ACK}(\text{ISN}_C + 1)$
- X → S (spoof): $\text{ACK}(\text{ISN}_S + 1)$
- How can X know ISN_S ?



Security problems in r-services

- a machine is running rsh, how can we pretend to be another machine to gain access?

Attack

- Source routing
- False routing table updates
- Session hijacking
- ICMP redirects
- False ARP packets
- TCP session stealing

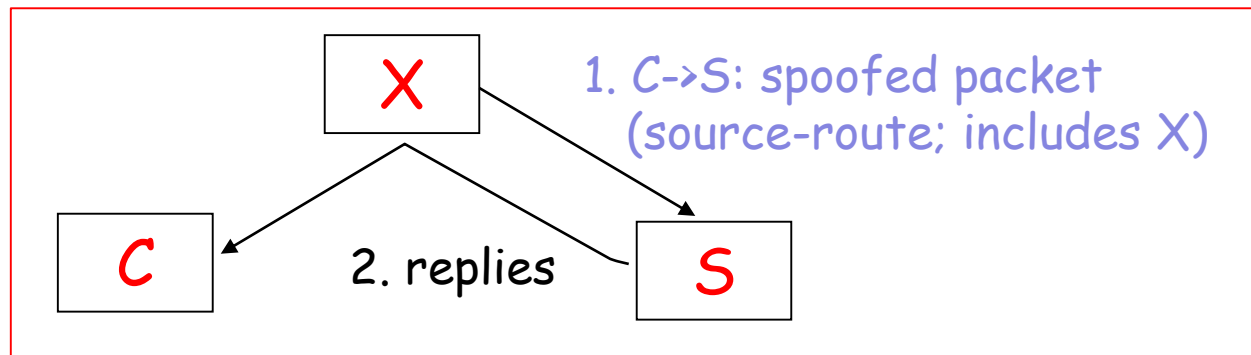
Defense

ignore source routes
secure routing protocols

ssh / secure connection
ignore ICMP redirects
publish ARP tables
ssh/ secure connection

Security problems in r-services

- ❑ Exploiting trusted relationships: C is a trusted host to S
- ❑ Source routing:
 - IP source-route option
 - The responder uses the source-route on the reply
 - Open a TCP connection to **rshd** spoofing the address of a trusted host, but include yourself in the source route




- ❑ Countermeasure: ignore source routes

Session hijacking

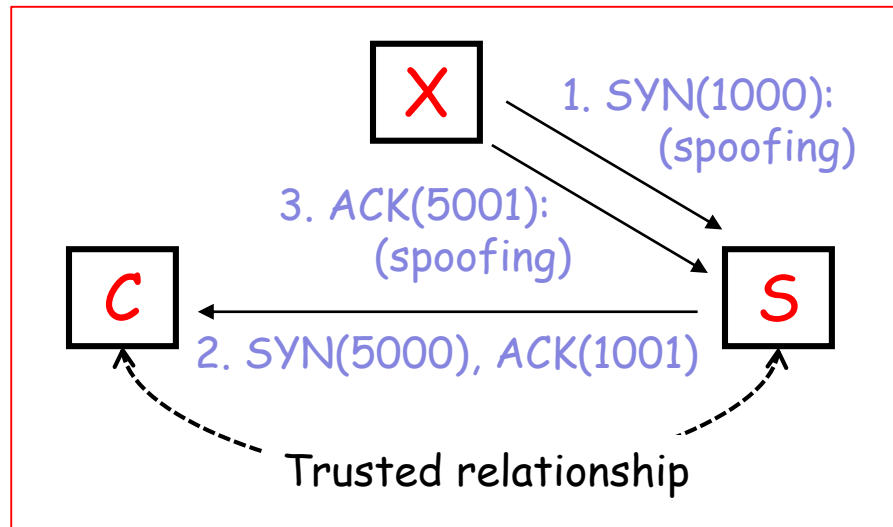
□ Normal TCP operation from client, C , to server, S

- $C \rightarrow S$: $\text{SYN}(\text{ISN}_C)$
- $S \rightarrow C$: $\text{SYN}(\text{ISN}_S), \text{ACK}(\text{ISN}_C + 1)$
- $C \rightarrow S$: $\text{ACK}(\text{ISN}_S + 1)$
- Client and Server exchange data
- **ISN number generation**
- 4.2BSD: increments 128/sec
(1 for 7.8ms)
- 4.3BSD: increments 125,000/sec
(1 for 8us)



Session hijacking

- Attacker X knows
 - S provides R-service
 - C is a trusted host of S
- X wants to disguise as C and access to S



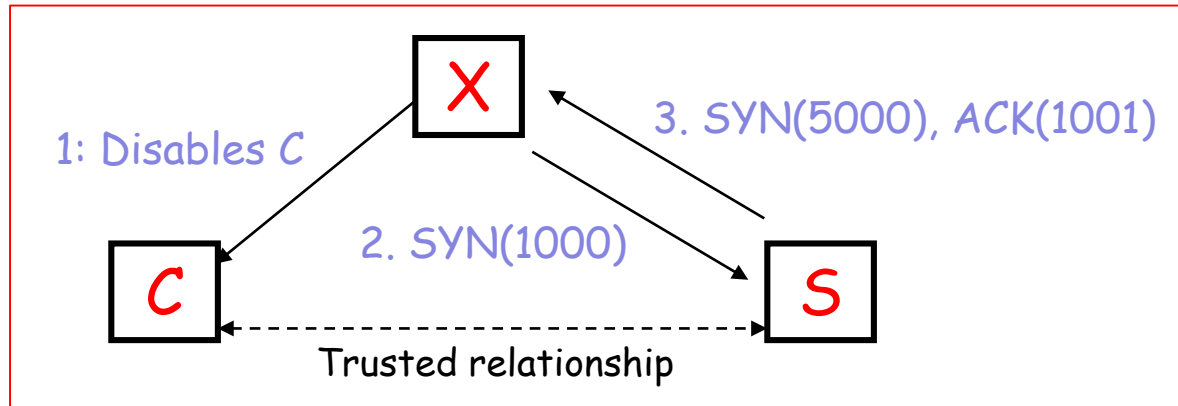
Session hijacking

□ Session hijacking: ISN estimation

(1) $X \rightarrow S$: $\text{SYN}(\text{ISN}_X)$ S: rshd server

(2) $S \rightarrow X$: $\text{SYN}(\text{ISN}_{S1})$, $\text{ACK}(\text{ISN}_X + 1)$

1. ISN estimation:



Session hijacking

□ Session hijacking: session hijacking

(3) X → S: SYN(ISN_X) [spoofs C] S: rshd server

(4) S → C: SYN(**ISN_{S2}**), ACK(ISN_X + 1)

(5) X → S: ACK(ISN_{S2} + 1)

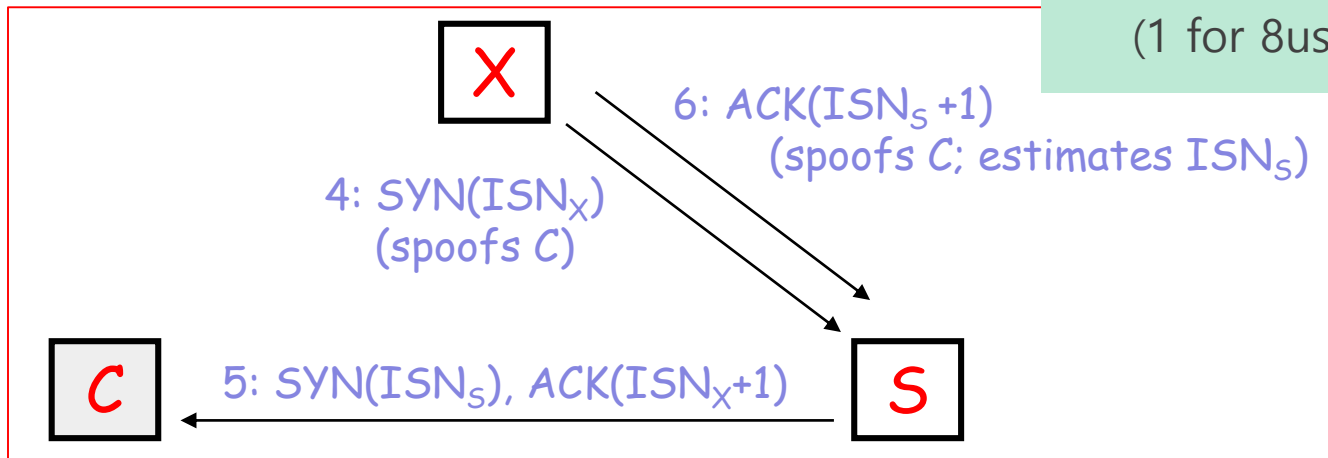
[spoofs C; **estimates ISN_{S2}**]

2. Session hijacking:

□ ISN number generation

■ 4.2BSD: increments 128/sec
(1 for 7.8ms)

■ 4.3BSD: increments 125,000/sec
(1 for 8us)




Session hijacking

□ Session hijacking: execute remote commands

(6) X->S: [`echo "* *" >> ~/.rhosts`] [spoofs C]

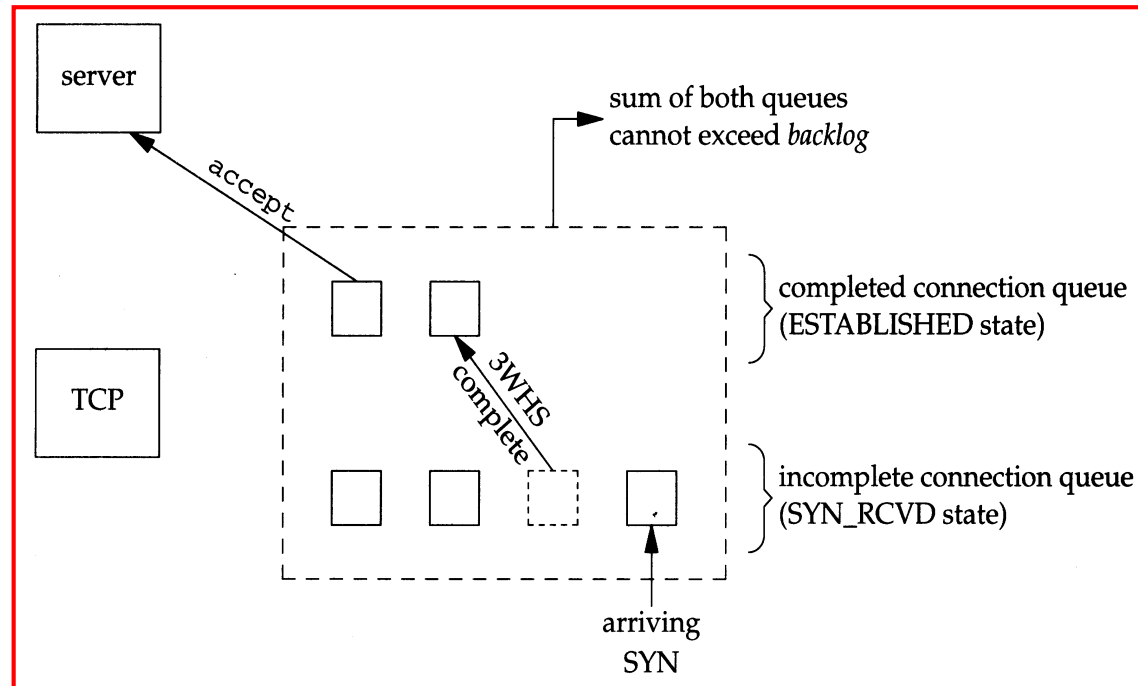
(7) X->S: RESET [spoofs C]

3. **Executes remote commands:** X can rlogin from anywhere in the world



Disabling hosts: SYN Flooding DoS

- Send lots of spoofed SYN packets to a victim host
- TCP connection queue
 - Kernel maintains two queues for each listening socket:
 - **Incomplete conn. queue:** an entry for each conn. in SYN_RCVD state
 - **Completed conn. queue:** an entry for each conn. in ESTABLISHED state
 - **accept** returns first entry on the completed queue



Attacking Routing to Exploit rsh

- Dynamic routing updates
 - OSPF: link-state algorithm
 - RIP: distance vector algorithm
- Attacker injects a fake RIP update msg. stating it has a good path to host C
 - All subsequent packets to C will be routed to the attacker
 - The attacker initiates connection to rshd of the server (spoofing C)
- Defense: uses secure routing protocols
 - Only accept authenticated updates
 - Requires key management

ICMP Attack

- ICMP redirect: forces a machine to route thru you
 - Requires an existing connection
 - Open a spoofed connection to the host you want to attack
 - Then send a spoofed ICMP redirect to the victim redirecting it to the gateway you've compromised
- Others
 - ICMP destination unreachable
 - Frequent ICMP source quenches

ARP Attacks

- When a machine sends an ARP request out, you could answer that you own the address
- Unfortunately, ARP just accepts replies without requests!
 - Just send a spoofed reply message saying your MAC address owns a certain IP address
 - Repeat frequently so that cache doesn't timeout
- Messages are routed through you to sniff or modify

ARP Spoofing - Countermeasures

- “Publish” MAC address of router/default gateway and trusted hosts to prevent ARP spoofing
 - Statically defining the IP to MAC address mapping
 - (e.g.) `arp -s 203.250.77.254 00:01:02:03:04:ab pub`

```
C:\Windows\system32>arp -a


Interface: 192.168.43.160 --- 0x15
Internet Address      Physical Address      Type
192.168.43.1          70-b7-aa-90-a3-07     dynamic
192.168.43.160        00-aa-00-62-c6-09     static
192.168.43.255        ff-ff-ff-ff-ff-ff     static
224.0.0.22            01-00-5e-00-00-16     static
224.0.0.252           01-00-5e-00-00-fc     static
239.255.255.250       01-00-5e-7f-ff-fa     static
255.255.255.255       ff-ff-ff-ff-ff-ff     static
```

TCP Session Stealing

- Reference: "A Simple Active Attack Against TCP" by Laurent Joncheray, *USENIX Symposium*, June 1995
- Active attack using TCP desynchronized states
 - attacker is in the path b/w the client and server (attacker can sniff all the packets and inject some spoofed packets)
 - Steps:
 1. attacker sniffs the communication b/w the two
 2. attacker disables the communication by desynchronizing the client and the server
 3. attacker injects spoofed packets acceptable for both ends

TCP Session Stealing

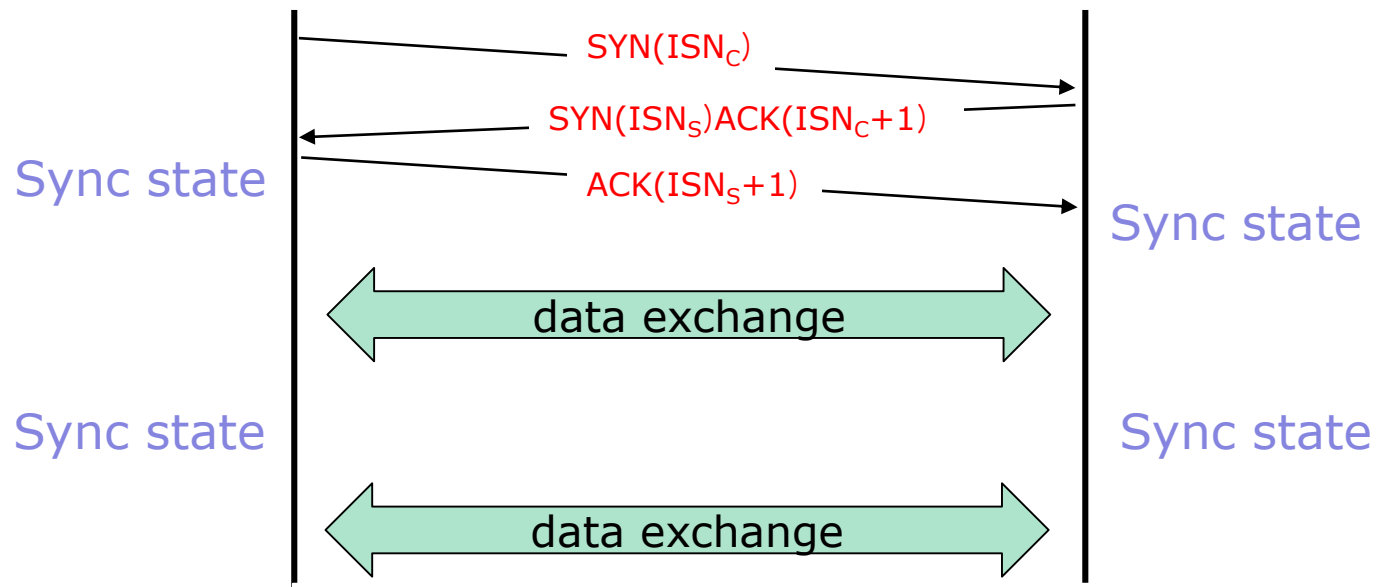
- Initial state after connection setup: **synchronized state** b/w client C and server S
 - $S_SEQ = C_ACK$ and $C_SEQ = S_ACK$



TCP Session Stealing

□ Desynchronized state b/w client C and server S

- Both in "Established state"
- No data is being sent (stable state)
- $S_SEQ \neq C_ACK$ and $C_SEQ \neq S_ACK$



TCP Session Stealing

□ In Desynchronized state b/w client C and server S

$C \rightarrow S$ connection: $C_SEQ \neq S_ACK$

□ When $(C_SEQ > S_ACK + W)$ or $(C_SEQ < S_ACK)$:

- packet is dropped

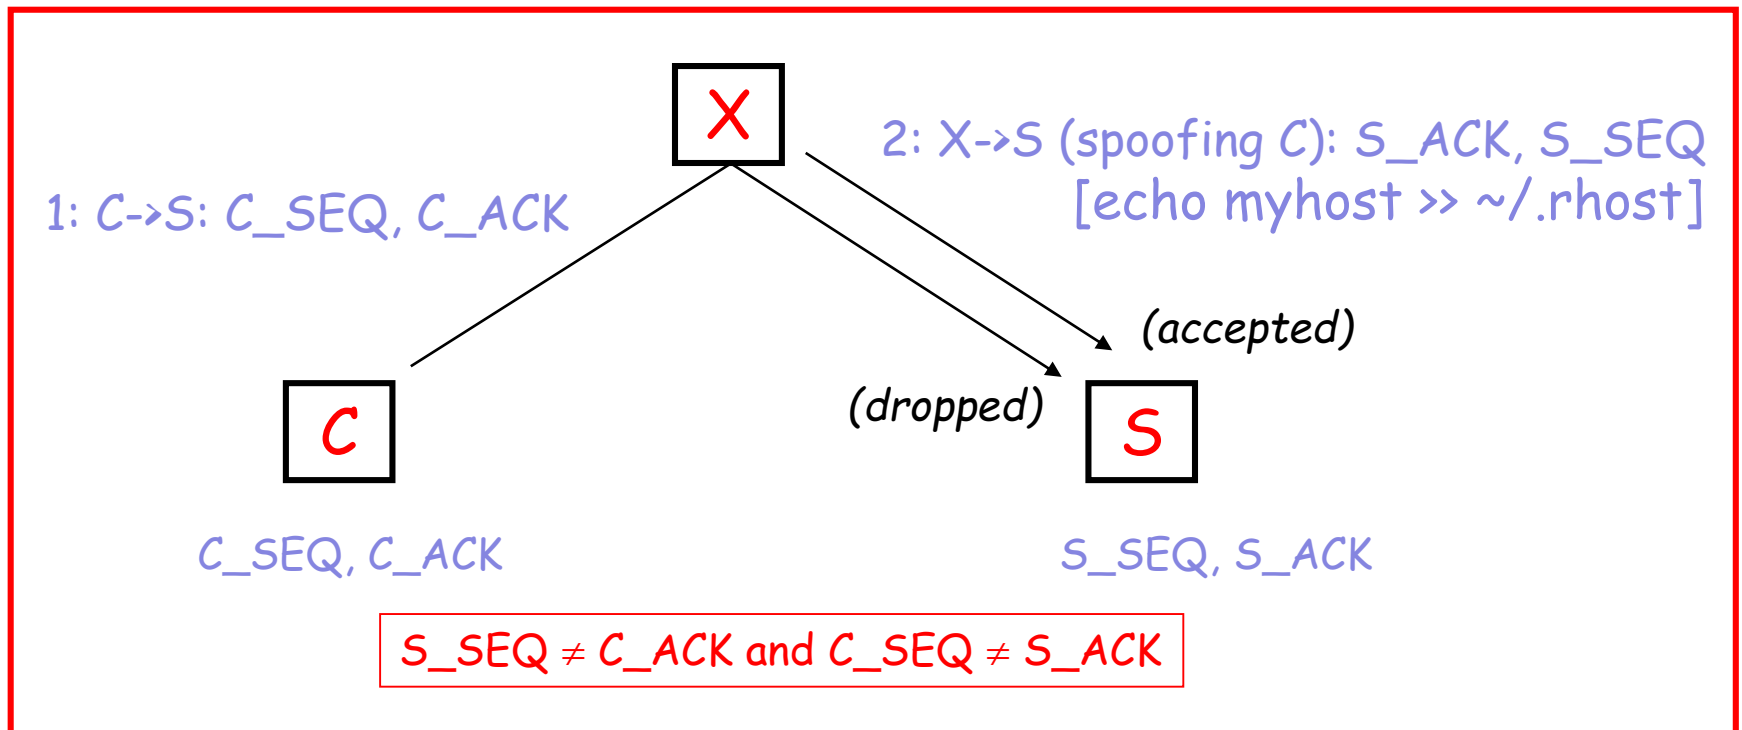
□ When $S_ACK < C_SEQ < S_ACK + W$:

- packet is accepted (buffered) but not sent to process

□ In both cases, $ACK(S_ACK)$ is sent : ACK packet with (S_SEQ, S_ACK)

TCP Session Stealing

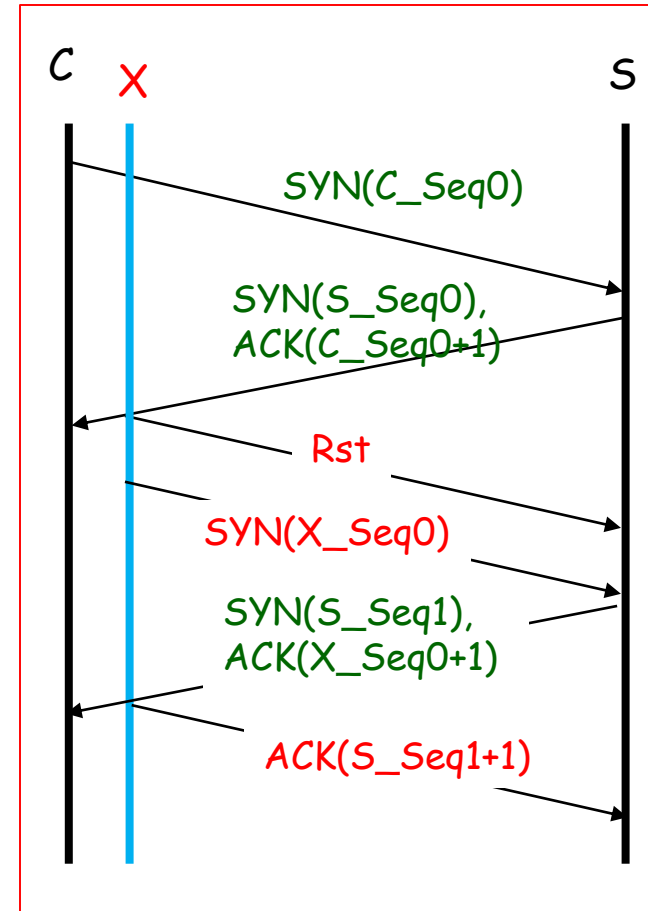
- the attacker knows the desynchronized state, and can send any acceptable data to the server
 - E.g. `[echo myhost >> ~/.rhost]` for rlogin



Desynchronization

□ Early desynchronization

1. C->S(Syn): C_Seq0 ; C: Syn_Sent
2. S->C(Syn/Ack): S_Seq0, C_Seq0+1
; S: Syn_Rcvd
; C: Established (C_Seq0+1, S_Seq0+1)
(before the packet C->S(Ack): S_Seq0+1)
3. X->S(spoofing C, Rst)
4. X->S(spoofing C, Syn): X_Seq0
; the same port # used in (1)
5. S->C(Syn/Ack): S_Seq1, X_Seq0+1
6. X->S(spoofing C, Ack): S_Seq1+1
; S: Established (S_Seq1+1, X_Seq0+1)



The Attack

□ Null data desynchronization

1. The attacker watches the session without interfering
2. During a quiet period, the attacker sends a large amount of null data (IAC,NOP for telnet): nothing happens, server only changes the TCP Ack number
3. Now, when the client sends data, it is dropped by the server because it's lower than the server's window
4. The attacker does the same with the client

□ Defense: ssh connection, or IPsec