

Chapter 3:

Transport Layer

- Principles of Transport Layer
- TCP (Transmission Control Protocol)
- UDP (User Datagram Protocol)

Transport Layer

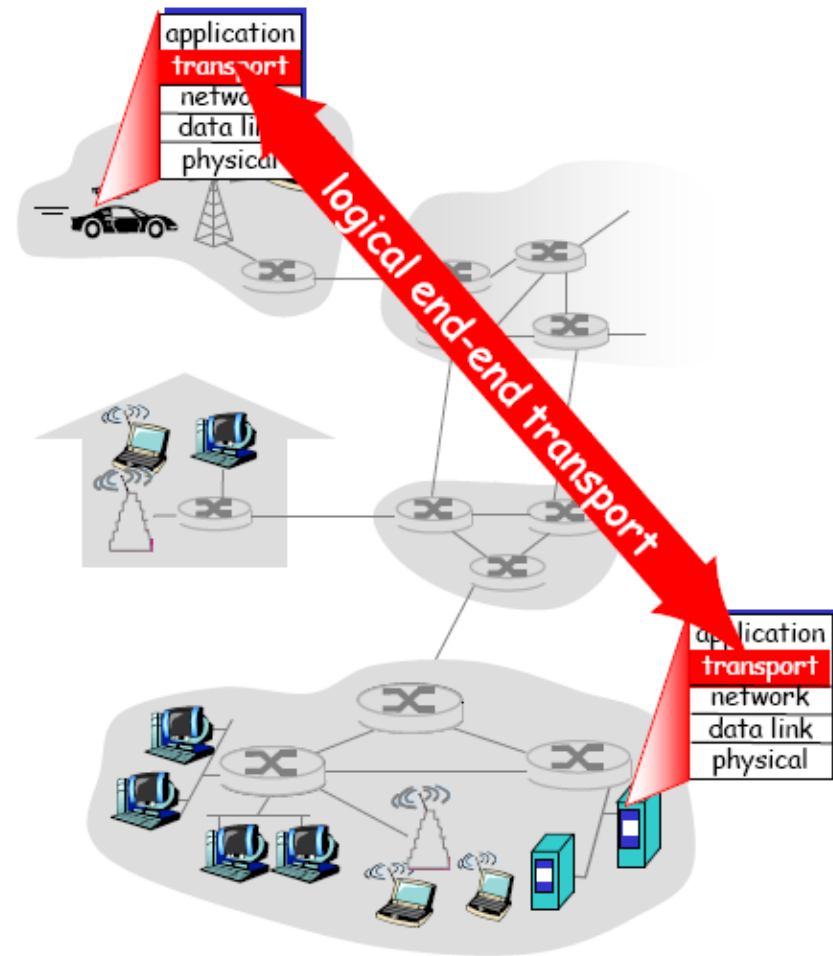
- Understand principles behind transport layer services
 - multiplexing/demultiplexing
 - reliable data transfer
 - flow control
 - congestion control
- Transport layer protocol
 - TCP
 - UDP

Chap 3

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- TCP congestion control

Transport Services and Protocols

- Provide *logical communication* between app processes
- Transport protocols in end system
 - send side: breaks app msg into *segments*, passes to IP layer
 - rcv side: reassembles segments into messages, passes to app layer
- More than one transport protocol available to apps
 - Internet: TCP and UDP



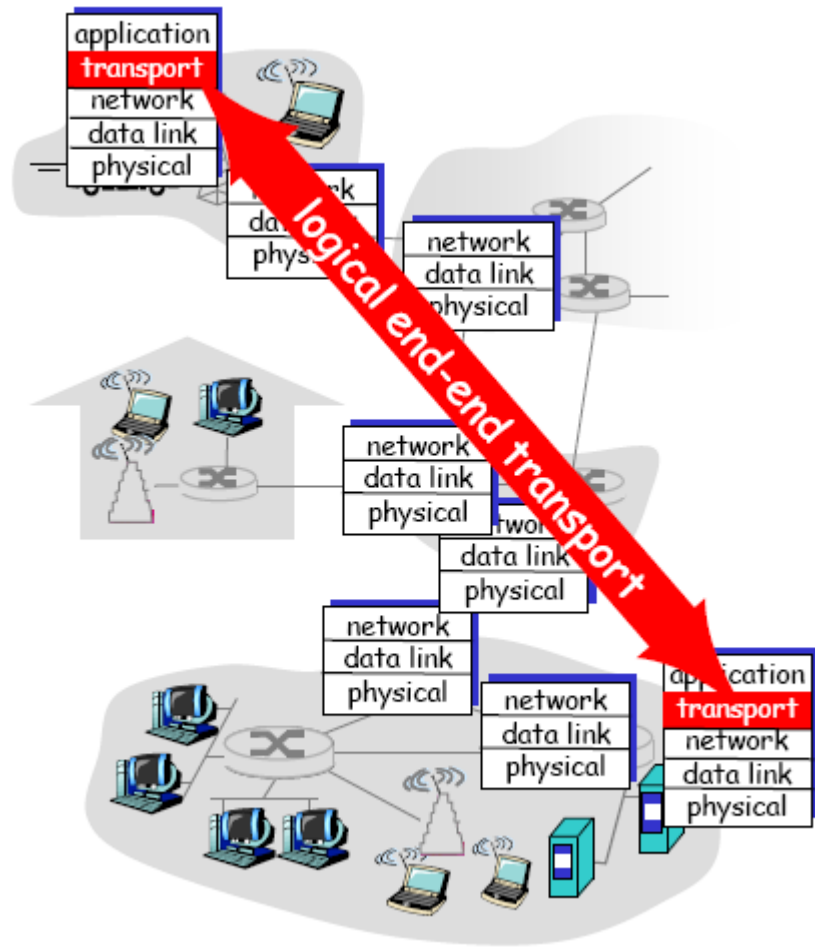
Transport vs. Network Layer

□ Network layer

- logical communication between hosts
- IP (Internet Protocol)

□ Transport layer

- logical communication between processes
- relies on, enhances, network layer services
- TCP and UDP



Chap 3

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- TCP congestion control

Multiplexing and Demultiplexing

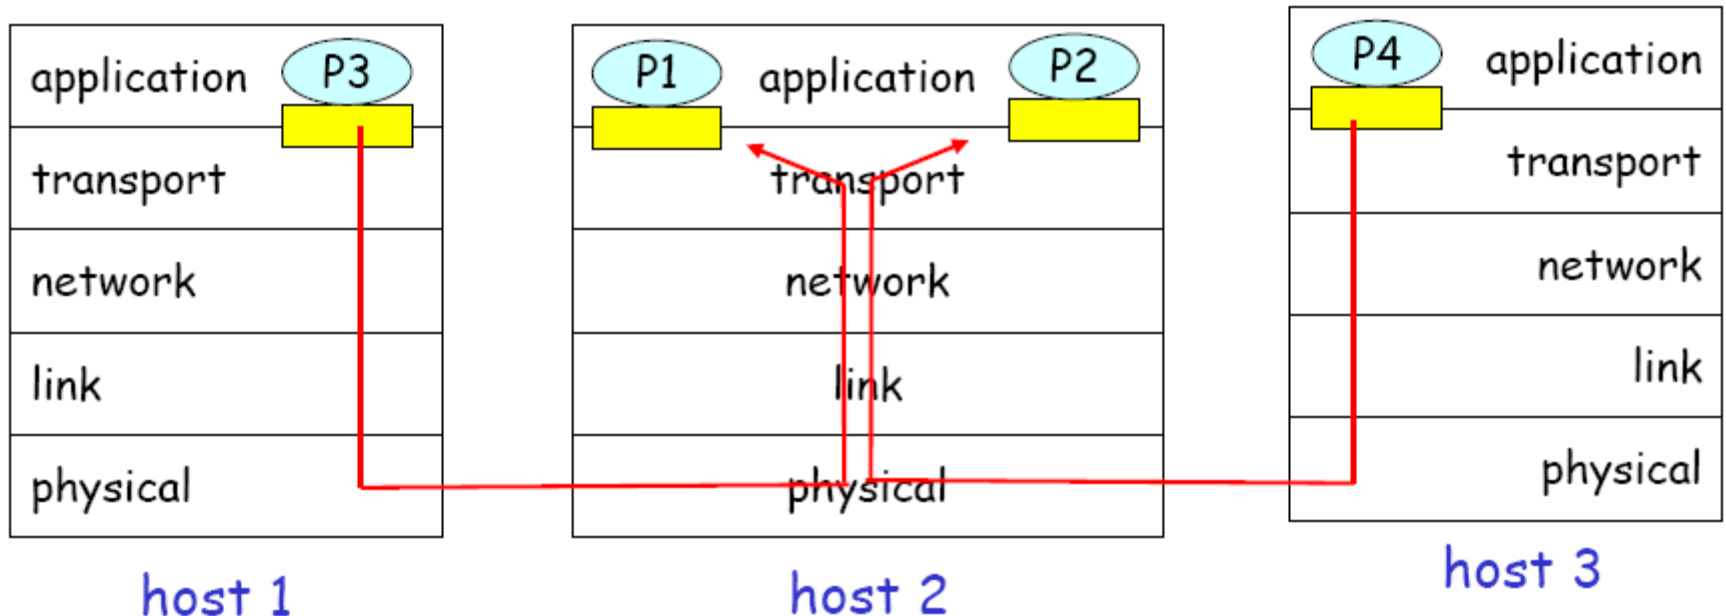
Demultiplexing at rcv host:

delivering received segments
to correct socket

Multiplexing at send host:

gathering data from multiple
sockets, enveloping data with
header (later used for
demultiplexing)

 = socket  = process

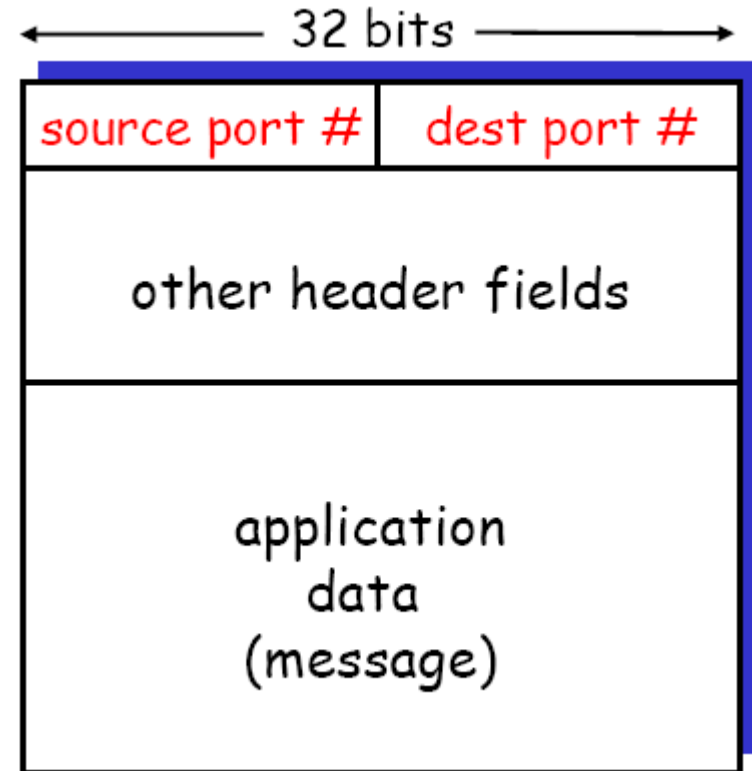


Multiplexing and Demultiplexing

□ host receives IP datagrams

- each datagram has source IP address, destination IP address, and
- each datagram has source, destination port number

□ host uses (IP addr, port number) to direct segment to appropriate socket



Connectionless Demultiplexing

- UDP socket identified by two-tuple:

`(local_IP, local_Port, *, *)`

- UDP datagram delivery: (SrcIP, DstIP, SrcPort, DstPort)

- checks destination (IP, port number) in segment:

`(DstIP==local_IP) && (DstPort==local_Port)`


- directs UDP segment to socket with that port number

- IP datagrams with different source IP and/or source port numbers directed to a UDP socket

Connectionless Demultiplexing

□ UDP socket

```
DatagramSocket svrSocket = new DatagramSocket(6428);
```

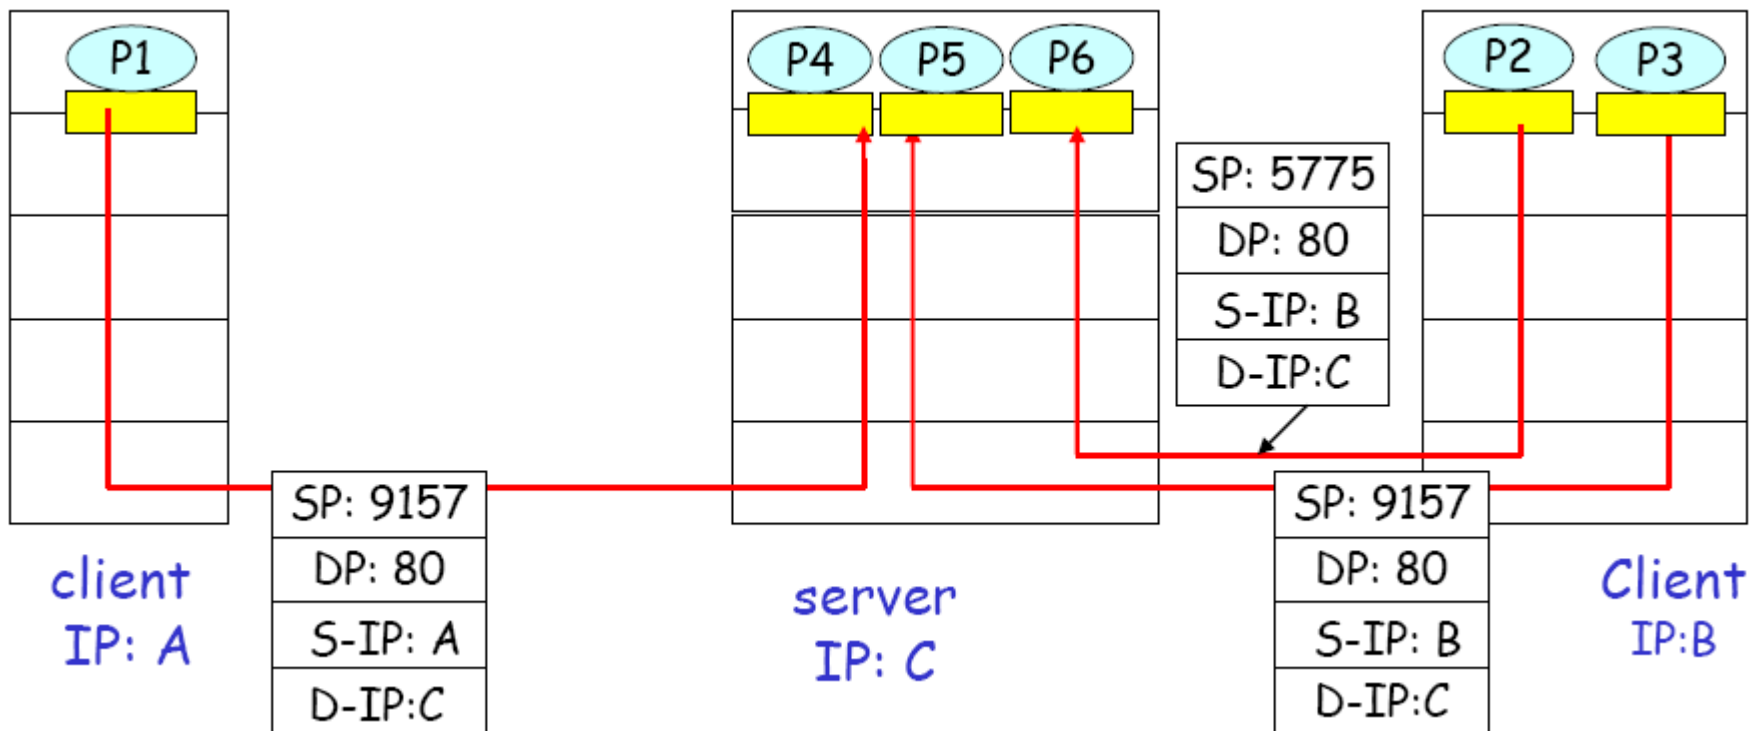


Connection-oriented Demultiplexing

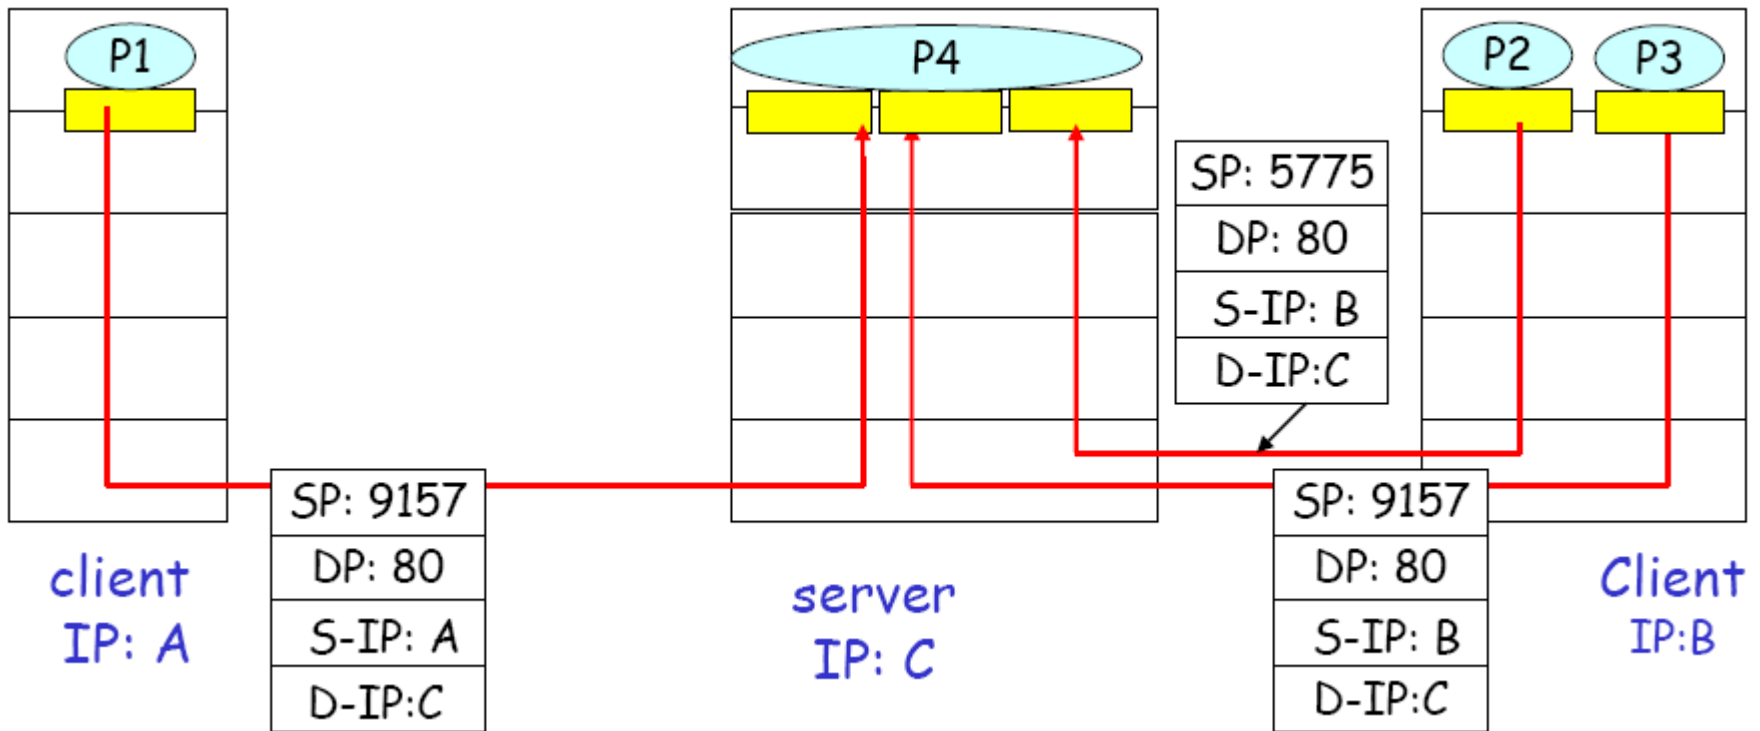
- TCP socket identified by 4-tuple:
 - `(local_IP, local_Port, remote_IP, remote_Port)`
- TCP segment delivery: `(SrcIP, DstIP, SrcPort, DstPort)`
 - checks source and destination (IP, port number) in segment:
`(DstIP==local_IP) && (DstPort==local_Port) && (SrcIP==remote_IP) && (SrcPort==remote_Port)`
- Server support many simultaneous TCP sockets:
 - each socket identified by its own 4-tuple
 - Listening socket: `(local_IP, local_Port, *, *)`
 - Connected socket: `(local_IP, local_Port, remote_IP, remote_Port)`

Connection-oriented Demultiplexing

□ TCP socket



Connection-oriented Demux: Threaded Web Server



Chap 3

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- TCP congestion control

UDP: User Datagram Protocol

□ UDP [RFC 768]

- “best effort” service
- UDP segments may be lost and delivered out of order to appl

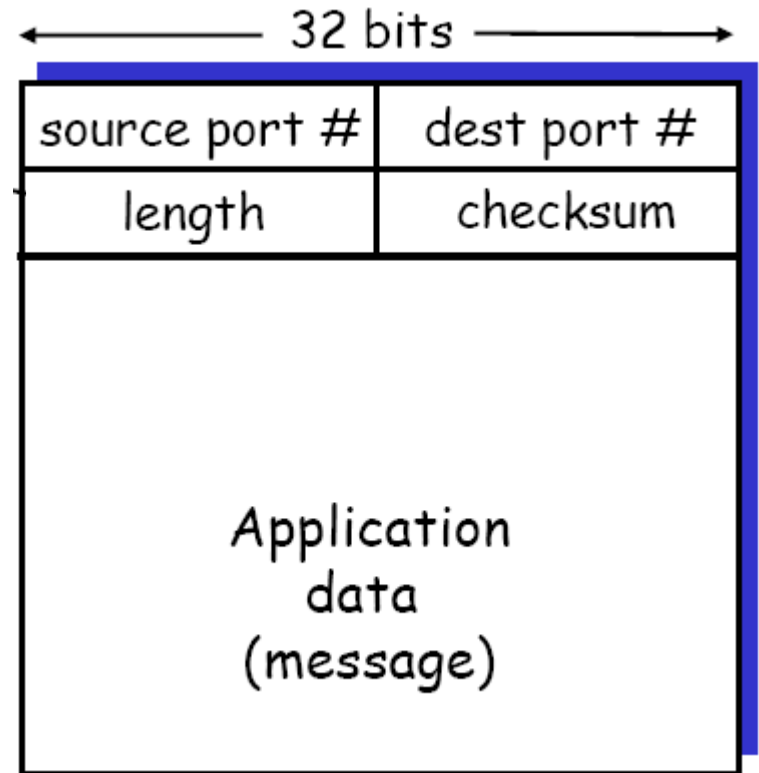
□ *connectionless*:

- no handshaking between UDP sender, receiver
- each UDP segment handled independently of others

□ UDP has smaller protocol overhead than TCP

UDP

- Often used for streaming multimedia apps
 - loss tolerant
 - rate sensitive
- other UDP uses
 - DNS
 - SNMP
- reliable transfer over UDP
 - need to implement reliability at application layer



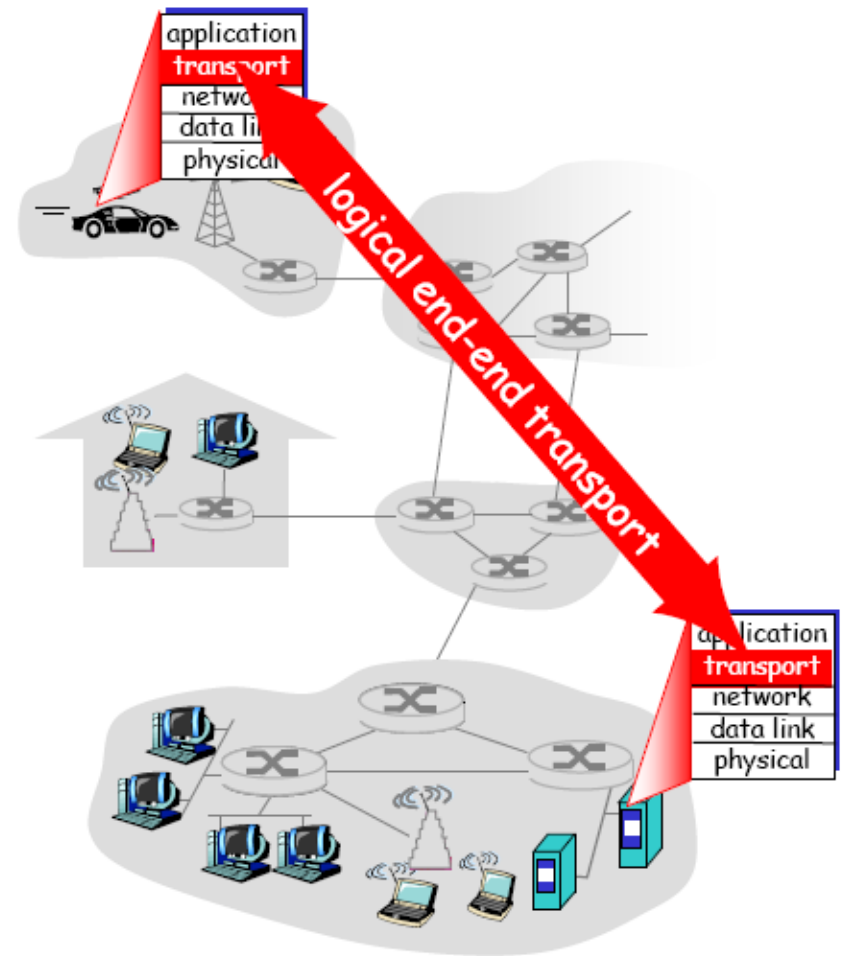
UDP datagram format

Chap 3

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- TCP congestion control

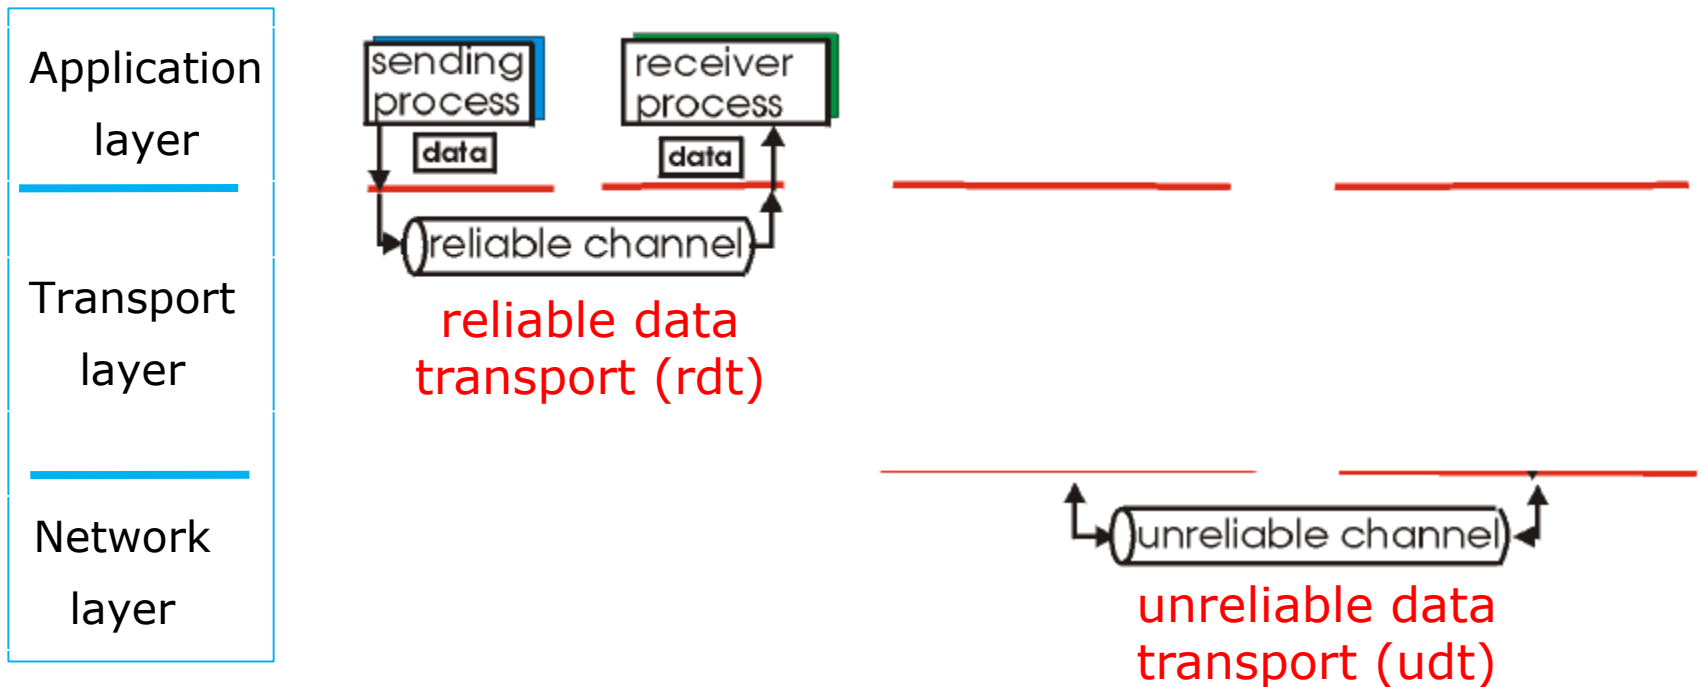
Principles of Reliable data transfer

- Applications needs reliable channel but network provides unreliable channel
- Network layer
 - unreliable communication between hosts
- Transport layer
 - end-to-end reliable communication between processes



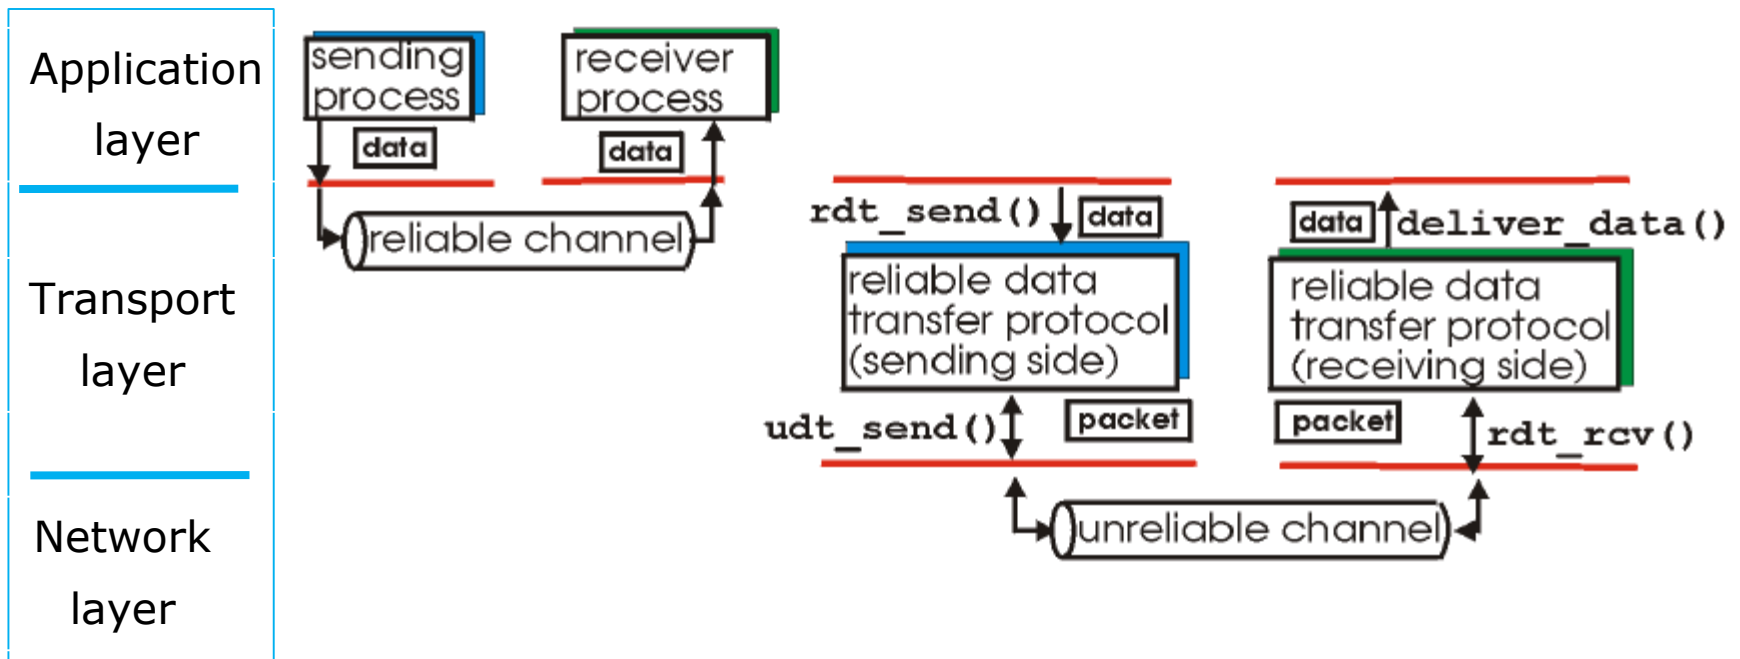
Principles of Reliable data transfer

□ Reliability needs in transport layer



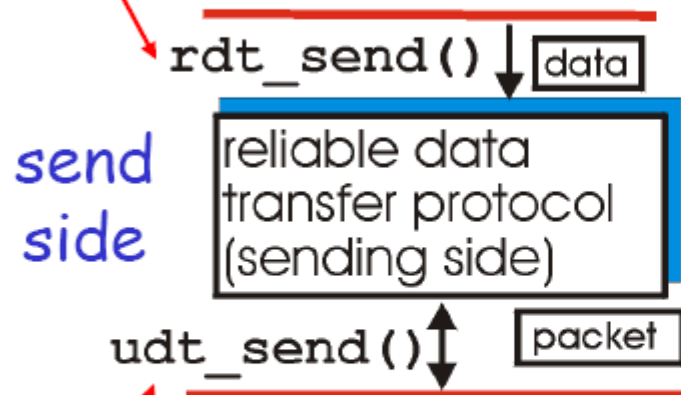
Principles of Reliable data transfer

□ Reliability in transport layer



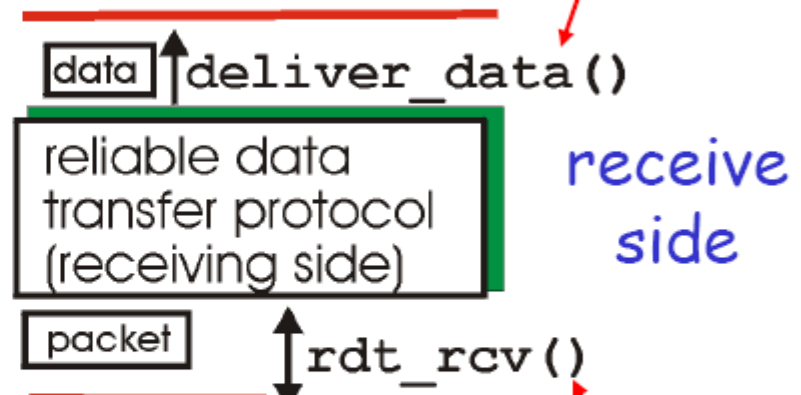
Reliable Data Transfer: getting started

rdt_send() : called from above, (e.g., by app.). Passed data to deliver to receiver upper layer



udt_send() : called by rdt, to transfer packet over unreliable channel to receiver

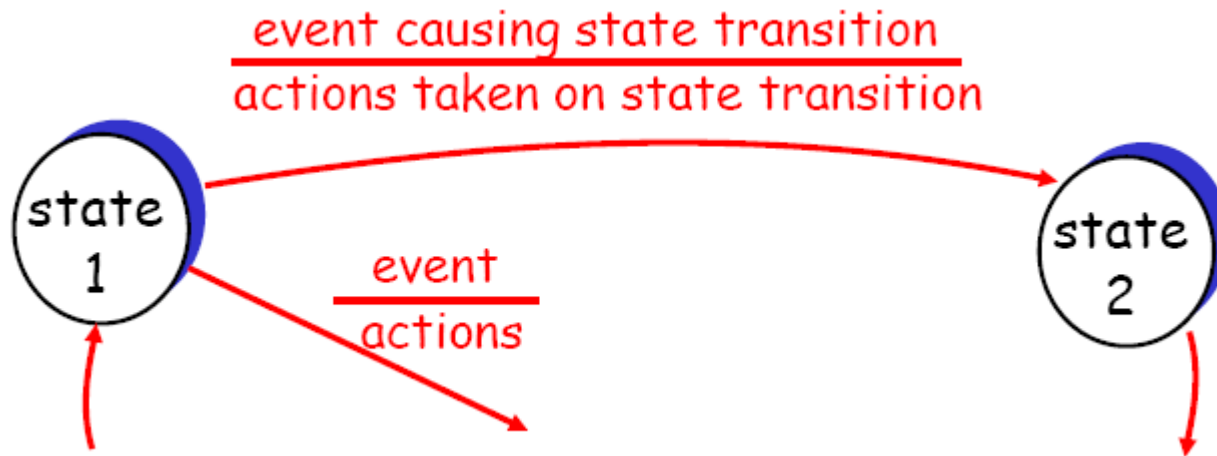
deliver_data() : called by rdt to deliver data to upper



rdt_rcv() : called when packet arrives on rcv-side of channel

Reliable Data Transfer: getting started

- incrementally develop sender, receiver sides of reliable data transfer protocol
- use **finite state machines (FSM)** to specify sender, receiver




Rdt1.0: reliable transfer over a reliable channel

□ underlying channel perfectly reliable


- no bit errors
- no loss of packets

□ separate FSMs for sender, receiver:

- sender sends data into underlying channel



sender



receiver

Rdt2.0: channel with bit errors, but no packet loss

- underlying channel may flip bits in packet


- **checksum** to detect bit errors

- Q: how to recover from errors:

- **acknowledgements (ACKs)**: receiver explicitly tells sender that pkt received OK
 - **negative acknowledgements (NAKs)**: receiver explicitly tells sender that pkt had errors
 - sender **retransmits pkt on receipt of NAK**

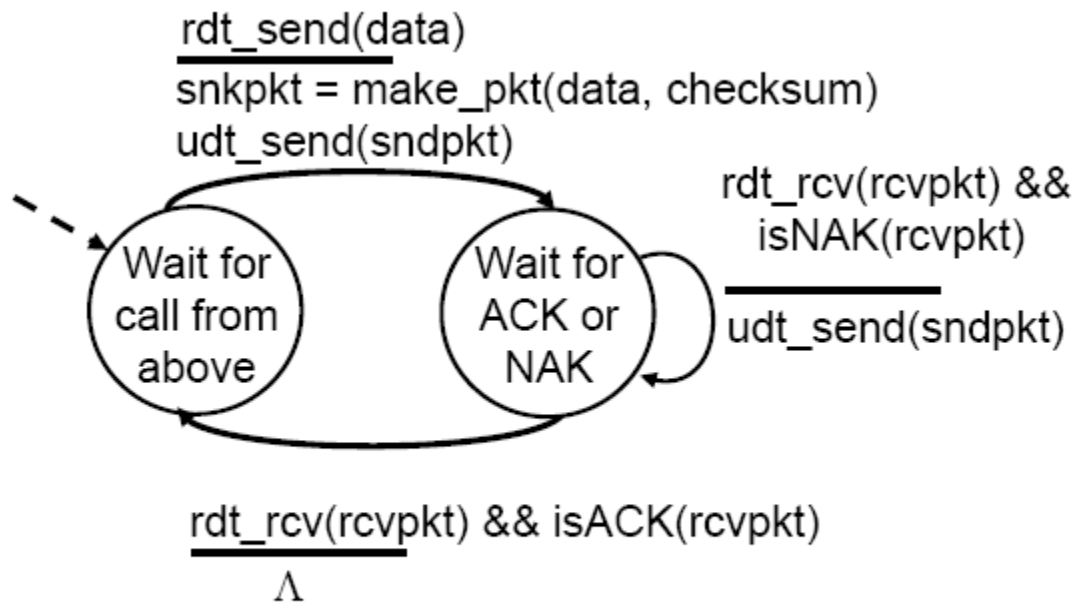
- new mechanisms in **rdt2.0**

- error detection
 - receiver feedback: rcv sends control msgs (ACK,NAK) to sender




Rdt2.0: channel with bit errors, but no packet loss

□ RDT2.0 FSM: Sender and receiver



sender


receiver



Rdt2.0: problems

□ What happens if ACK/NAK corrupted?

- needs checksum in ACK or NAK packet
- if corrupted, sender doesn't know what happened at receiver
- can't just retransmit: possible duplicate



Rdt2.0: problems



□ Handling duplicates:

- sender retransmits current pkt if ACK/NAK garbled
- sender adds *sequence number* to each pkt
- receiver discards duplicate pkt

□ Stop-and-wait

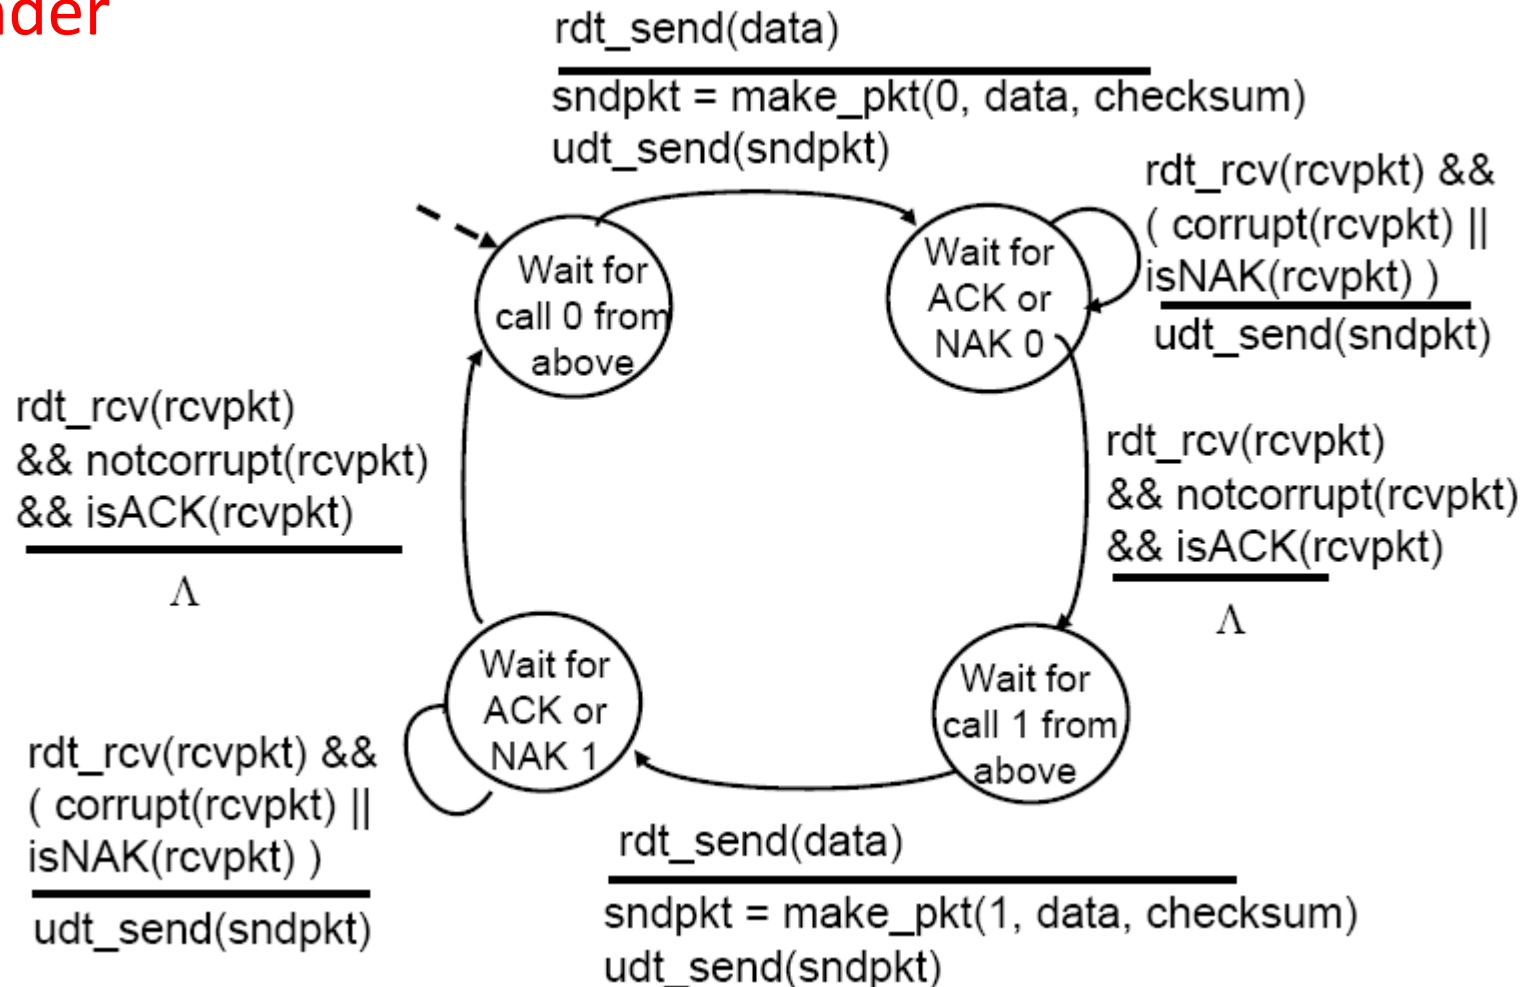
- Sender sends one packet, then waits for receiver response
- Needs 1-bit sequence number

sequence number:



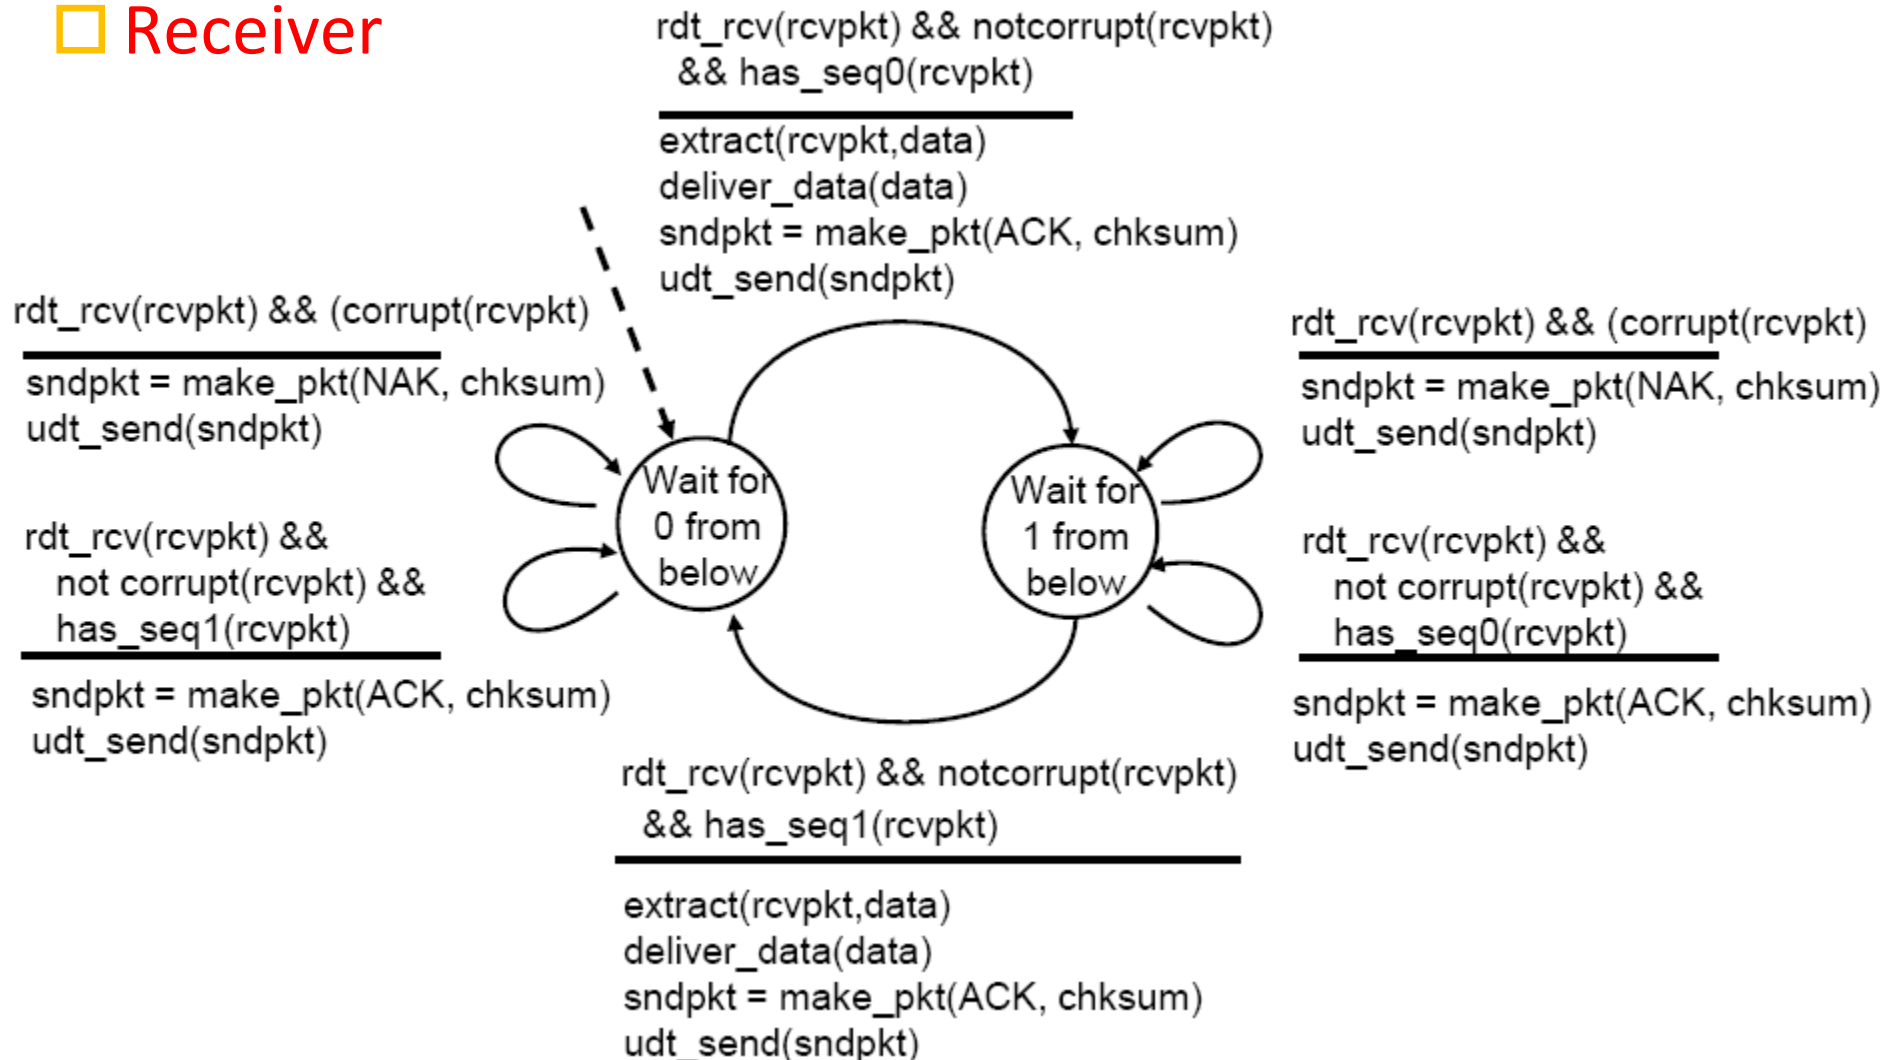
Rdt2.1: handles garbled ACK/NAKs

□ Sender



Rdt2.1: handles garbled ACK/NAKs

□ Receiver



Rdt2.1: discussion

□ Sender:

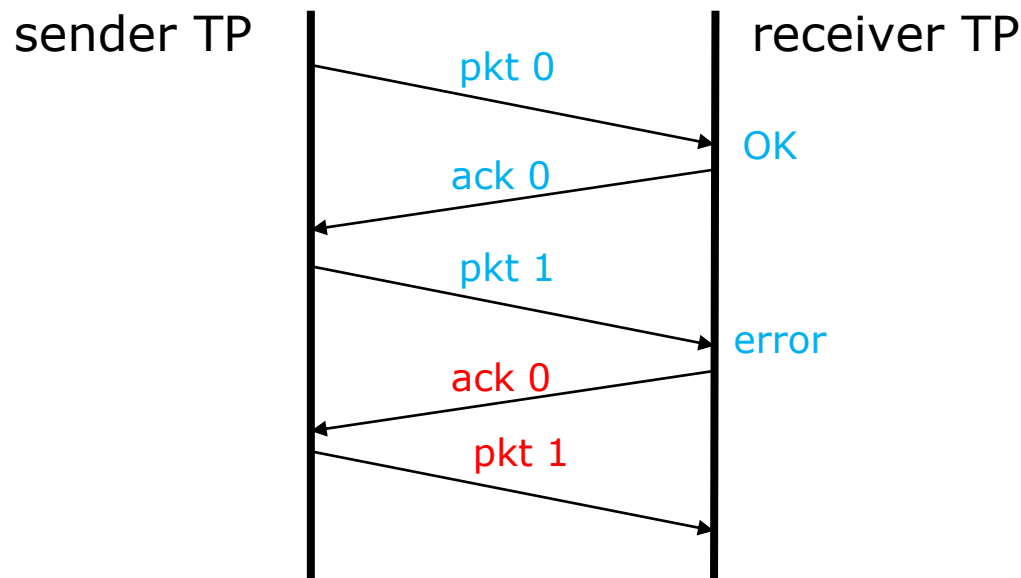
- seq # added to pkt: needs two seq. #'s (0,1) in stop-and-wait
- must check if received ACK/NAK packet corrupted
- needs 4 states: state must “remember” whether “current” pkt has 0 or 1 seq. #

□ Receiver:

- must check if received packet is duplicate: state indicates whether 0 or 1 is expected pkt seq #
- note: receiver can *not* know if its last ACK/NAK received OK at sender

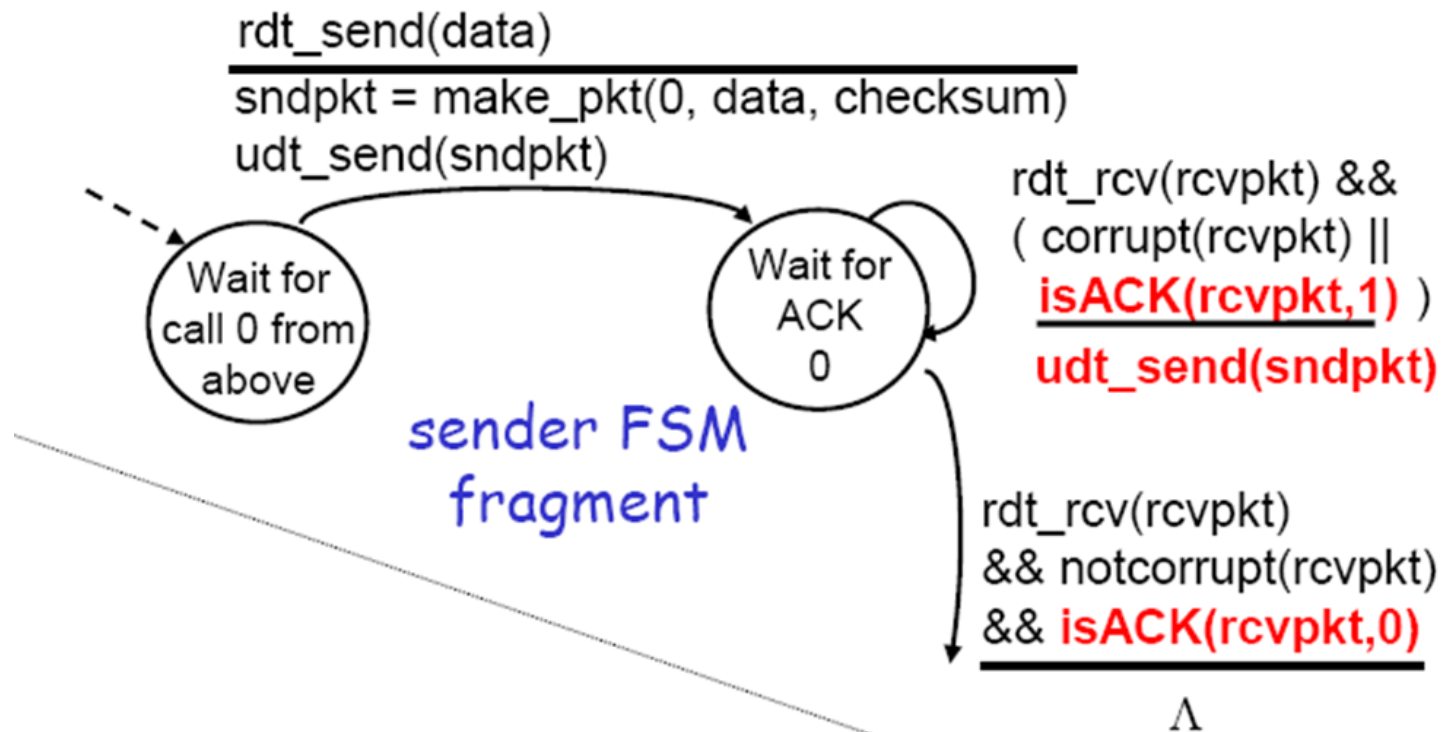
Rdt2.2: a NAK-free protocol

- same functionality as rdt2.1, using ACKs only
- instead of NAK, receiver sends ACK for last pkt received OK
 - ACK has seq # to indicate pkt being ACKed
- duplicate ACK at sender results in same action as NAK:
retransmit current pkt



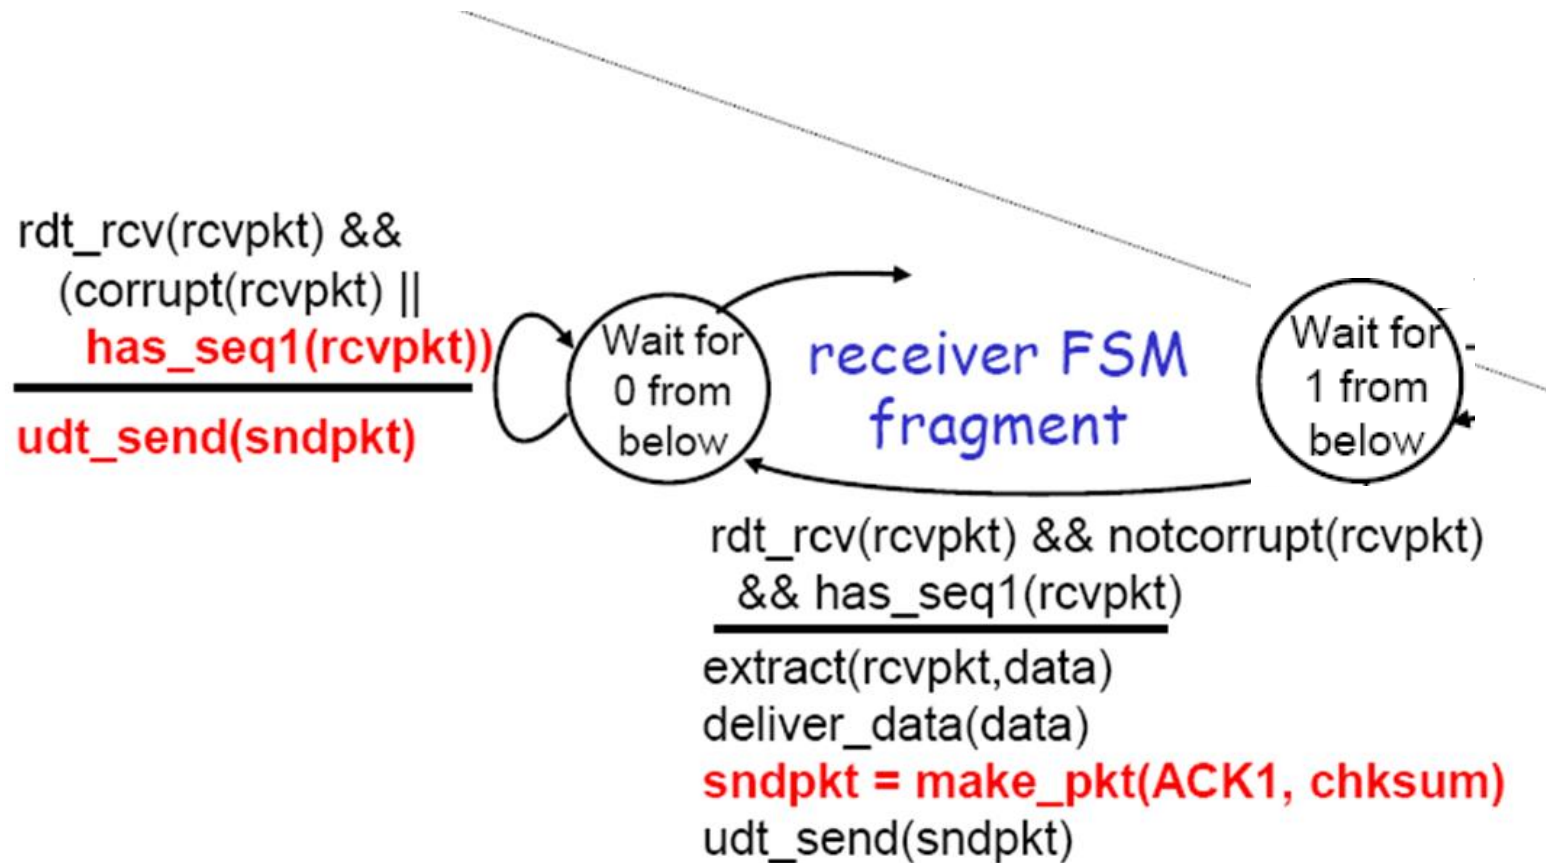
Rdt2.2: sender and receiver

□ Sender:




Rdt2.2: sender and receiver

□ Receiver:



Rdt3.0: channels with errors *and* loss

- underlying channel can lose packets (data or ACKs)
 - checksum, seq. #, ACKs, retransmissions will be of help, but not enough



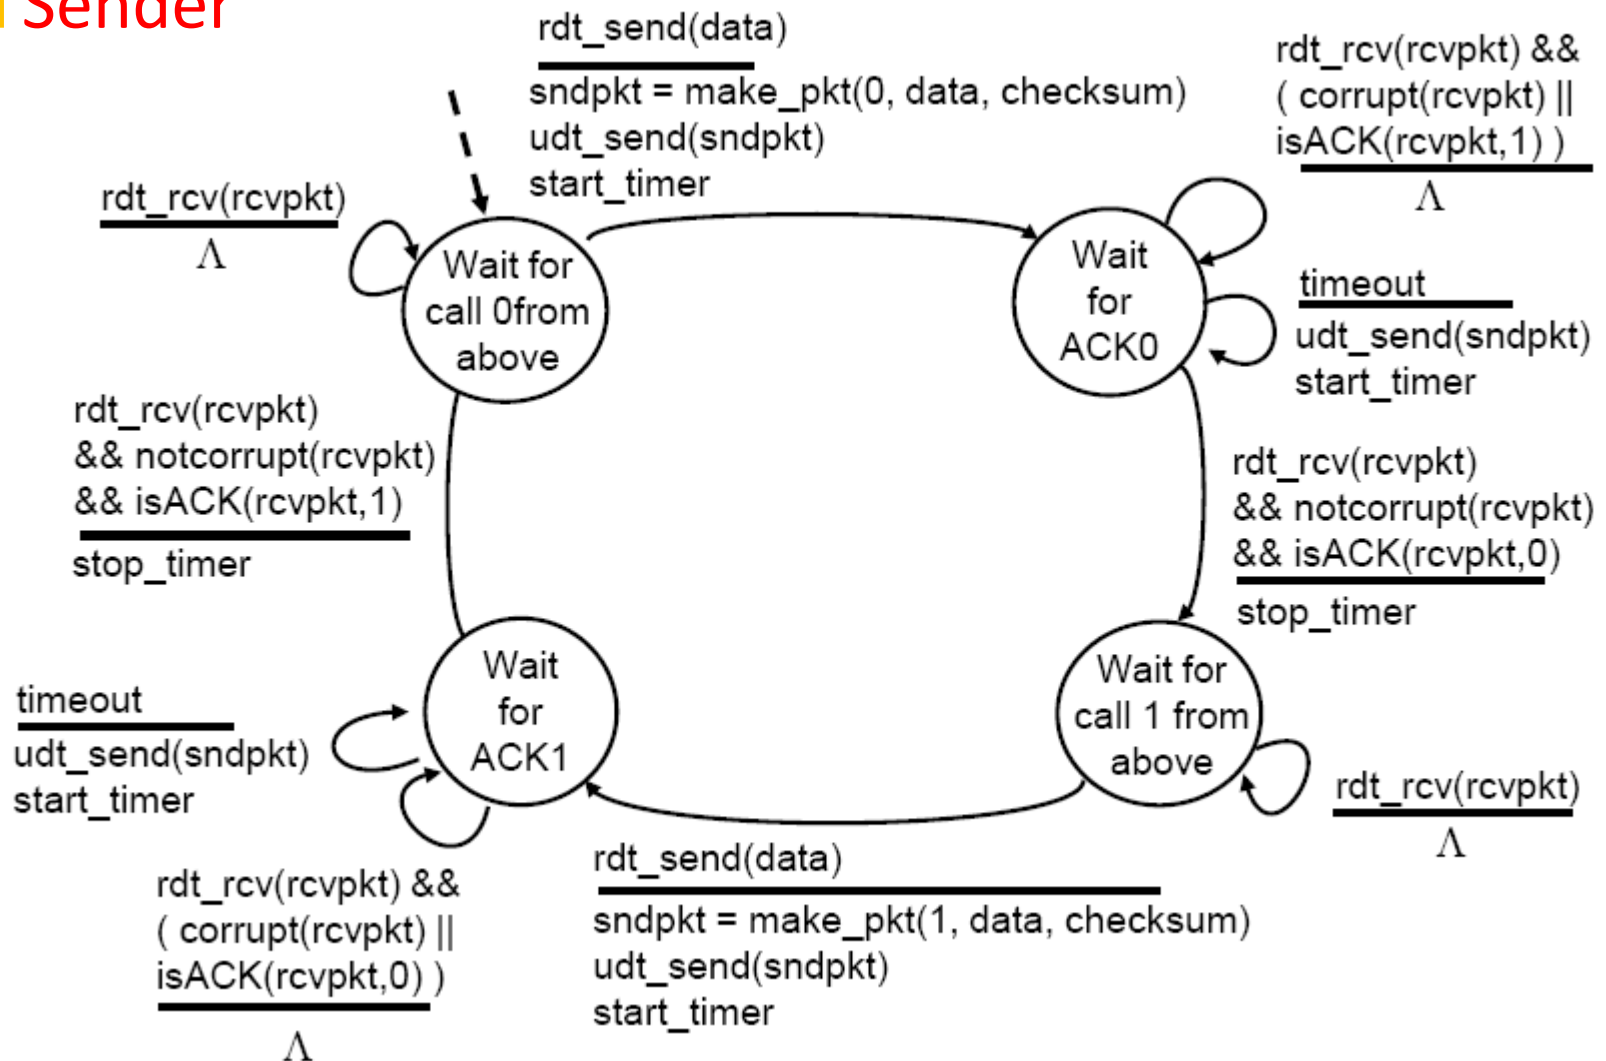
Rdt3.0: channels with errors *and* loss

□ Approach:

- sender waits “reasonable” amount of time for ACK → needs **timeout timer (TO)** for waiting ACK packet
- **retransmits if no ACK received in this time**
- what happens if pkt (or ACK) just delayed (not lost):
 - retransmission will be duplicate, but use of seq. #'s already handles this
 - receiver must specify seq # of pkt being ACKed


Rdt3.0:

□ Sender

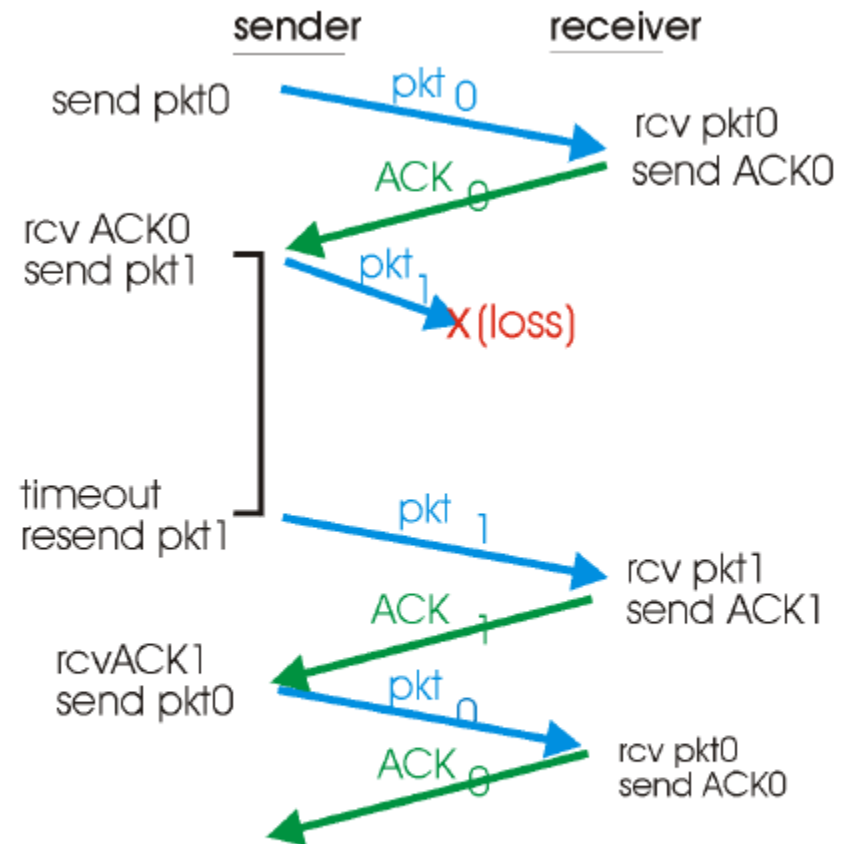


Rdt3.0 Operation

□ No loss



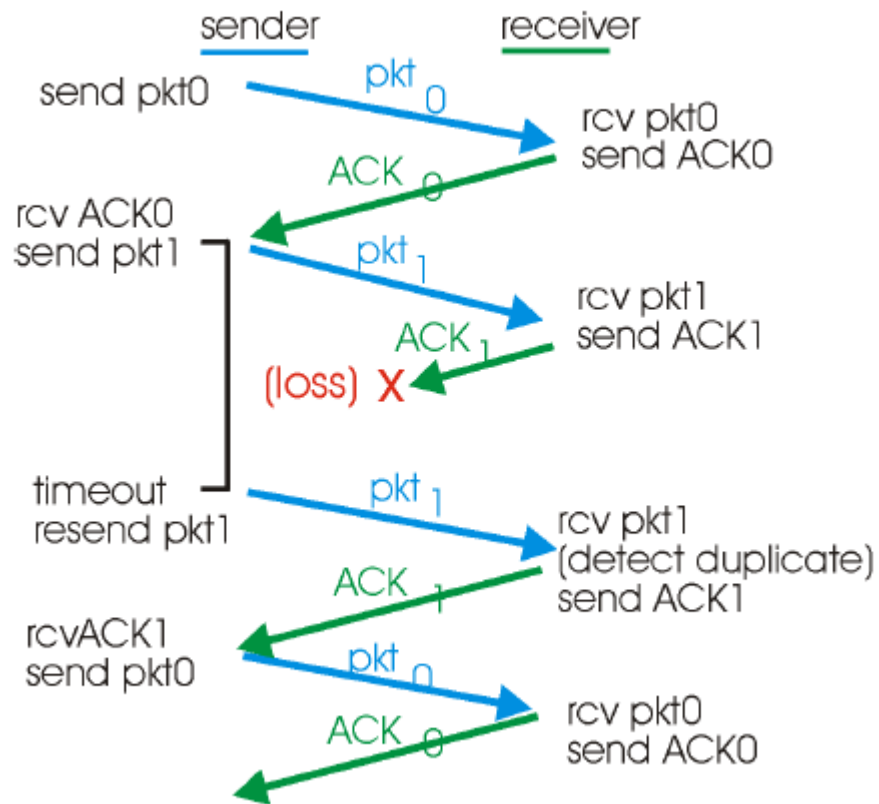
□ Lost packet



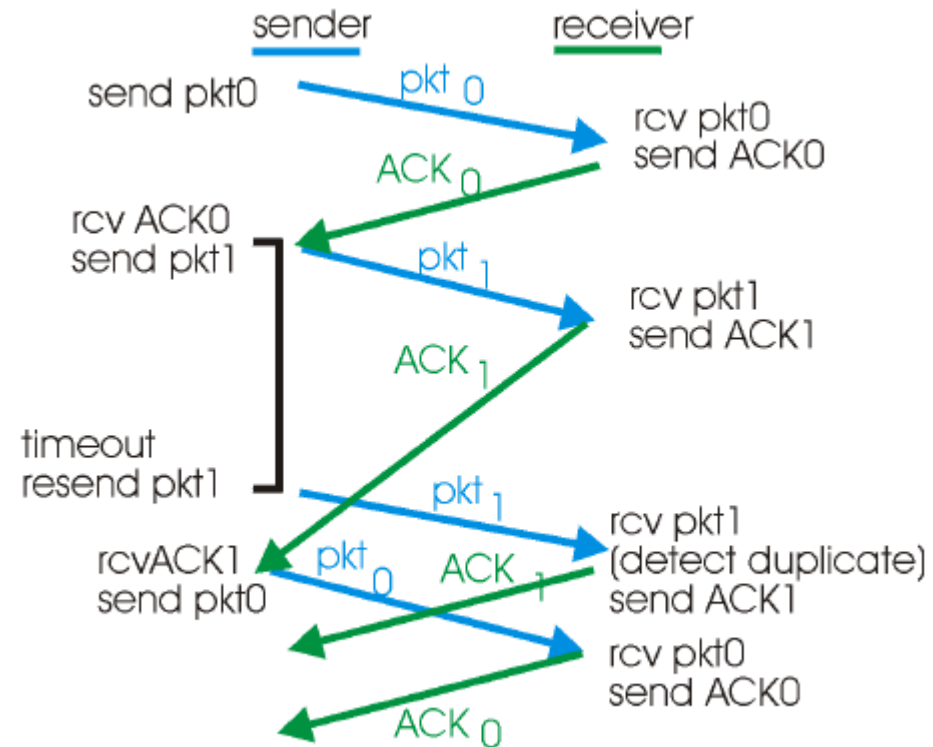
(b) lost packet

Rdt3.0 Operation

□ No ACK



□ Premature timeout



Performance of Rdt3.0

□ **Stop-and-wait protocol:** rdt3.0 works, but very low performance

- Ex) 1 Gbps link, 15 ms prop. delay, 8000 bit packet


$$d_{trans} = \frac{L}{R} = \frac{8000\text{bits}}{10^9\text{bps}} = 8\text{microseconds}$$

- **U_{sender} : utilization** – fraction of time sender busy sending

$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

- 1KB pkt every 30 msec → 33kB/s throughput over 1 Gbps link

Rdt3.0: stop-and-wait operation

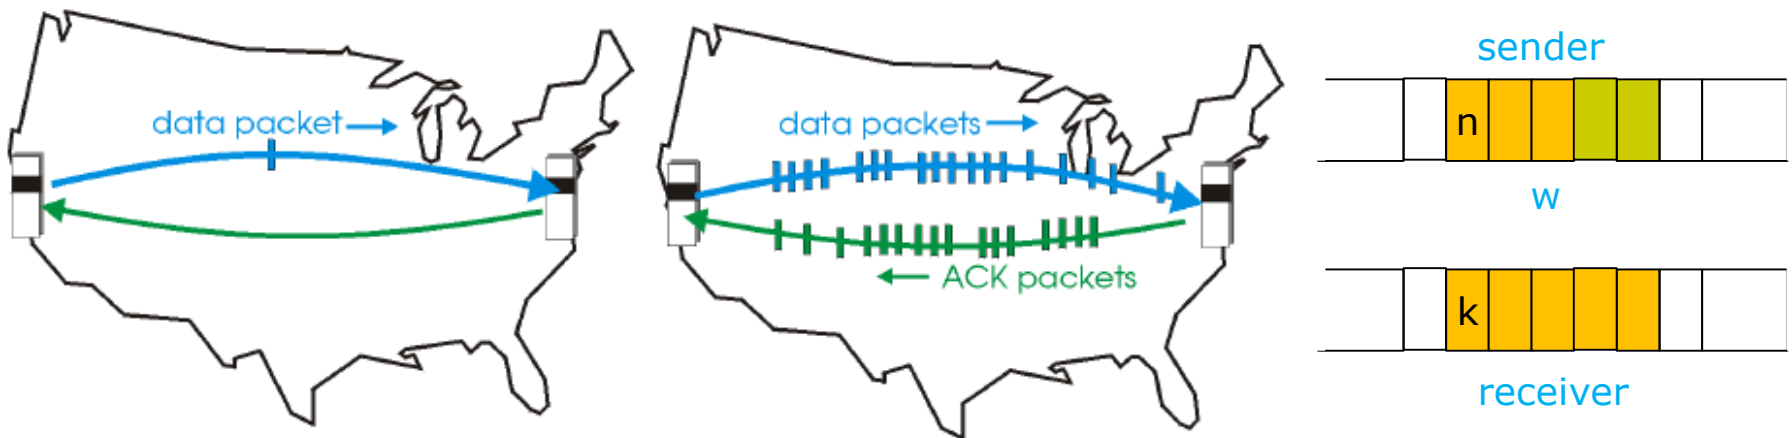


$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

Pipelined Protocols

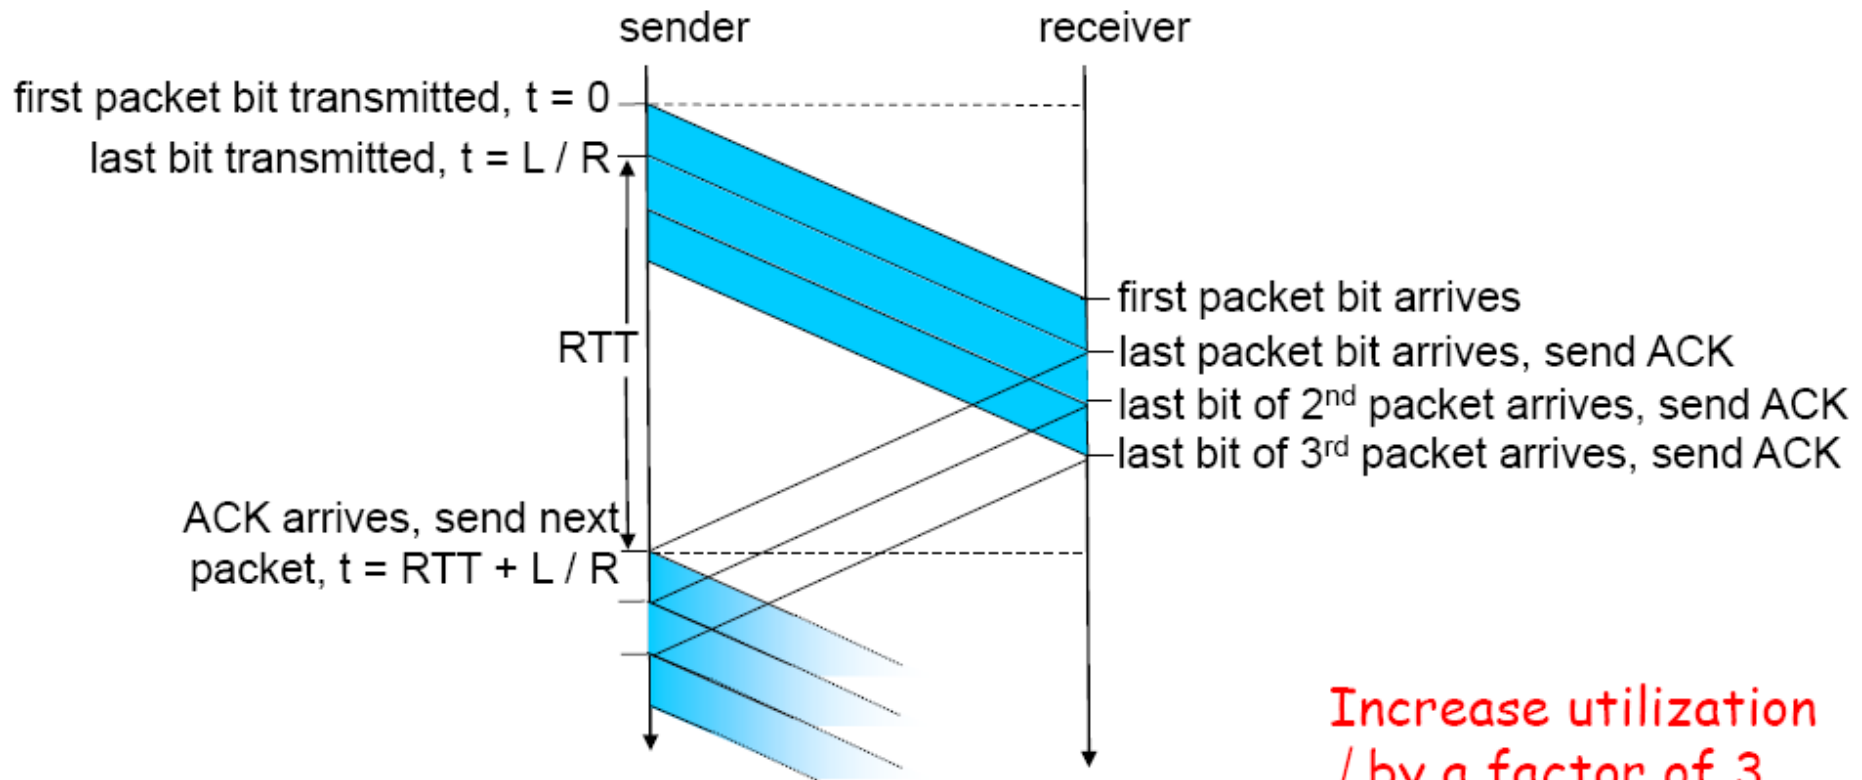
□ Pipelining (sliding window) protocol :

- sender allows multiple, “in-flight”, yet-to-be-acknowledged pkts
- window of packets (“in-flight” and “can-transmit”) is managed
- window can slide while transmitting or receiving packets
- buffering of packets in window at sender and/or receiver



□ Two pipelined protocols: *go-Back-N, selective repeat*

Pipelining: increased utilization



$$U_{\text{sender}} = \frac{3 * L / R}{RTT + L / R} = \frac{.024}{30.008} = 0.0008$$

Increase utilization
/ by a factor of 3

Pipelining Protocols

□ Go-back-N:

- Sender can have up to N un-acked packets in pipeline
- Rcvr receives packets only in the order of the original sequence
- Rcvr sends cumulative ACKs
- Sender keeps timer for the oldest un-acked packet (the first packet in the window)
 - If timer expires, retransmit all un-acked packets

Pipelining Protocols


□ Selective Repeat:

- Sender can have up to N un-acked packets in pipeline
- Rcvr can receive packets (in the window) in an arbitrary order
- Rcvr **acks individual packets**
- Sender maintains **timer for each un-acked packet**
 - When timer expires, **retransmit only un-acked packet**

Go-Back-N

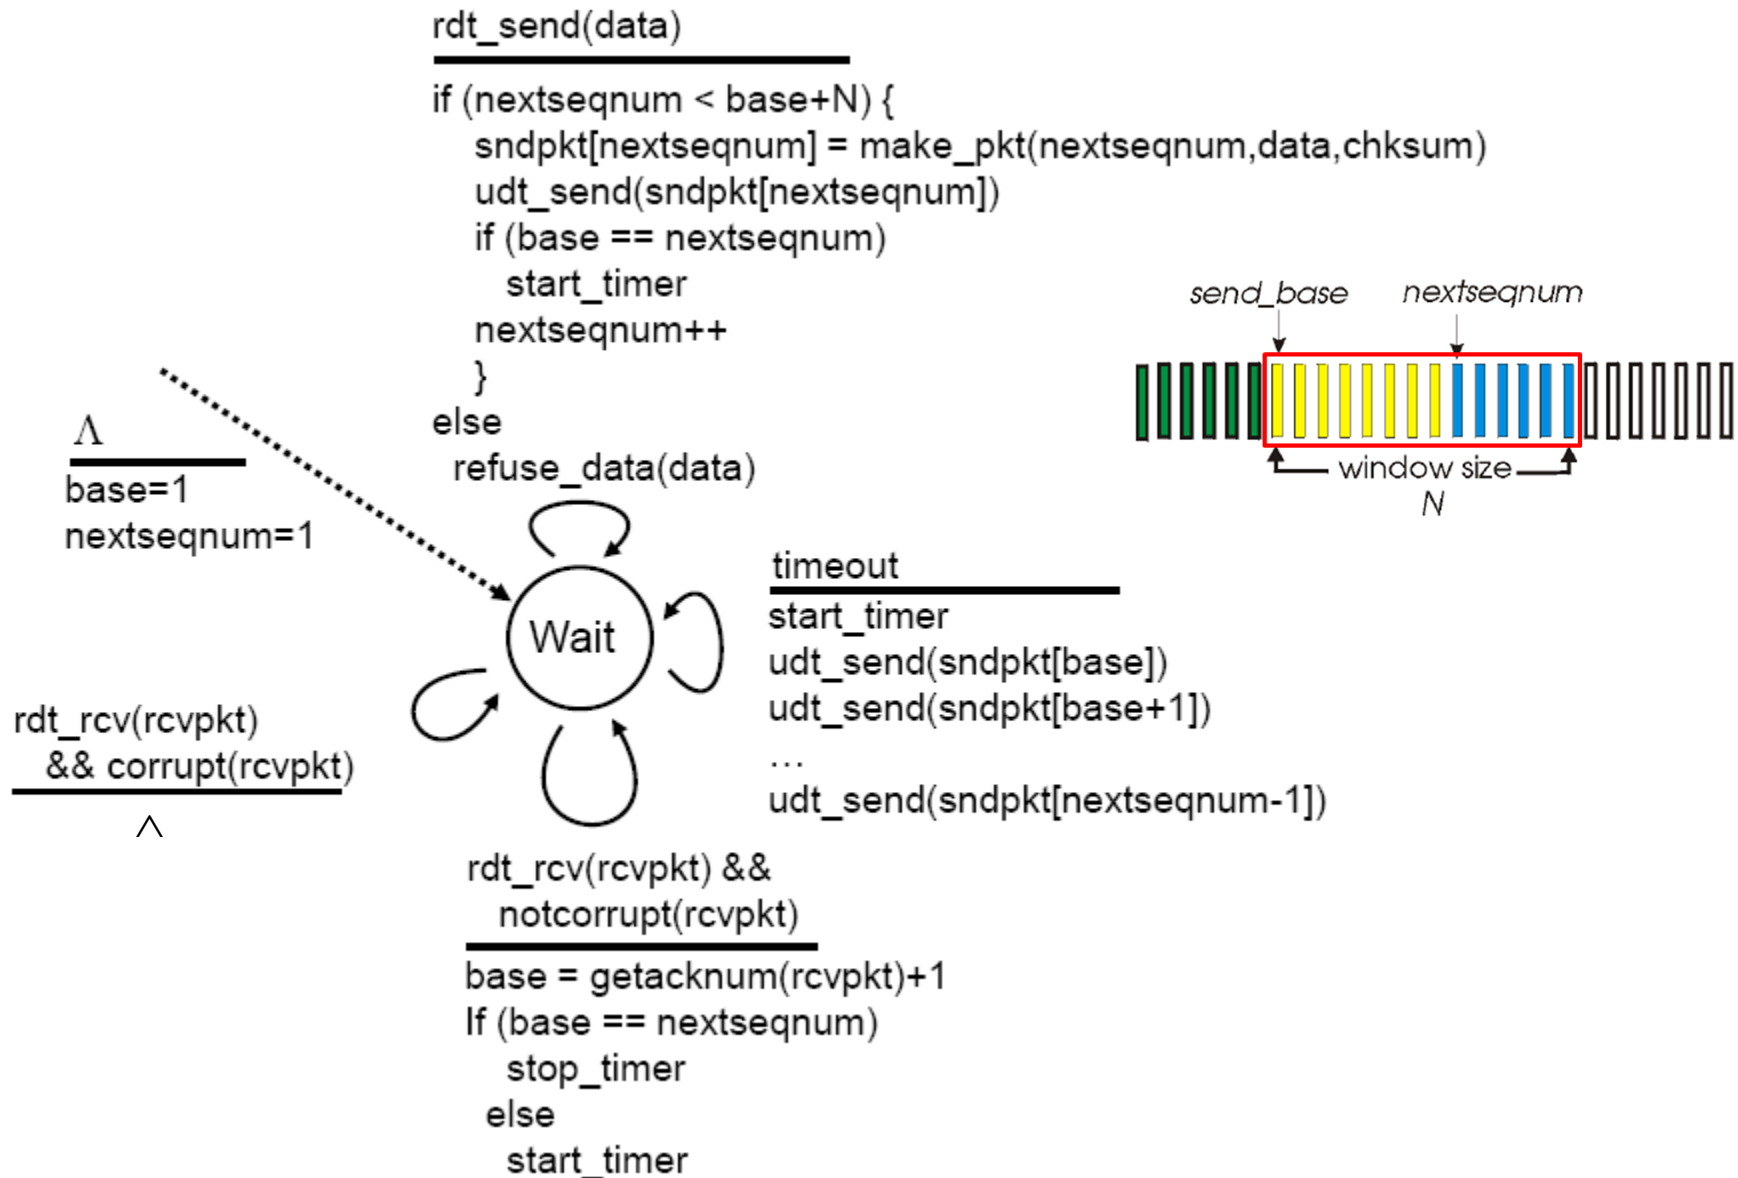
□ Sender:

- k-bit seq # in pkt header
- “window” of up to N, consecutive un-ack’ed pkts allowed


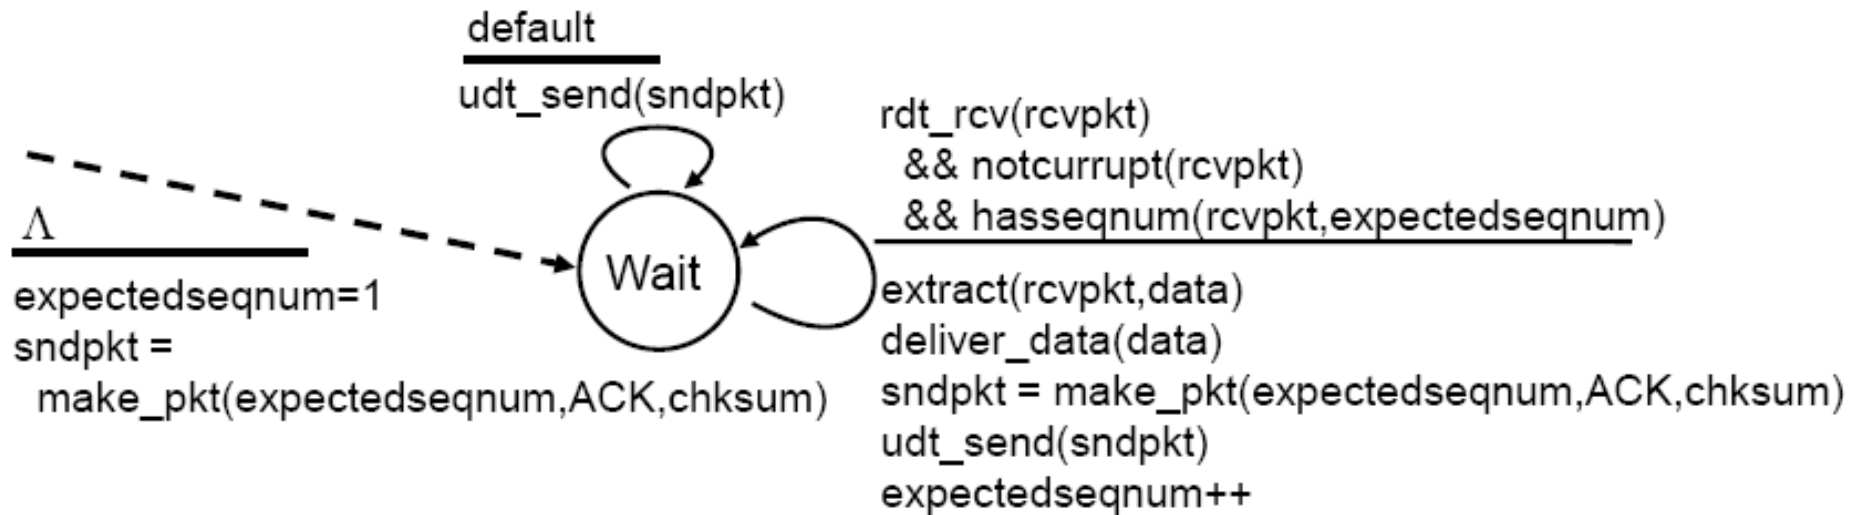


- **cumulative ACK** : ACK(n): ACKs all pkts up to, including seq # n
- **timer for the first packet in the window**
- **timeout(n)**: retransmit pkt n and all un-ack'ed pkts in window

GBN: sender extended FSM




GBN: receiver extended FSM



GBN: receiver extended FSM

□ Receiver:

- ACK-only: always send ACK for correctly-received pkt with highest in-order seq #
- may generate duplicate ACKs
- need only remember expectedseqnum
- out-of-order pkt:
 - discard (don't buffer) -> no receiver buffering!
 - Re-ACK pkt with highest in-order seq # ← duplicate ACK

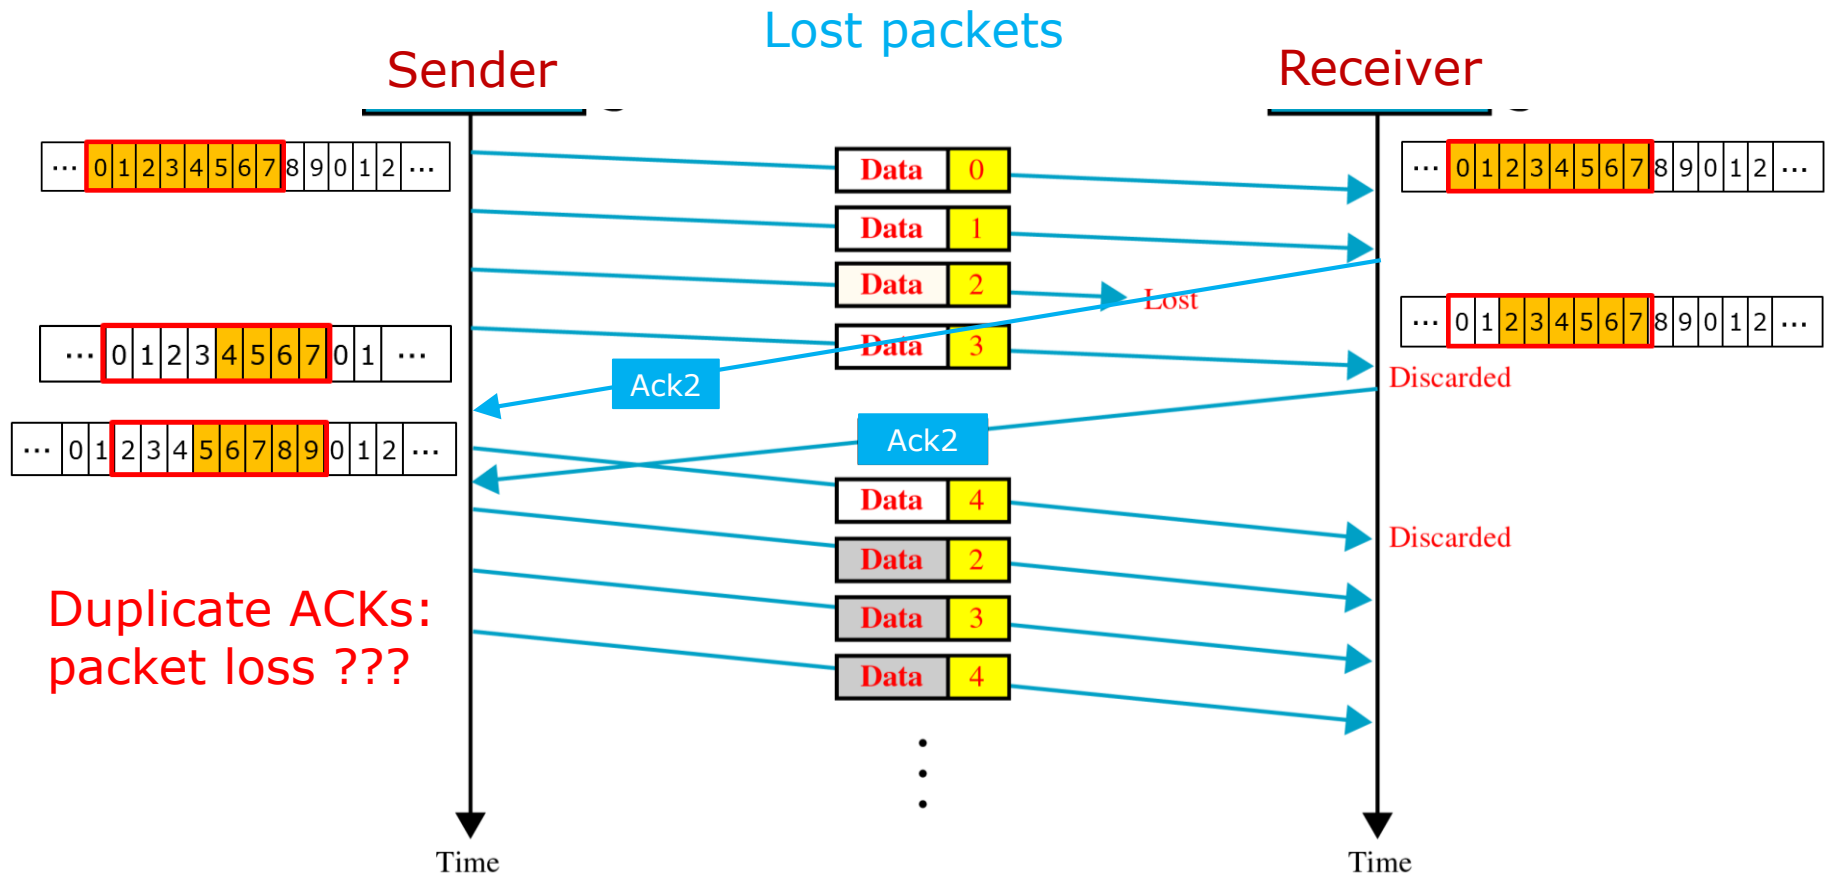


□ Go-back- n

-
- The diagram illustrates a Stop-and-Wait protocol scenario with packet errors and retransmission. It shows a Sender and a Receiver with sequence buffers and a timeline of packet transmissions.
- Sequence Buffers:**
- Sender Buffer:**
 - Initial state: ... 0 1 2 3 4 5 6 7 8 9 0 1 2 ... (Packets 0-7 are highlighted in yellow, 8 is red, 9-12 are white)
 - After packet 4: ... 0 1 2 3 4 5 6 7 8 9 0 1 2 ... (Packets 0-7 are highlighted in yellow, 8 is red, 9-12 are white)
 - After packet 5: ... 0 1 2 3 4 5 6 7 8 9 0 1 2 ... (Packets 0-2 are white, 3-9 are highlighted in yellow, 10 is red, 11-12 are white)
 - Receiver Buffer:**
 - Initial state: ... 0 1 2 3 4 5 6 7 8 9 0 1 2 ... (Packets 0-7 are highlighted in yellow, 8 is red, 9-12 are white)
 - After packet 5: ... 0 1 2 3 4 5 6 7 8 9 0 1 2 ... (Packets 0-2 are white, 3-9 are highlighted in yellow, 10 is red, 11-12 are white)
- Packet Transmissions and Timeline:**
- Packet 0:** Sent successfully.
 - Packet 1:** Sent successfully.
 - Packet 2:** Sent successfully.
 - Packet 3:** Sent successfully.
 - Packet 4:** Sent successfully.
 - Packet 5:** Sent successfully but discarded by the receiver due to an error.
 - Ack3:** Received by the sender but discarded due to an error.
 - Retransmission:** The sender resends packets 3, 4, and 5.
 - Successful Reception:** The receiver successfully receives packets 3, 4, and 5.
 - Acknowledgment:** The receiver sends back Ack3.
- Text Annotations:**
- Packet error:** Indicated by a red box around the 'Ack3' packet received by the sender.
 - Error, Discarded:** Label for the discarded packet 5.
 - Discarded:** Labels for the discarded Ack3 and packet 5.
 - Resent:** Labels for the retransmitted packets 3, 4, and 5.
 - Ack3:** Label for the acknowledgment packet sent by the receiver.
 - Time:** Vertical arrows on the left and right indicate the progression of time.
- Bottom Text:**
- P2 까지 에러없이 모두 받았고 다음에 P3을 받을 차례임

Pipelined Protocol: Go-back N

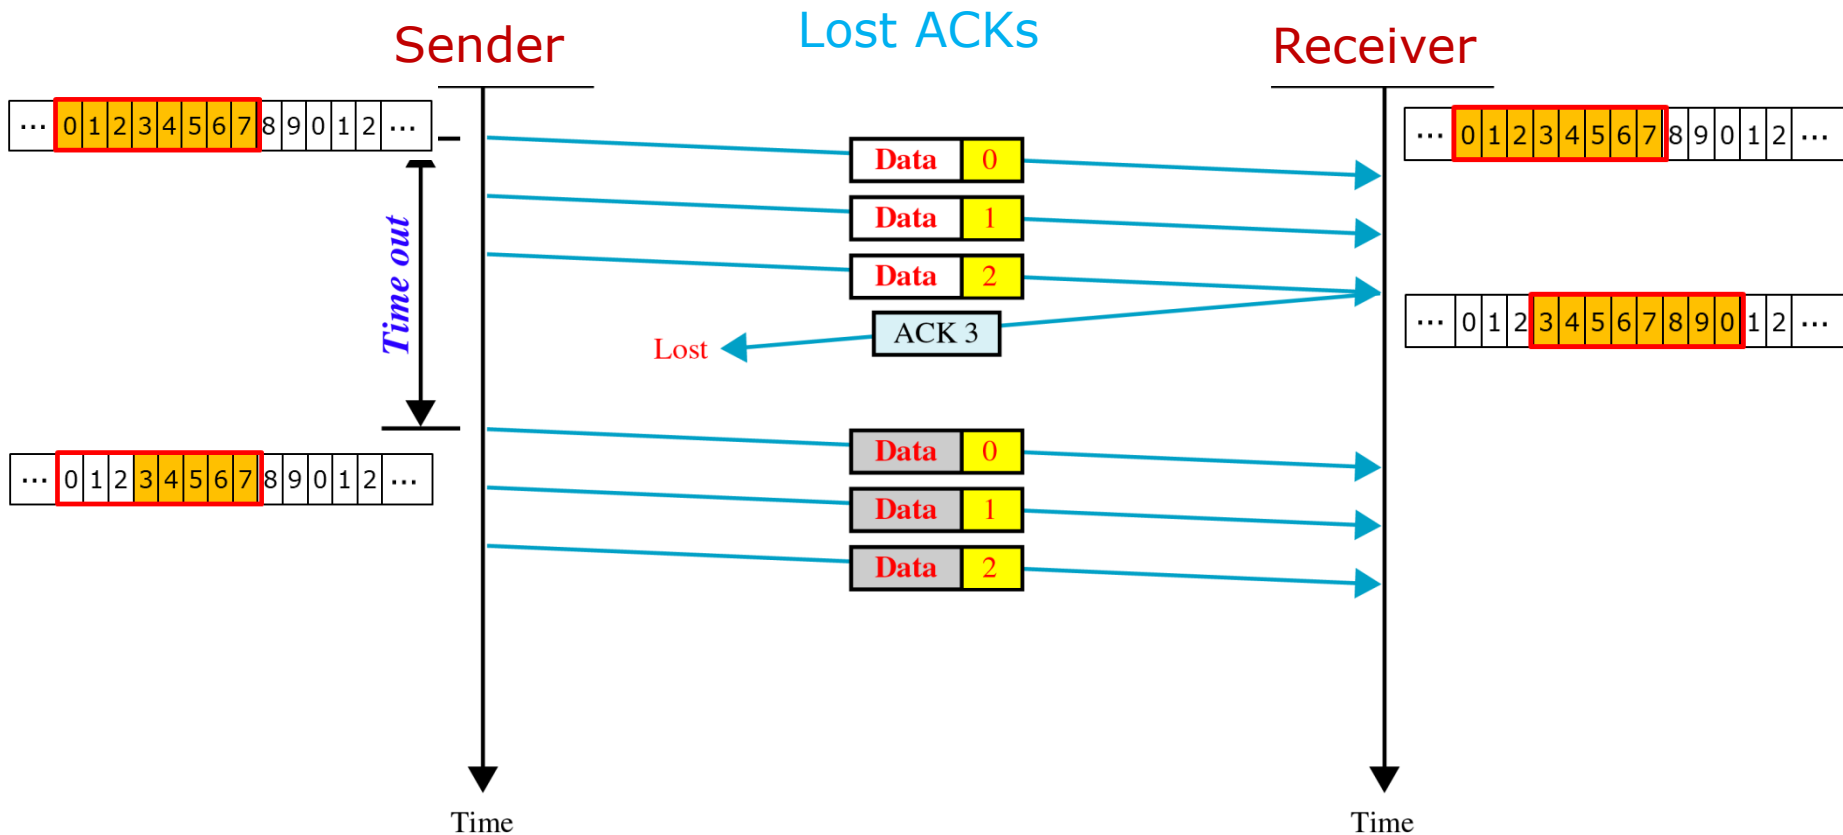
Go-back- n



Pipelined Protocol: Go-back N

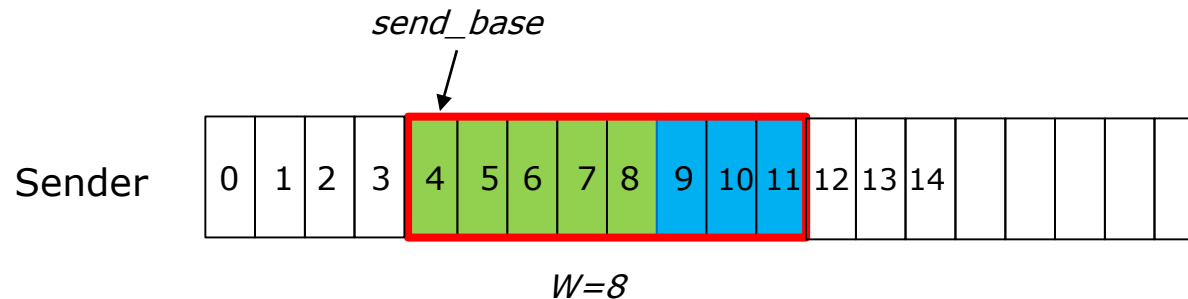
□ Go-back- n

- need packet **timeout timer**



Pipelined Protocol: Go-back N

□ Sender




ACK 5: P4까지 모두 다 받았고, 다음에 P5를 받을 차례임

What happen if the following event occur?

- 1) receive ACK 5 2) receive ACK 6 3) receive ACK 4 4) TO happen

Pipelined Protocol: Go-back N

□ Receiver



What happen if the following event occur?

- 1) receive Pkt 5
- 2) receive Pkt 6
- 3) receive Pkt 4

Pipelined Protocol: Go-back N

□ Go-back-n : window size


- Max window size: $N - 1$ (max number of seq num $- 1$)
- When the window size becomes N
 - Sender: sends P_0 and receives $ACK\ 0$
 - Sender: sends next $P_1, P_2, \dots, P_7, P_0$
 - Sender: receives $ACK\ 0$
 - The sender can not determine whether the $ACK\ 0$ is the duplicate of the previous $ACK\ 0$ or is a new one

ACK 0: P0까지 에러없이 받았음

Pipelined Protocol: Go-back N

□ Go-back-n : window size

- when $W=8$ in 3-bit SN



Pipelined Protocol: Selective-reject

□ Selective-reject (Selective-repeat) protocol


■ Receiver

- can receive packets out of sequence
- Sends “ACK n” when successfully received “packet n”
- sends “NAK n” when an error detected in “packet n”

■ Sender

- can send packets out of sequence
- keeps separate TOs for each packet transmitted
- Re-transmits only the requested packet when it receives a NAK or TO happened

Selective repeat: sender, receiver



Pipelined Protocol: Selective-reject

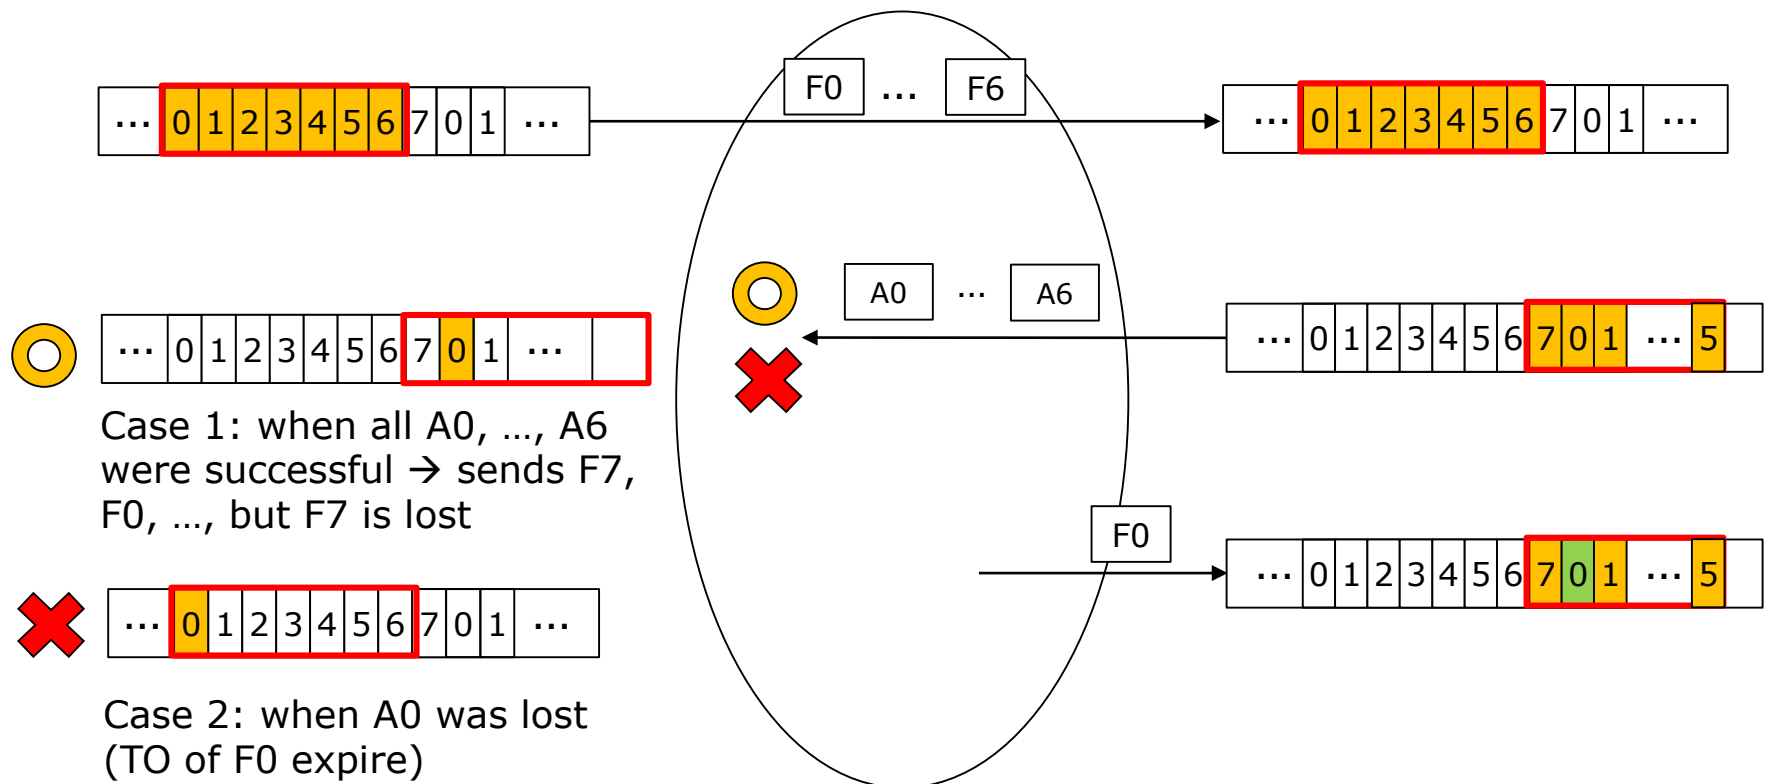
□ Selective-reject

- Max window size: $\lfloor (N+1)/2 \rfloor$
- When the window size = $(N-1)$ (e.g. $N = 8$)
 - Station A: sends **P0, P1, ..., P6** to station B
 - Station B: sends **ACK0, ACK1, ..., ACK6** (but all lost) => expands the window to accept P7, P0, P1, ..., P5
 - Station A: **TO timer of P0** expires and retransmits **P0**
 - Station B thinks the received **P0** as a new one and accepts it (wrong !!)

Pipelined Protocol: Selective-reject

□ Selective-reject : window size

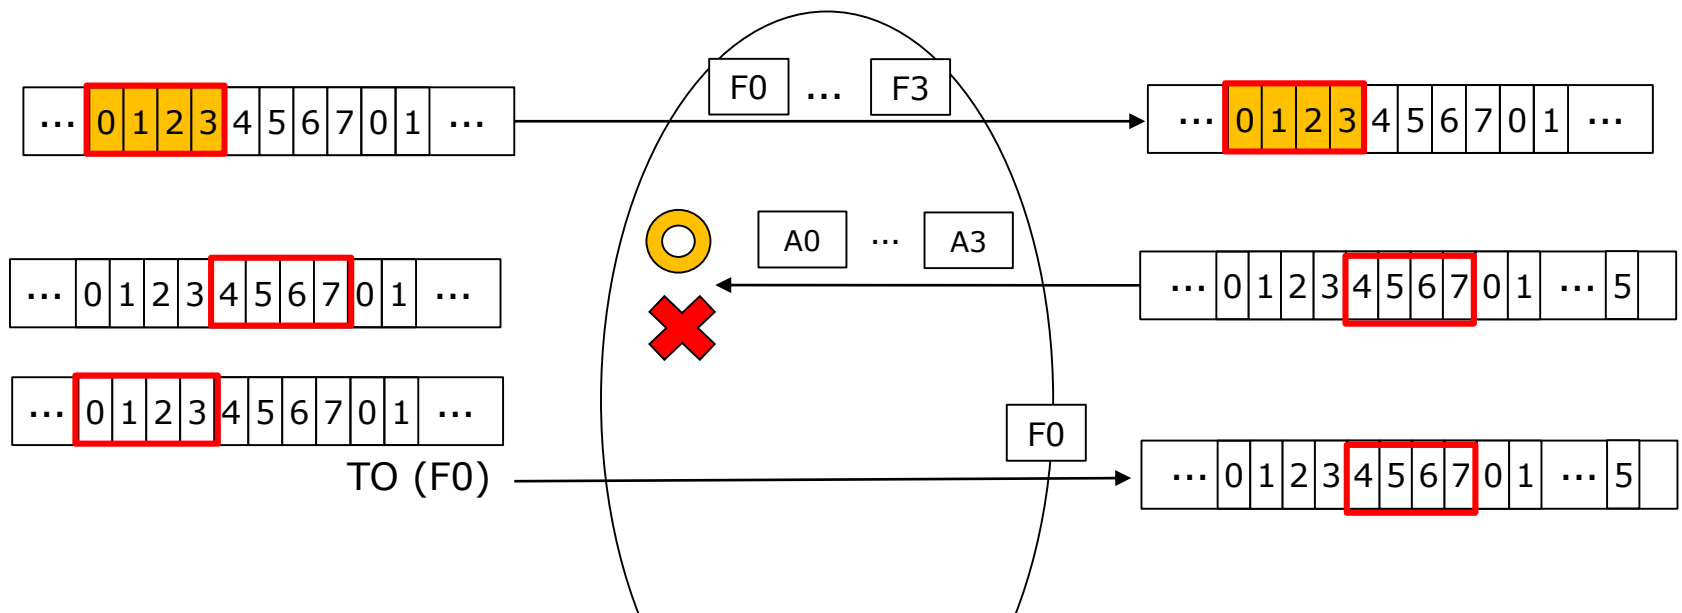
- Problem: when $W=7$ in 3-bit SN



Pipelined Protocol: Selective-reject

□ Selective-reject : window size


- when $W=7$ in 3-bit SN



There has to be no packet with the same ID between the current window and the window sliding after receiving all packets in the window successfully

Selective Repeat

□ Sender



ACK 4: P4를 성공적으로 받았음

What happen if the following event occur?


1) ACK4

2) ACK6

3) TO of P4

Selective Repeat

□ Receiver



What happen if the following event occur?

- 1) receive Pkt 6
- 2) receive Pkt 8
- 3) receive Pkt 6 (error)
- 4) receive Pkt 6 after Pkt 8
- 5) receive Pkt 5

Selective Repeat

sender

data from above :

- ❑ if next available seq # in window, send pkt

timeout(n):

- ❑ resend pkt n, restart timer

ACK(n) in [sendbase, sendbase+N]:

- ❑ mark pkt n as received
- ❑ if n smallest unACKed pkt, advance window base to next unACKed seq #

receiver

pkt n in [rcvbase, rcvbase+N-1]

- ❑ send ACK(n)
- ❑ out-of-order: buffer
- ❑ in-order: deliver (also deliver buffered, in-order pkts), advance window to next not-yet-received pkt

pkt n in [rcvbase-N, rcvbase-1]

- ❑ ACK(n)

otherwise:

- ❑ ignore

Chap 3

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- TCP congestion control

TCP: Overview

[RFCs: 793, 1122, 1323, 2018, 2581]

□ point-to-point connection-oriented

- one sender, one receiver
- 3-way handshaking: initialize sender and receiver state before data exchange

□ pipelined flow control:

- TCP congestion and flow control set window size

□ *sender & receiver side buffering*

□ reliable, in-order byte stream:


- no message boundaries

□ full duplex data transmission:

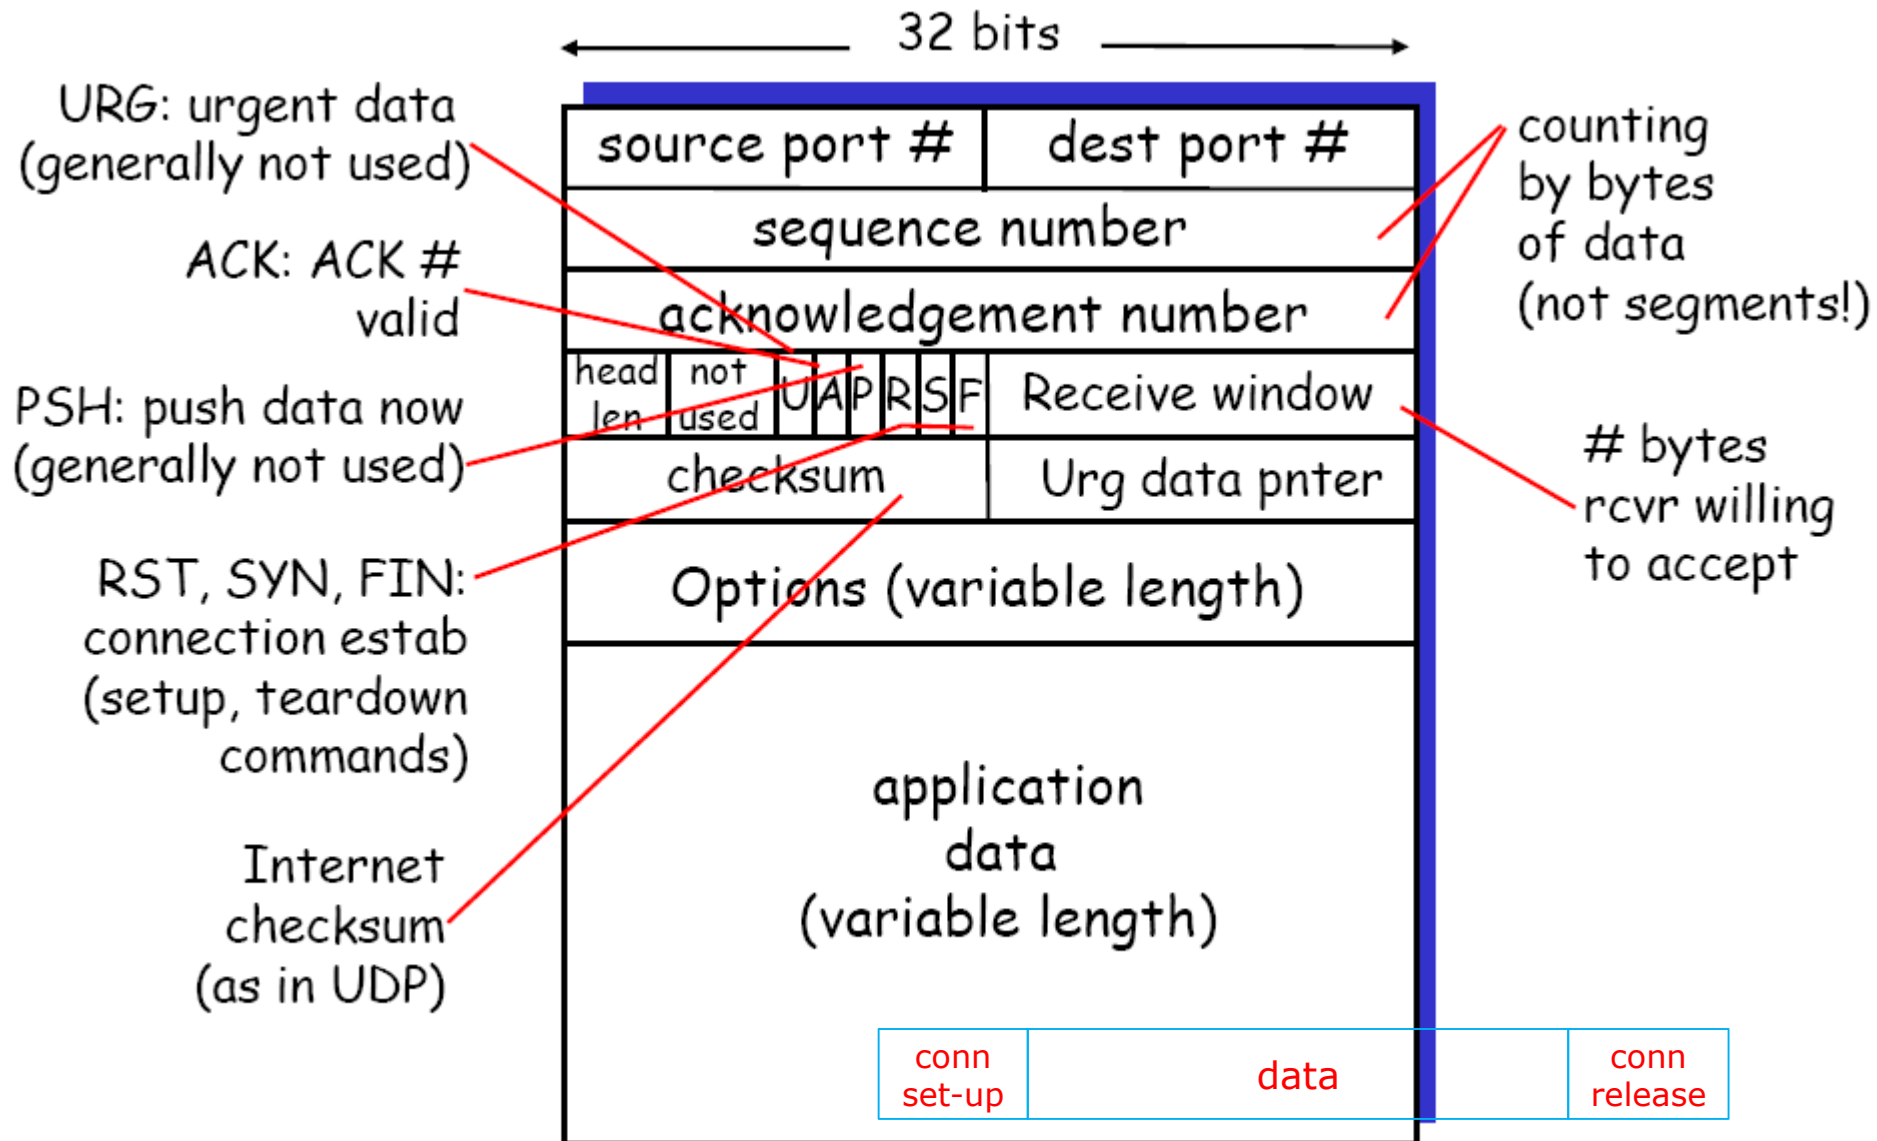
- bi-directional data flow in same connection
- MSS: maximum segment size

□ flow controlled :

- sender will not overwhelm receiver



TCP Segment Format



TCP Seq. #'s and ACKs

Seq. #'s:

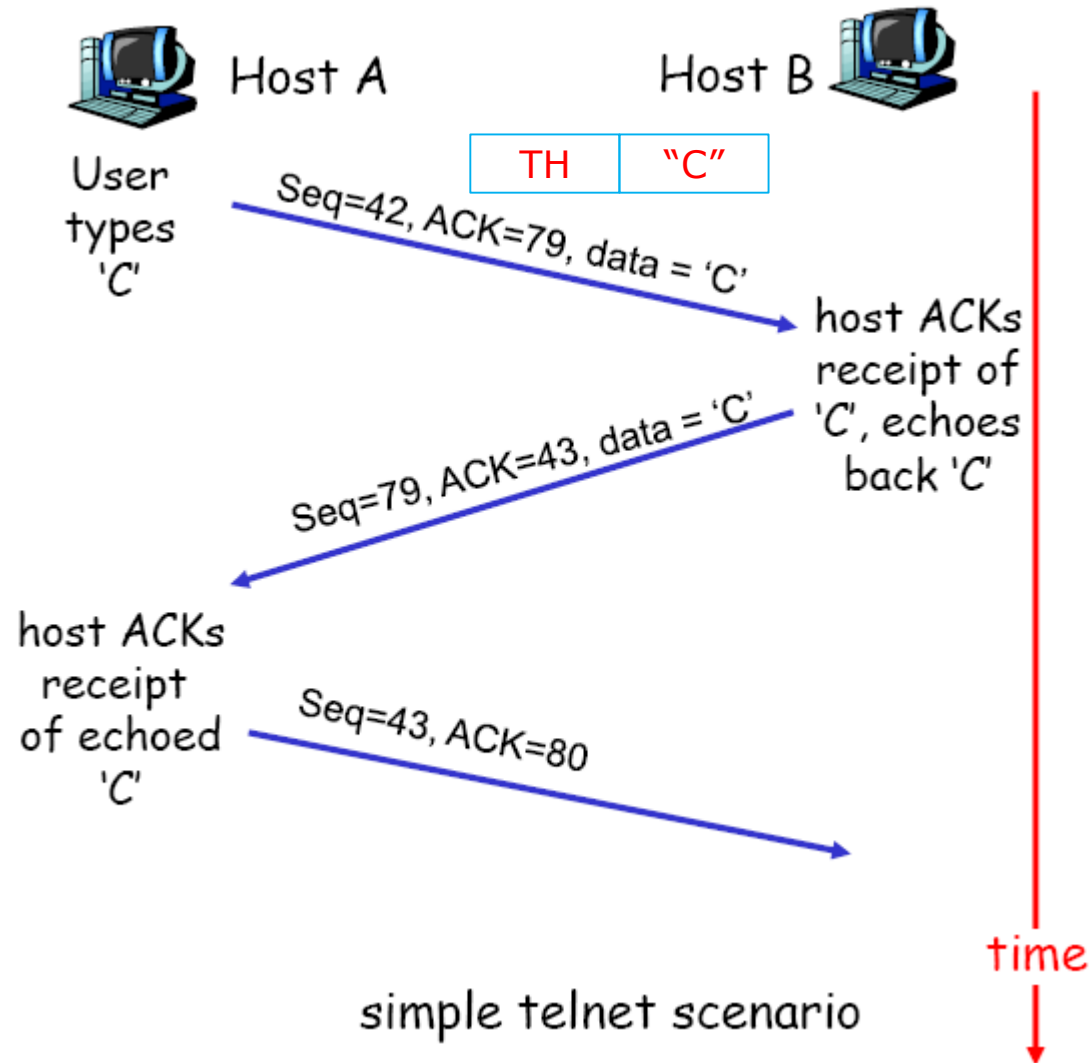
- byte stream “number” of first byte in segment’s data

ACKs:

- seq # of next byte expected from other side
- cumulative ACK

Q: how receiver handles out-of-order segments

- A: TCP spec doesn’t say, - up to implementor



TCP Round Trip Time and Timeout

Q: how to set timeout value?


□ must be $TO > RTT$

- but RTT is variable

□ too short: premature timeout

- unnecessary retransmissions

□ too long: slow reaction to segment loss





TCP Round Trip Time and Timeout

Q: how to estimate RTT?

□ **SampleRTT**: measured time from segment transmission until ACK receipt

- ignore retransmissions




TCP Round Trip Time and Timeout

- SampleRTT will vary, want estimated RTT “smoother”
 - average several recent measurements, not just current SampleRTT

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- Exponential weighted moving average
- influence of past sample decreases exponentially fast
- typical value: $\alpha = 0.125$

Example RTT estimation



TCP Round Trip Time and Timeout

Setting the timeout

- **TO = EstimatedRTT** plus “safety margin”
 - larger variation in **EstimatedRTT** → larger safety margin

- DevRTT : weighted average of | **SampleRTT-deviation** |

$$\text{DevRTT} = (1-\beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically, $\beta = 0.25$)

- Setting timeout interval

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$


TCP reliable data transfer

- Pipelined segments: [sliding window protocol](#)
- [Cumulative ACKs](#); ACK n means receiver received up to $(n-1)$
- Receiver receives out-of-order segments
- TCP uses [single retransmission timer; the first segment in window](#)
- Retransmissions are triggered by:
 - [timeout](#) events, [duplicate ACKs](#)
 - [sender retransmits only the lost segment](#)
- Initially consider simplified TCP sender:
 - ignore duplicate ACKs, flow control, congestion control

Simplified TCP sender events:

□ data received from application

- maintains a sliding window for pipelined transmission



- create segment with `cur_seq #`
- `cur_seq #` is **byte-stream num** of first byte in segment
- start timer if not already running
- expiration interval: `TimeoutInterval`

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$


Simplified TCP sender events:

□ timeout:

- retransmit segment that caused timeout
- restart timer

□ ACK rcvd: ACK k

- If acknowledges previously un-acked segments
 - Update send_base
 - start timer if there are outstanding segments




TCP Sender

NextSeqNum = InitialSeqNum

SendBase = InitialSeqNum

```
loop (forever) {  
    switch(event)
```





event: data received from application above
create TCP segment with sequence number NextSeqNum
if (timer currently not running)
 start timer
pass segment to IP
NextSeqNum = NextSeqNum + length(data)

event: timer timeout
retransmit not-yet-acknowledged segment with
 smallest sequence number
start timer


event: ACK received, with ACK field value of y
if (y > SendBase) {
 SendBase = y
 if (there are currently not-yet-acknowledged segments)
 start timer
}

```
} /* end of loop forever */
```

TCP: retransmission scenarios



TCP retransmission scenarios



Cumulative ACK scenario

TCP Receiver: ACK Generation


Event at Receiver	TCP Receiver action
<u>Arrival of in-order segment</u> with expected seq #. All data up to expected seq # already ACKed	<u>Delayed ACK</u> . Wait up to 500ms for next segment. If no next segment, send ACK
<u>Arrival of in-order segment</u> with expected seq #. One other segment has <u>ACK pending</u>	Immediately send single cumulative ACK, ACKing both in-order segments
<u>Arrival of out-of-order segment</u> higher-than-expect seq. # . Gap detected	Immediately send <u>duplicate ACK</u> , indicating seq. # of next expected byte
<u>Arrival of segment that partially or completely fills gap</u>	Immediate send ACK, provided that segment starts at lower end of gap

TCP Receiver: ACK Generation

□ In-order segment with no ACK pending

■ Receive Pkt 5

- deliver Pkt 5 to appl
- $\text{recv_base} \leftarrow 6$ and delay (ACK 6)




□ In-order segment with 1-ACK pending

■ Receive Pkt 6

■ 1 delayed ACK: (ACK 6)


- deliver Pkt 6 to appl
- Send ACK 7 and $\text{recv_base} \leftarrow 7$



TCP Receiver: ACK Generation

□ Out-order segment with higher seq-no than recv_base


- Receive Pkt 6



- Buffer Pkt 6 (mark Pkt 6) and send ACK 5 (dup-ACK)

□ Out-order segment with higher seq-no than recv_base

- Receive Pkt 7




- Buffer Pkt 7 (mark Pkt 7) and send ACK 5 (dup-ACK)

TCP Receiver: ACK Generation

□ In-order segment that fills gap


- Receive Pkt 5



- Deliver Pkt 5, Pkt 6, and pkt 7 to upper layers
- send ACK 8 and $\text{recv_base} \leftarrow 8$

Fast Retransmit

- Timeout value relatively large: cause long delay before resending lost packet
- Detect segment loss via duplicate ACKs
 - Sender sends packets back-to-back
 - If segment is lost, there will likely be many duplicate ACKs
- If sender receives 3 duplicate ACKs, that segment may be lost with high probability
 - TCP Reno (fast retransmit): resend segment before timer expires



Fast Retransmit Algorithm

event: ACK received, with ACK field value of y

if ($y > \text{SendBase}$) {

$\text{SendBase} = y$

 if (there are currently not-yet-acknowledged segments)
 start timer

}


else {

 increment count of dup ACKs received for y

 if (count of dup ACKs received for $y = 3$) {

 resend segment with sequence number y

}




a duplicate ACK for
already ACKed segment

fast retransmit

TCP Flow Control

- receiver side buffering: TCP receiver has a receive buffer
- appl process may be slow at reading from buffer



- need to match the speed b/w sending rate and the receiving appl's drain rate
- Flow control
 - sender won't overflow receiver's buffer by transmitting too much, too fast

TCP Flow control



□ spare room in buffer:

= **RcvWindow**

□ Receiver sends **RcvWindow** in segments (**receiveWindow** in TH):
piggybacking

□ **Sender limits un-ACKed data to RcvWindow**

■ guarantees receive buffer doesn't overflow



TCP Connection Management

Connection set-up: 3-way handshaking


Step 1: client host sends **TCP SYN** segment to server

- specifies initial seq # (**ISN**)
- no data

Step 2: server host receives SYN, replies with **SYNACK** segment

- server allocates buffers
- specifies server initial seq. #

Step 3: client receives SYNACK, replies with **ACK** segment, which may contain data



TCP Connection Management


Closing a connection:

client closes socket:

```
clientSocket.close();
```

Step 1: client end system sends **TCP FIN** control segment to server

Step 2: server receives FIN, replies with **ACK**; Closes connection, sends FIN




TCP Connection Management

Step 3: client receives FIN, replies with ACK

- Enters “**timed wait**” - will respond with ACK to received FINs

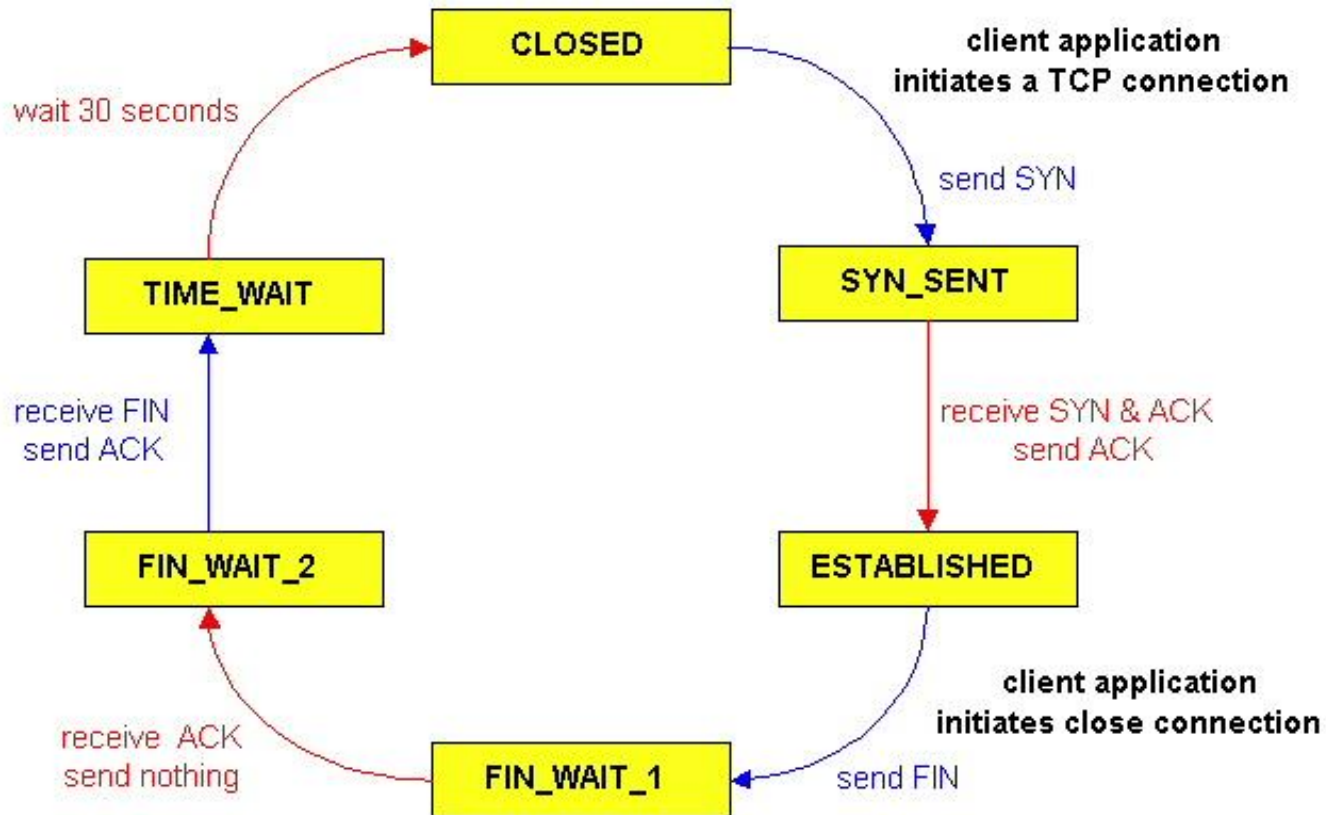
Step 4: server, receives ACK
Connection closed

Note: with small modification, can handle simultaneous FINs



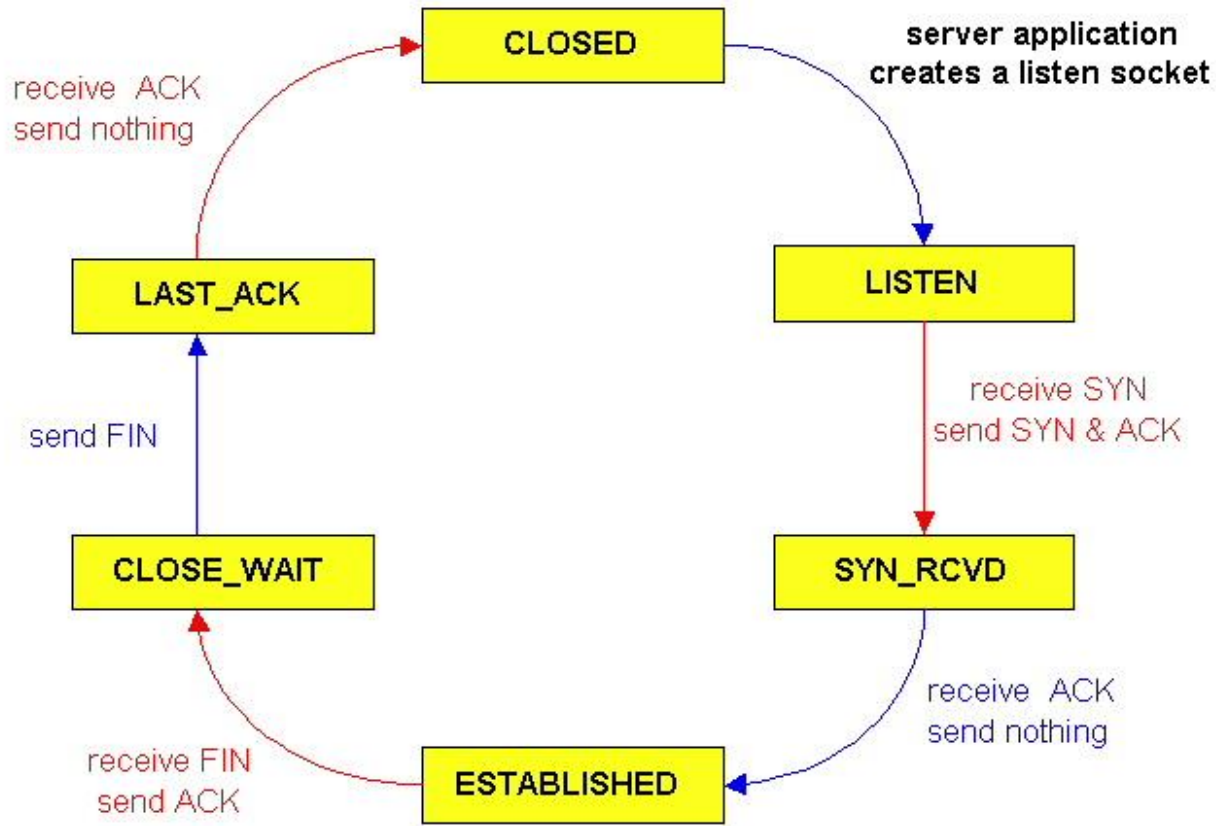
TCP Connection Management

□ TCP client life cycle



TCP Connection Management

□ TCP server life cycle



Chap 3

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- TCP congestion control




Principles of Congestion Control

Congestion:

- informally: “too many sources sending too much data too fast for *network* to handle”
- different from flow control!
- manifestations:
 - long delays (queueing in router buffers)
 - lost packets (buffer overflow at routers)

Congestion: scenario 1


- two senders, two receivers
- one router, infinite buffers, output link capacity C
- no retransmission



- large delays when congested
- maximum achievable throughput




Congestion: scenario 2

- one router, *finite* buffers
- sender retransmission of lost packet



Congestion: scenario 2

- (1) When A sends a packet only when a buffer is free: $\lambda_{in} = \lambda_{out}$ (goodput)
- (2) “perfect” retransmission only when loss: $\lambda'_{in} > \lambda_{out}$
- (3) retransmission of delayed (not lost) packet makes λ'_{in} larger (than perfect case) for same λ_{out}




costs of congestion:


- more work (retransmission) for given “goodput”
- unneeded retransmissions: link carries multiple copies of pkt

Congestion: scenario 3

- four senders, multihop paths
- timeout/retransmit
- Q: what happens as λ_{in} and λ'_{in} increase ?



Congestion: scenario 3



□ Another “cost” of congestion:

- when packet dropped, any “upstream transmission capacity used for that packet was wasted

Approaches towards congestion control

Network-assisted congestion control:


- routers provide feedback to end systems
 - single bit indicating congestion (TCP/IP ECN, ATM)
 - explicit rate sender should send at

End-end congestion control:

- no explicit feedback from network
- congestion inferred from end-system observed loss, delay
- approach taken by TCP

Approaches towards congestion control

Network-assisted congestion control:



Chap 3


- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- TCP congestion control

TCP Congestion Control

Additive increase, multiplicative decrease (AIMD)

□ *Approach:* increase transmission rate (congestion window size: **cwnd**), probing for usable bandwidth, until loss occurs

- *additive increase:* increase **cwnd** by 1 MSS every RTT until loss detected
- *multiplicative decrease:* cut **cwnd** in half after loss



TCP Congestion Control

- sender limits transmission:
(LastByteSent -
LastByteAked) \leq cwnd

- Roughly,

$$\text{rate} = \frac{\text{cwnd}}{\text{RTT}} \text{ Bytes/sec}$$

- **cwnd** is dynamic, indicates how much we can transmit before congestion

How does sender perceive congestion?

- loss event : timeout *or* 3 duplicate ACKs
- TCP sender reduces rate (cwnd) after loss event

three mechanisms: AIMD

- slow start
- congestion avoidance
- congestion control

TCP Slow Start


- When connection begins, **cwnd = 1*MSS**
 - Example: MSS = 500 bytes & RTT = 200 msec
 - initial rate = 20 kbps
- available bandwidth may be \gg MSS/RTT
 - desirable to quickly ramp up to available rate
- When connection begins, **increase rate exponentially fast until first loss event**

TCP Slow Start

□ **slow start**: increase rate exponentially until loss event

- increments **cwnd** by $(1 * MSS)$ for each ACK received until **threshold** (**ssthreshold**)
- roughly double **cwnd** for each RTT

□ initial rate is slow but ramps up exponentially fast



TCP Congestion Avoidance


□ congestion avoidance: when
($\text{cwnd} \geq \text{ssthreshold}$)

- increments **cwnd** by $\text{MSS} * (\text{MSS} / \text{cwnd})$ for each ACK received
- increase rate linearly for each RTT until loss event
- when loss happens, goes to congestion control step

$\text{cwnd} = 4 * \text{MSS}$

$\text{cwnd} = 4\text{MSS} + \text{MSS}/4$

$\text{cwnd} = 5 * \text{MSS}$



TCP Congestion Control

□ loss: 3 dup ACKs

- $ssthreshold = cwnd/2$
- $cwnd = ssthreshold$
- goes to congestion avoidance: window then grows linearly


□ loss: timeout event

- $ssthreshold = cwnd/2$
- $cwnd = 1 * MSS$
- goes to slow start

Philosophy:

- 3 dup ACKs indicates network capable of delivering some segments
- timeout indicates a "more alarming" congestion scenario

TCP Congestion Control



Implementation:

- Variable Threshold
- At loss event, Threshold is set to $1/2$ of **cwnd** just before loss event

Summary: TCP Congestion Control


- When **cwnd** is below **Threshold**, sender in **slow-start** phase, window grows exponentially
- When **cwnd** is above **Threshold**, sender is in **congestion-avoidance** phase, window grows linearly
- When a **3 duplicate ACK** occurs, **Threshold** set to **cwnd/2** and **cwnd** set to **Threshold**
- When **timeout** occurs, **Threshold** set to **cwnd/2** and **cwnd** is set to 1 MSS

TCP sender congestion control

State	Event	TCP Sender Action	Commentary
Slow Start (SS)	ACK receipt for previously unacked data	$\text{CongWin} = \text{CongWin} + \text{MSS}$, If ($\text{CongWin} > \text{Threshold}$) set state to "Congestion Avoidance"	Resulting in a doubling of CongWin every RTT
Congestion Avoidance (CA)	ACK receipt for previously unacked data	$\text{CongWin} = \text{CongWin} + \text{MSS} * (\text{MSS} / \text{CongWin})$	Additive increase, resulting in increase of CongWin by 1 MSS every RTT
SS or CA	Loss event detected by triple duplicate ACK	$\text{Threshold} = \text{CongWin} / 2$, $\text{CongWin} = \text{Threshold}$, Set state to "Congestion Avoidance"	Fast recovery, implementing multiplicative decrease. CongWin will not drop below 1 MSS.
SS or CA	Timeout	$\text{Threshold} = \text{CongWin} / 2$, $\text{CongWin} = 1 \text{ MSS}$, Set state to "Slow Start"	Enter slow start
SS or CA	Duplicate ACK	Increment duplicate ACK count for segment being acked	CongWin and Threshold not changed


TCP Throughput

- average throughput of TCP as a function of window size and RTT?
 - Let W be the window size when loss occurs
 - When window is W , throughput is W/RTT
- Just after loss, window drops to $W/2$, throughput to $W/2RTT$
- Average throughput:
 $0.75 * (W/RTT)$



TCP Fairness


- **Fairness goal:** if K TCP sessions share a bottleneck link of bandwidth R , each should have average rate of R/K



Why is TCP fair?

Two competing sessions:

- Additive increase gives slope of 1, as throughput increases
- multiplicative decrease decreases throughput proportionally



Fairness more

Fairness and parallel TCP connections

- nothing prevents app from opening parallel connections between 2 hosts
- Web browsers use parallel connections
- Example: link of rate C supporting 9 applications
 - 8 appl's ask for 1 TCP conn and 1 appl asks for 2 conn
 - 8 appl's with 1 TCP conn, gets rate $C/10$ for each appl
 - one appl with 2 TCP conns, gets $C/5$

