

Computer Organization

Lecture 2 - Performance

Reading: 1.4



Photo credit:
Jurvetson via flickr



Roadmap for the term: major topics

- ▶ Computer Systems Overview
- ▶ **Performance** ◀
- ▶ Assembly Language
- ▶ Instruction sets (and Software)
- ▶ Logic & Arithmetic
- ▶ Processor Implementation
- ▶ Memory Systems

Performance Outline

- ▶ **Motivation** **3**
- ▶ Defining Performance
- ▶ Common Performance Metrics
- ▶ Benchmarks
- ▶ Amdahl's Law

Motivation

- ▶ Goal: Learn to “**measure, summarize and report**” performance of a computer system
- ▶ Why study performance?
 - ▶ To make intelligent decisions when choosing a system
 - ▶ **To make intelligent decisions when designing a system**
 - ▶ **Understand impact of implementation decisions**
- ▶ Challenges
 - ▶ How do we measure performance **accurately**?
 - ▶ How do we compare performance **fairly**?

What's a good measure of performance?

- ▶ **Response Time**

- ▶ How long does it take to complete a single task?

- ▶ **Throughput**

- ▶ **How many tasks** are completed per unit time

- ▶ **The measure we use depends on the application**

Execution Time vs. Throughput

- ▶ Analogy: passenger airplanes (book Figure 1.13)
 - **Concorde** - fastest “**response time**” for an individual user
 - **Boeing 747** - highest passenger **throughput**

Airplane	Passenger Capacity	Cruising Range (miles)	Cruising Speed (mph)	Passenger Throughput (passengers x mph)
Boeing 777	375	4630	610	228,750
Boeing 747	470	4150	610	286,700
Concorde	132	4000	1350	178,200
DC-8-50	146	8720	544	79,424

Performance Outline

- ▶ Motivation
- ▶ **Defining Performance** ◀
- ▶ Common Performance Metrics
- ▶ Benchmarks
- ▶ Amdahl's Law

Relative Performance

- ▶ For a given program on machine X:

$$\text{Performance}_X = \frac{1}{\text{Execution time}_X}$$

- ▶ Comparing performance of machines:

$\text{Performance}_X > \text{Performance}_Y$ if
 $\text{Execution Time}_X < \text{Execution Time}_Y$

Comparing Performance

► We say “**X is n times faster than Y**” if:

$$\frac{\text{Performance}_X}{\text{Performance}_Y} = n$$

$$\frac{\text{Performance}_X}{\text{Performance}_Y} = \frac{\text{Execution time}_Y}{\text{Execution time}_X} = n$$

Example - Performance

Machine	Execution Time
A	15 seconds
B	20 seconds



► Which machine is faster?

► By how much?

Performance Outline


- ▶ Motivation
- ▶ Defining Performance
- ▶ **Common Performance Metrics** ◀
- ▶ Benchmarks
- ▶ Amdahl's Law

Elapsed Time / Execution Time

- ▶ **Elapsed Time (“Wall Clock” Time)**
 - ▶ How long does it take the program to complete?
 - ▶ **System Performance** – based on elapsed time
- ▶ **Execution Time (CPU Time)**
 - ▶ How much time the CPU spent executing the program
 - User time – time CPU spends executing program instructions
 - System time – time CPU spends in OS on behalf of program
 - ▶ **CPU Performance** – based on CPU time

Measuring CPU Time with the OS

- ▶ OS runs multiple programs
 - ▶ **Timer** interrupts programs at fixed intervals (e.g. 10ms)
 - ▶ OS decides whether to **switch** which program runs
 - ▶ Interval counter – counts intervals when program is running in user mode (u) and system mode (s)
- ▶ Issues
 - ▶ Inaccuracy – interval counter can miss changes (cancels out for long programs)



Measuring CPU Time with Unix/Linux

- ▶ Interval-based timing in the “real world”
- ▶ Linux/Unix `time` command

`% time A`
...
90.7u 12.9s 2:39 65%


↑ ↑ ↑ ←

User Time System Time Wall-clock CPU Utilization = (UT+ST) / WCT
(sec) (sec) Time

$103.6 / 159 = 65\%$

Clocks and Performance

- ▶ Clock signal - controls sequential circuit operation




Clock frequency $f_{\text{clock}} = 1/T_{\text{clock}}$

초당 클락이
몇 번 바뀌는가?

Clocks and Performance - CPU Time

► How do we relate clock to CPU Time?



$$\text{CPU time} = \frac{\text{clock cycles}}{\text{program}} \times \frac{\text{seconds}}{\text{clock cycle}}$$

- $1 \text{ GHz} = 10^9 \text{ Hz}$
- $1 \text{ ps} = 10^{-12} \text{ s}$
- $1 \text{ ns} = 10^{-9} \text{ s}$

Measuring **Elapsed Time** with Counters

- ▶ Modern microprocessors include **cycle counters**

Ex) Intel Pentium & Later Processors

- 64-bit Time Stamp Counter (TSC)
- Counts clock cycles since reset
- `rdtsc` instruction – moves TSC to {`edx`, `eax`} registers
- Measure elapsed time by
 - ✓ Reading TSC at start
 - ✓ Reading TSC at end
 - ✓ Subtract & multiply by clock period

Get the number of clocks during program execution

▶ Issues

- ▶ Context changes
- ▶ Impact of caching

Example - Clock Cycles


Machine	Execution Time	Clock Freq.
A	15 seconds	3 GHz
B	20 seconds	2.5 GHz

- ▶ How many clock cycles does A execute?

- ▶ How many clock cycles does B execute?

Clock Cycles per Instruction (CPI)

- ▶ Consider the 68HC11 ...
 - ▶ ADDA - 3 cycles (IMM) -> 5 cycles (IND, Y)
 - ▶ MUL - 10 cycles
 - ▶ IDIV - 41 cycles
- ▶ More complex processors have other issues...
 - ▶ Pipelining - parallel execution (CPI=1!), but stalls occur
 - ▶ Memory system: cache misses, page faults, etc.
- ▶ How can we combine these into an overall metric?



Clock Cycles per Instruction (CPI)

- ▶ **Average number of clock cycles per instruction**
- ▶ **Measured for an entire program**

$$\text{CPU Clock Cycles} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Average Clock Cycles}}{\text{Instruction}}$$

$$\text{CPI} = \frac{\text{Average Clock Cycles}}{\text{Instruction}} = \frac{\text{CPU Clock Cycles/Program}}{\text{Instructions/Program}}$$

Example - CPI

Machine	Clock Cycles	Instructions
A	1.35×10^{10}	2.6×10^9
B	1.2×10^{10}	3.0×10^9

- ▶ What is the CPI of A?

- ▶ What is the CPI of B?

CPI Example

- ▶ Computer A: Cycle Time = 250ps, CPI = 2.0
- ▶ Computer B: Cycle Time = 500ps, CPI = 1.2 (**Same ISA**)
- ▶ Which is faster, and by how much?

$$\begin{aligned}\text{CPU Time}_A &= \text{Instruction Count} \times \text{CPI}_A \times \text{Cycle Time}_A \\ &= \text{IC} \times 2.0 \times 250\text{ps} = \text{IC} \times 500\text{ps} \end{aligned}$$

← A is faster...

$$\begin{aligned}\text{CPU Time}_B &= \text{Instruction Count} \times \text{CPI}_B \times \text{Cycle Time}_B \\ &= \text{IC} \times 1.2 \times 500\text{ps} = \text{IC} \times 600\text{ps} \end{aligned}$$

$$\frac{\text{CPU Time}_B}{\text{CPU Time}_A} = \frac{\text{IC} \times 600\text{ps}}{\text{IC} \times 500\text{ps}} = 1.2$$

← ...by this much

The Performance Equation

► The “Iron Law” of Performance

$$\begin{aligned}\text{CPU time} &= \frac{\# \text{ instructions}}{\text{program}} \times \frac{\text{clock cycles}}{\text{instruction}} \times \frac{\text{seconds}}{\text{clock cycle}} \\ &= \frac{\text{Instruction Count} \times \text{CPI}}{\text{Clock Rate}}\end{aligned}$$

CPI in More Detail

Class 1: CPI = 3

Class 2: CPI = 5

10 Ins

30 Ins

- ▶ If ***n* different instruction classes** take different numbers of cycles per instruction (CPI)

$$\text{Clock Cycles} = \sum_{i=1}^n (\text{CPI}_i \times \text{Instruction Count}_i)$$

$$\text{CC} = 10 \times 3 + 30 \times 5 = 180$$

$$\text{CPI} = 180/40 = 4.5$$

- ▶ **Weighted average CPI**

$$\text{CPI} = \frac{\text{Clock Cycles}}{\text{Instruction Count}} = \sum_{i=1}^n \left(\text{CPI}_i \times \underbrace{\frac{\text{Instruction Count}_i}{\text{Instruction Count}}}_{\text{Relative frequency}} \right)$$

Relative frequency

$$\text{CPI} = 3 \times 10/40 + 5 \times 30/40 = 180/40 = 4.5$$

Another CPI Example

- ▶ Consider two compiled code sequences using instructions in classes A, B, C

Class	A	B	C
CPI for class	1	2	3
IC in sequence 1	20	10	20
IC in sequence 2	40	10	10

- ▶ Which sequence executes more instructions?
- ▶ Which sequence is faster?
- ▶ What is the CPI for each sequence?

CPI Example

- ▶ **Solution: calculate total cycles for each sequence, then CPI**

$$\text{Clock Cycles} = \sum_{i=1}^n (\text{CPI}_i \times \text{Instruction Count}_i)$$

Class	A	B	C
CPI for class	1	2	3
IC in sequence 1	20	10	20
IC in sequence 2	40	10	10

- ▶ **Sequence 1: IC = 50**

- Clock Cycles
 $= 20 \times 1 + 10 \times 2 + 20 \times 3 = 100$
- Avg. CPI = $100/50 = 2.0$

- ▶ **Sequence 2: IC = 60**

- Clock Cycles
 $= 40 \times 1 + 10 \times 2 + 10 \times 3 = 90$
- Avg. CPI = $90/60 = 1.5$

Performance Summary

▶ The “**BIG PICTURE**”

$$\text{CPU time} = \frac{\# \text{ instructions}}{\text{program}} \times \frac{\text{clock cycles}}{\text{instruction}} \times \frac{\text{seconds}}{\text{clock cycle}}$$

▶ Performance depends on

- ▶ **Algorithm**: affects Instruction Count (IC), possibly CPI
- ▶ **Programming language**: affects IC, CPI
- ▶ **Compiler**: affects IC, CPI
- ▶ **Instruction set architecture**: affects IC, CPI, T_{clock}
- ▶ **Implementation Technology and Design**: affects CPI, T_{clock}

Clock Cycles and Performance - Example

- ▶ Program runs on Computer A:
 - ▶ CPU Time: 10 seconds
 - ▶ Clock (rate): 400MHz = 400×10^6 cycles/sec
- ▶ Computer B can run clock faster
 - ▶ But, requires 1.2 X clock cycles to perform same task
 - ▶ Desired CPU Time: 6 Seconds
- ▶ What should the clock frequency of Computer B be to reach this target?
- ▶ Key to approach: Performance equation

$$\text{CPU time} = \frac{\text{clock cycles}}{\text{program}} \times \frac{\text{seconds}}{\text{clock cycle}} = \frac{\text{clock cycles}}{\text{clock frequency}}$$

Clock Cycles and Performance - Example (cont'd)

- **First step: find clock cycles executed by Computer A**



- **Second step: find clock cycles executed by Computer B**



Clock Cycles and Performance - Example (cont'd)

- **Third step: given clock cycles and CPU time, solve for clock rate of Computer B**



Summary

- ▶ Response time vs. Throughput
- ▶ Elapsed time vs. execution time (CPU time)
- ▶ How does OS measure **CPU time**?
- ▶ What is the problem if **interrupt interval** is too short in measuring CPU time?
- ▶ Why is it difficult to measure the execution time of a program?
- ▶ Why do we use the number of clock cycles to measure the **elapsed time**?
- ▶ Why do we use CPI to measure **CPU time**?