

# Computer Organization

---

## Lecture 22 - Memory Hierarchy 2

Reading: 5.3-5.4 (cache), ~~5.7 (virtual memory)~~

Homework :

# Outline - Memory Systems

---

- ▶ **Overview**
  - ▶ Motivation
  - ▶ General Structure and Terminology
- ▶ **Memory Technology**
  - ▶ Static RAM
  - ▶ Dynamic RAM
  - ▶ Disks
- ▶ **Cache Memory** ◀
- ▶ **Virtual Memory**

# Four Key Cache Questions:

---

1. Where can block be placed in cache?  
(block placement)
2. How can block be found in cache? ...using a tag  
(block identification)
3. Which block should be replaced on a miss?  
(block replacement)
4. What happens on a write?  
(write strategy)

# Q1: Block Placement

---

## ► Where can block be placed in cache?

### ► In one predetermined place - direct-mapped

- Use fragment of address to calculate block location in cache
- Compare cache block with tag to test if block is present

### ► Anywhere in cache - fully associative


- Compare tag to every block in cache

### ► In a limited set of places - set-associative

- Use address fragment to calculate set (like direct-mapped)
- Place in any block in the set
- Compare tag to every block in set
- Hybrid of direct mapped and fully associative


# Concept

## (1) Direct-mapped




	tag	data
CI	00	
	01	
	10	
	11	

## (2) Fully associative



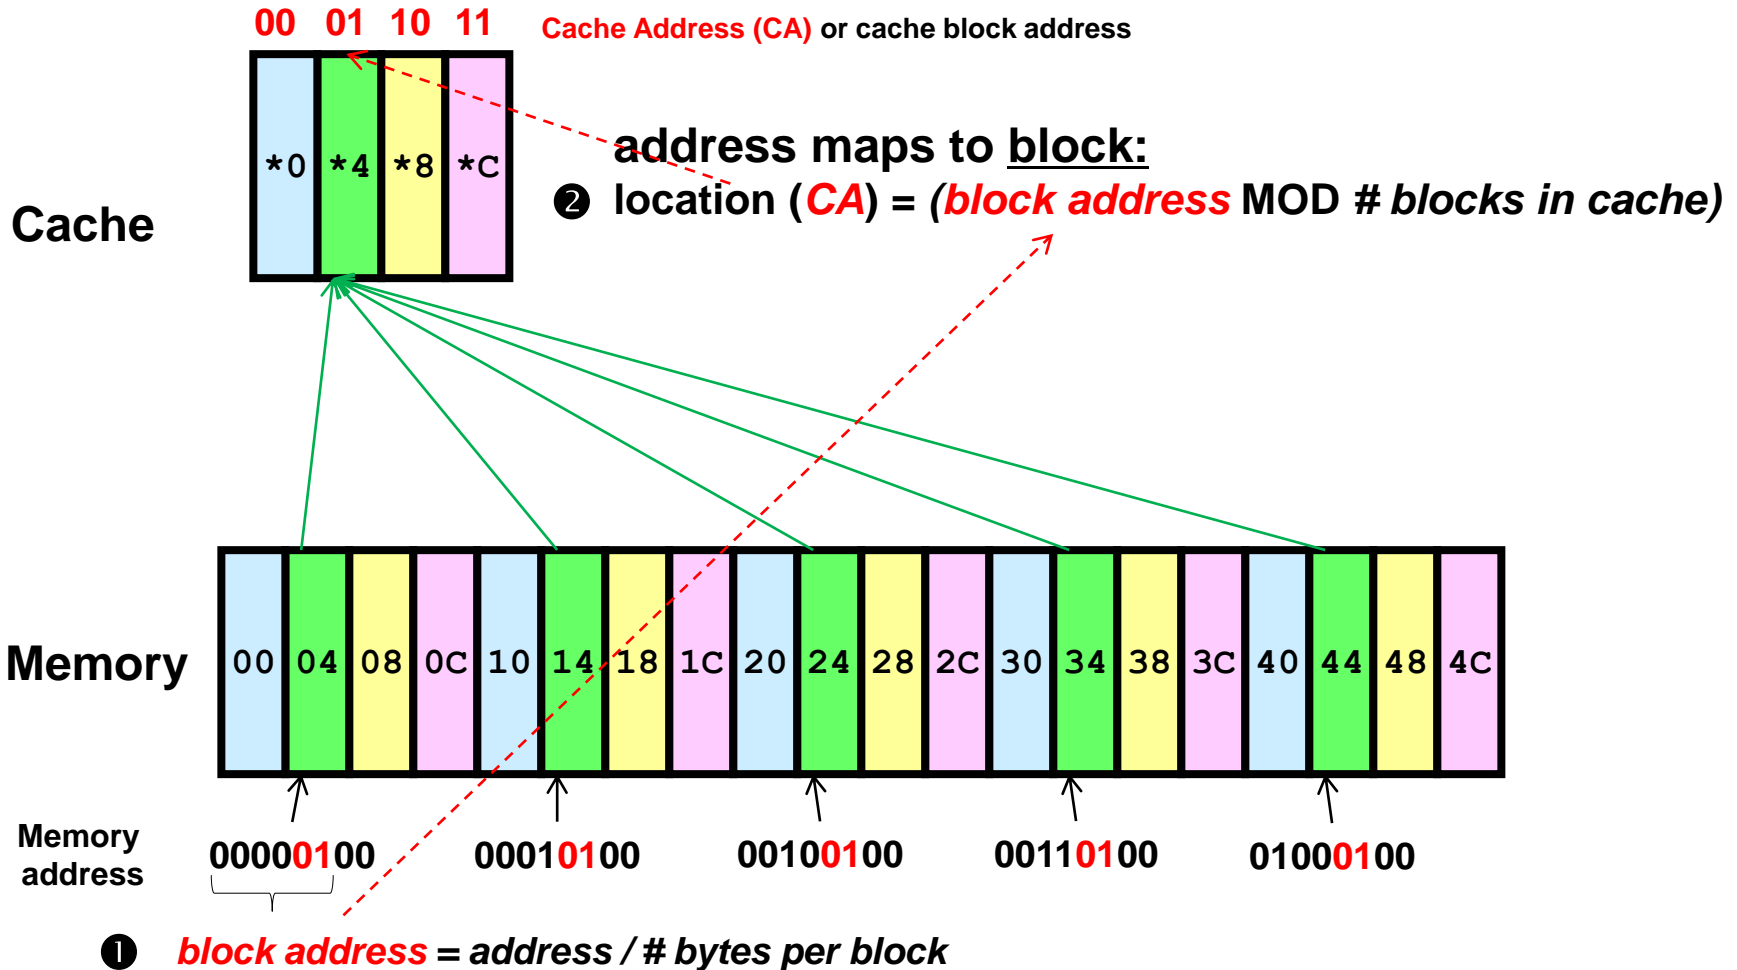
	tag	data
000		
001		
010		
011		
100		
101		
110		
111		

## (3) Set-associative



	tag	data
SI	00	
	01	
	10	
	11	


# Direct Mapped Block Placement




# Fully Associative Block Placement

---

Cache

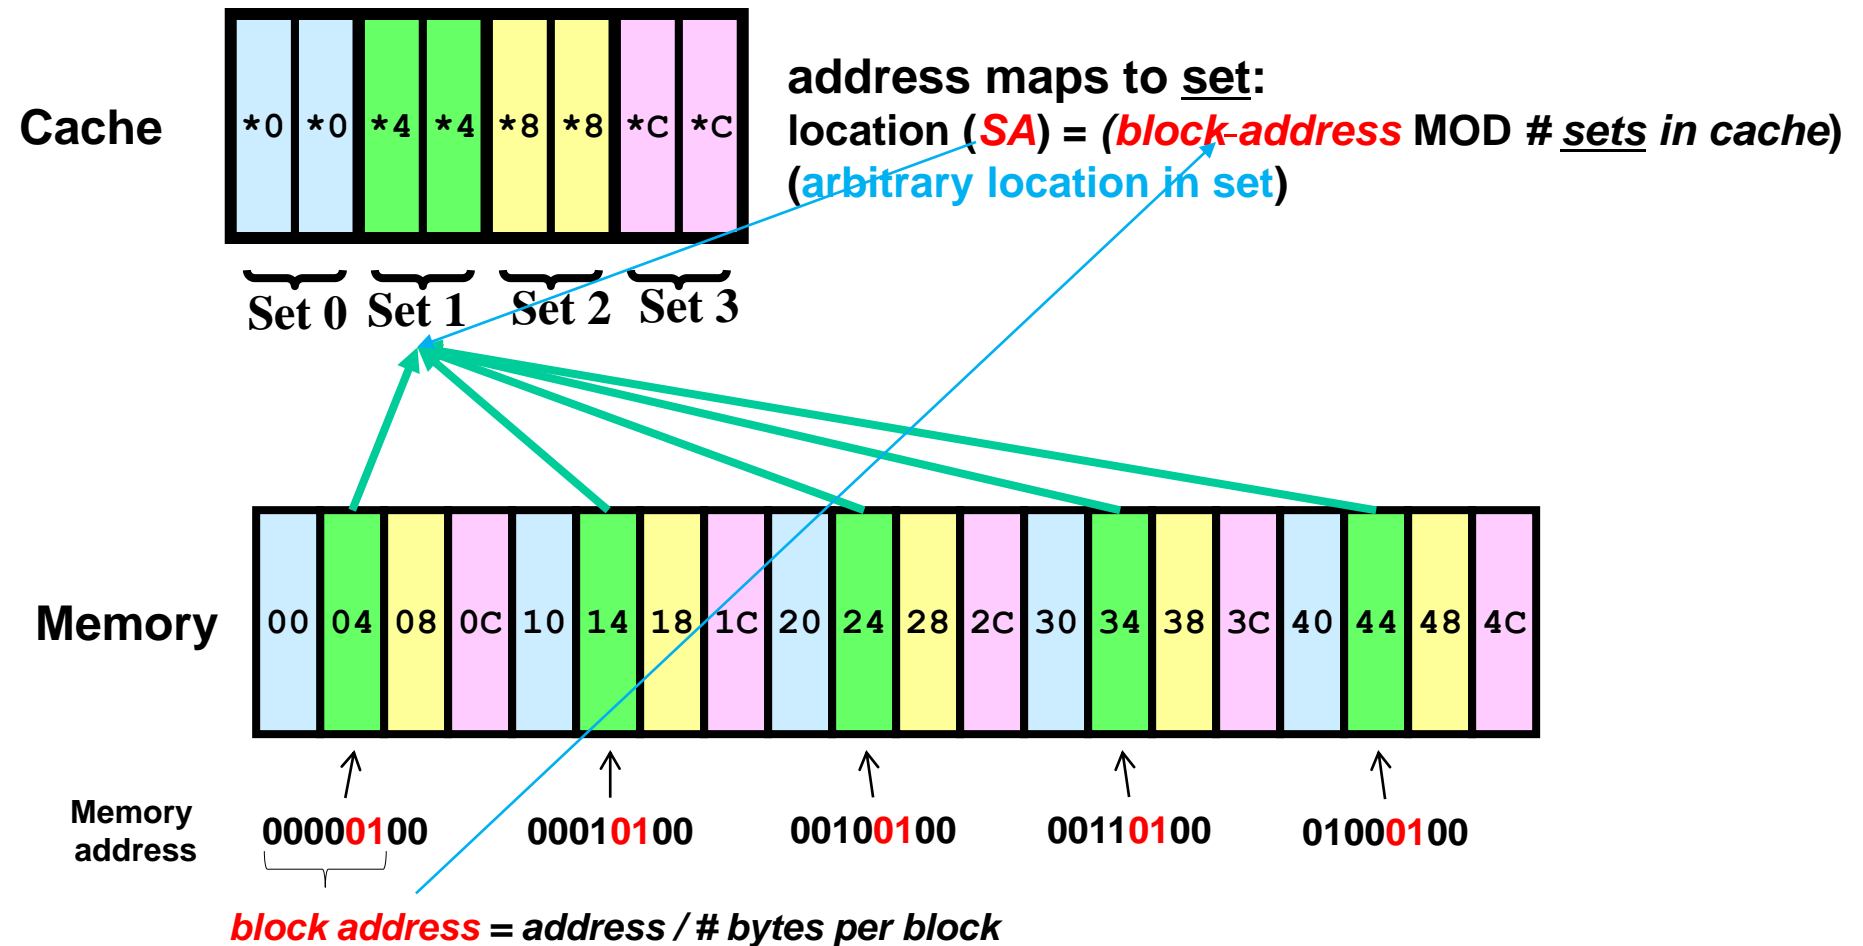


arbitrary block mapping  
location = *any*



Memory

# Set-Associative Block Placement





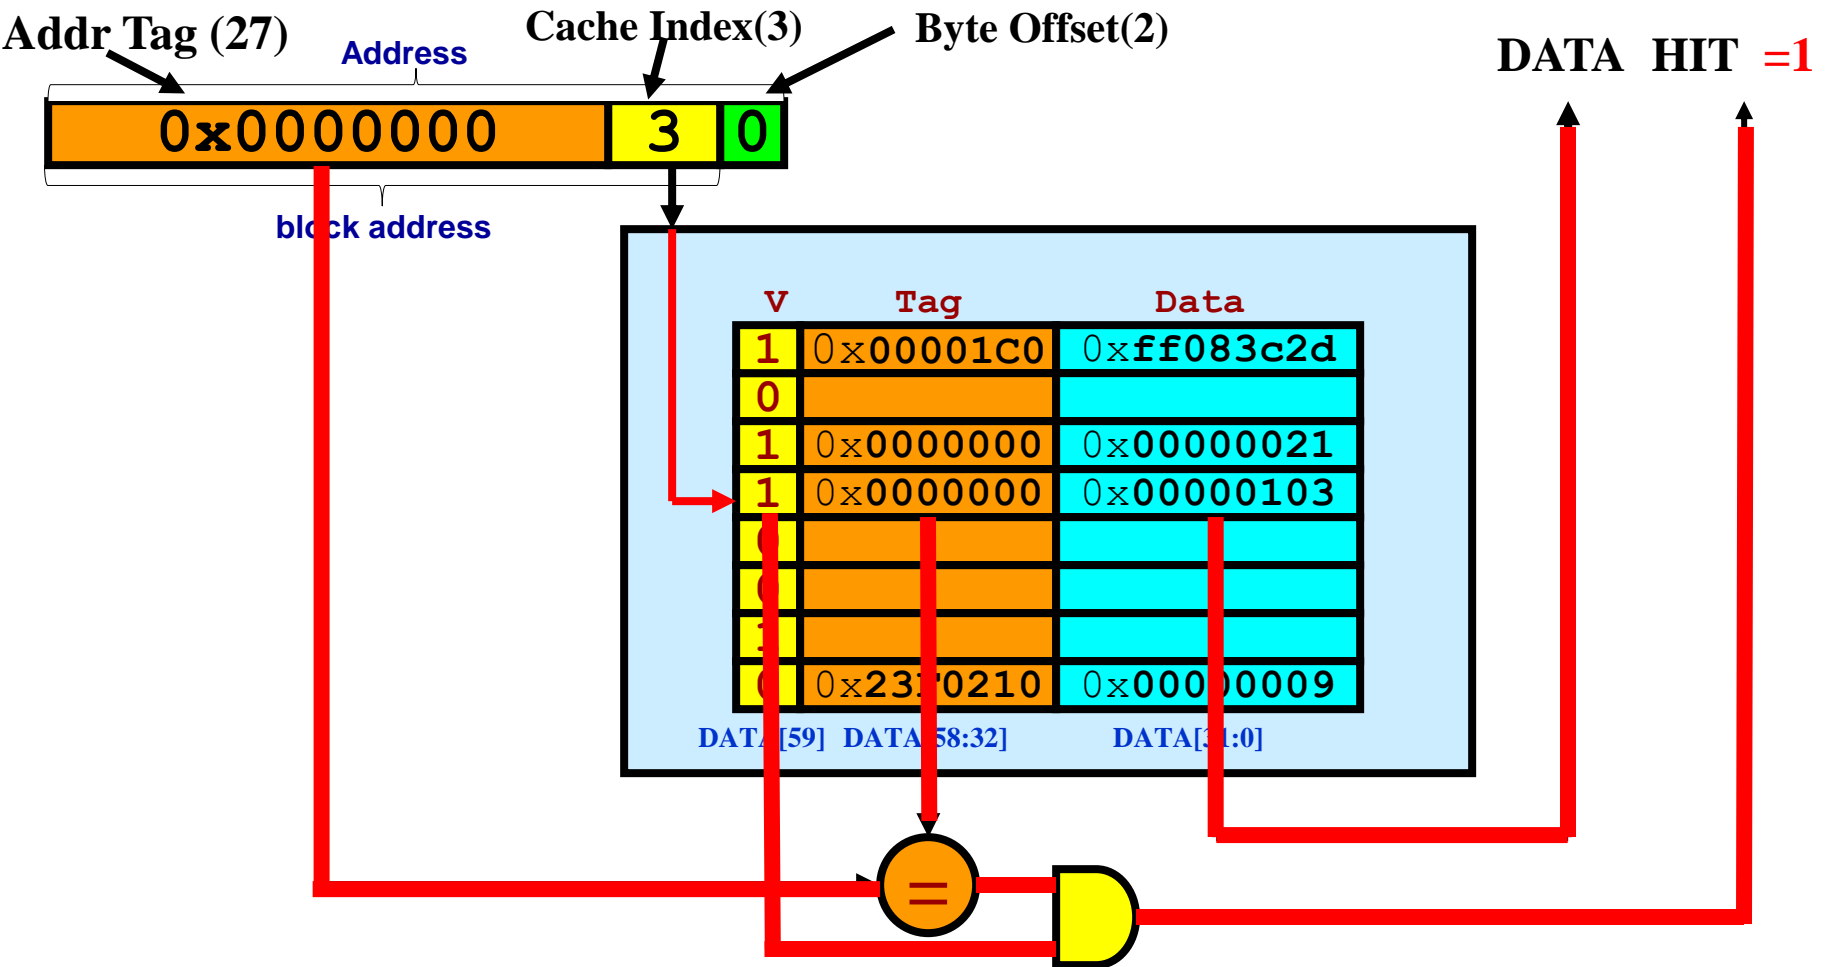
## Q2: Block Identification

PC 0x00001C0



- ▶ Every cache block has an address **tag** that identifies its location in memory
- ▶ Hit when tag and address of desired word **match** (**comparison by hardware**)
- ▶ Q: What happens when a cache block is empty?  
A: Mark this condition with a **valid bit**  
(**vaild => 1, invalid => 0**)

# Direct-Mapped Cache Design



# Cache Example

---

- ▶ 8-blocks, 1 word/block, direct mapped
- ▶ Initial state

	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	N		
111	N		

# Cache Example

Block address (b)	Hit/miss
10 <b>110</b>	Miss

Cache block address

	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
<b>110</b>	<b>Y</b>	<b>10</b>	<b>Mem[10110]</b>
111	N		

# Cache Example

---

Block address (b)	Hit/miss
11 010	Miss

	V	Tag	Data
000	N		
001	N		
010	Y	11	Mem[11010]
011	N		
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

# Cache Example

---

Block address (b)	Hit/miss
10 110	Hit
11 010	Hit

	V	Tag	Data
000	N		
001	N		
010	Y	11	Mem[11010]
011	N		
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

# Cache Example

---

Block address (b)	Hit/miss
10 000	Miss
00 011	Miss
10 000	Hit

	V	Tag	Data
000	Y	10	Mem[10000]
001	N		
010	Y	11	Mem[11010]
011	Y	00	Mem[00011]
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

# Cache Example

---

Block address (b)	Hit/miss
10 010	Miss

	V	Tag	Data
000	Y	10	Mem[10000]
001	N		
010	Y	10	Mem[10010]
011	Y	00	Mem[00011]
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		



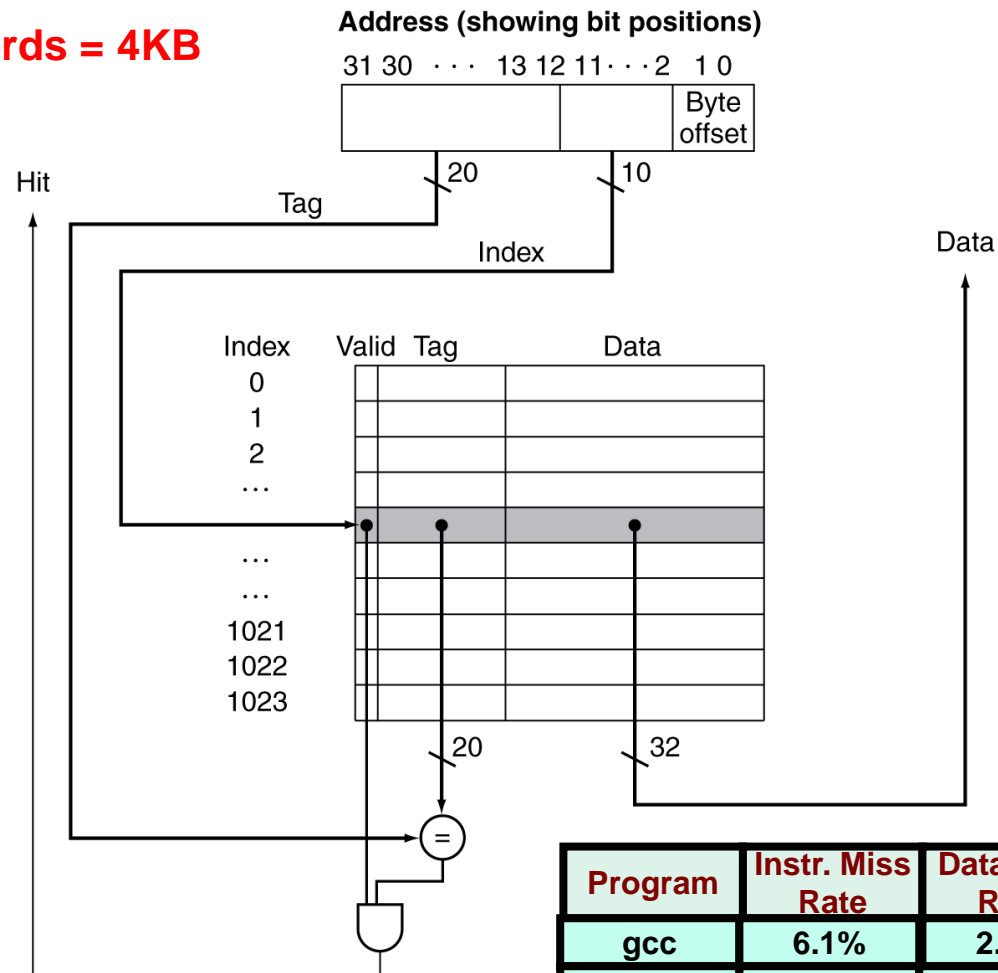
# Example: Larger Direct-Mapped Cache

**Cache Size = 1024 Words = 4KB**

**Byte Offset = 2 bits**

**Cache Index - 10 bits**

**Tag - 20 bits**



**Fig. 5.7**

Program	Instr. Miss Rate	Data Miss Rate	Combined Miss Rate
gcc	6.1%	2.1%	5.4%
spice	1.2%	1.3%	1.2%

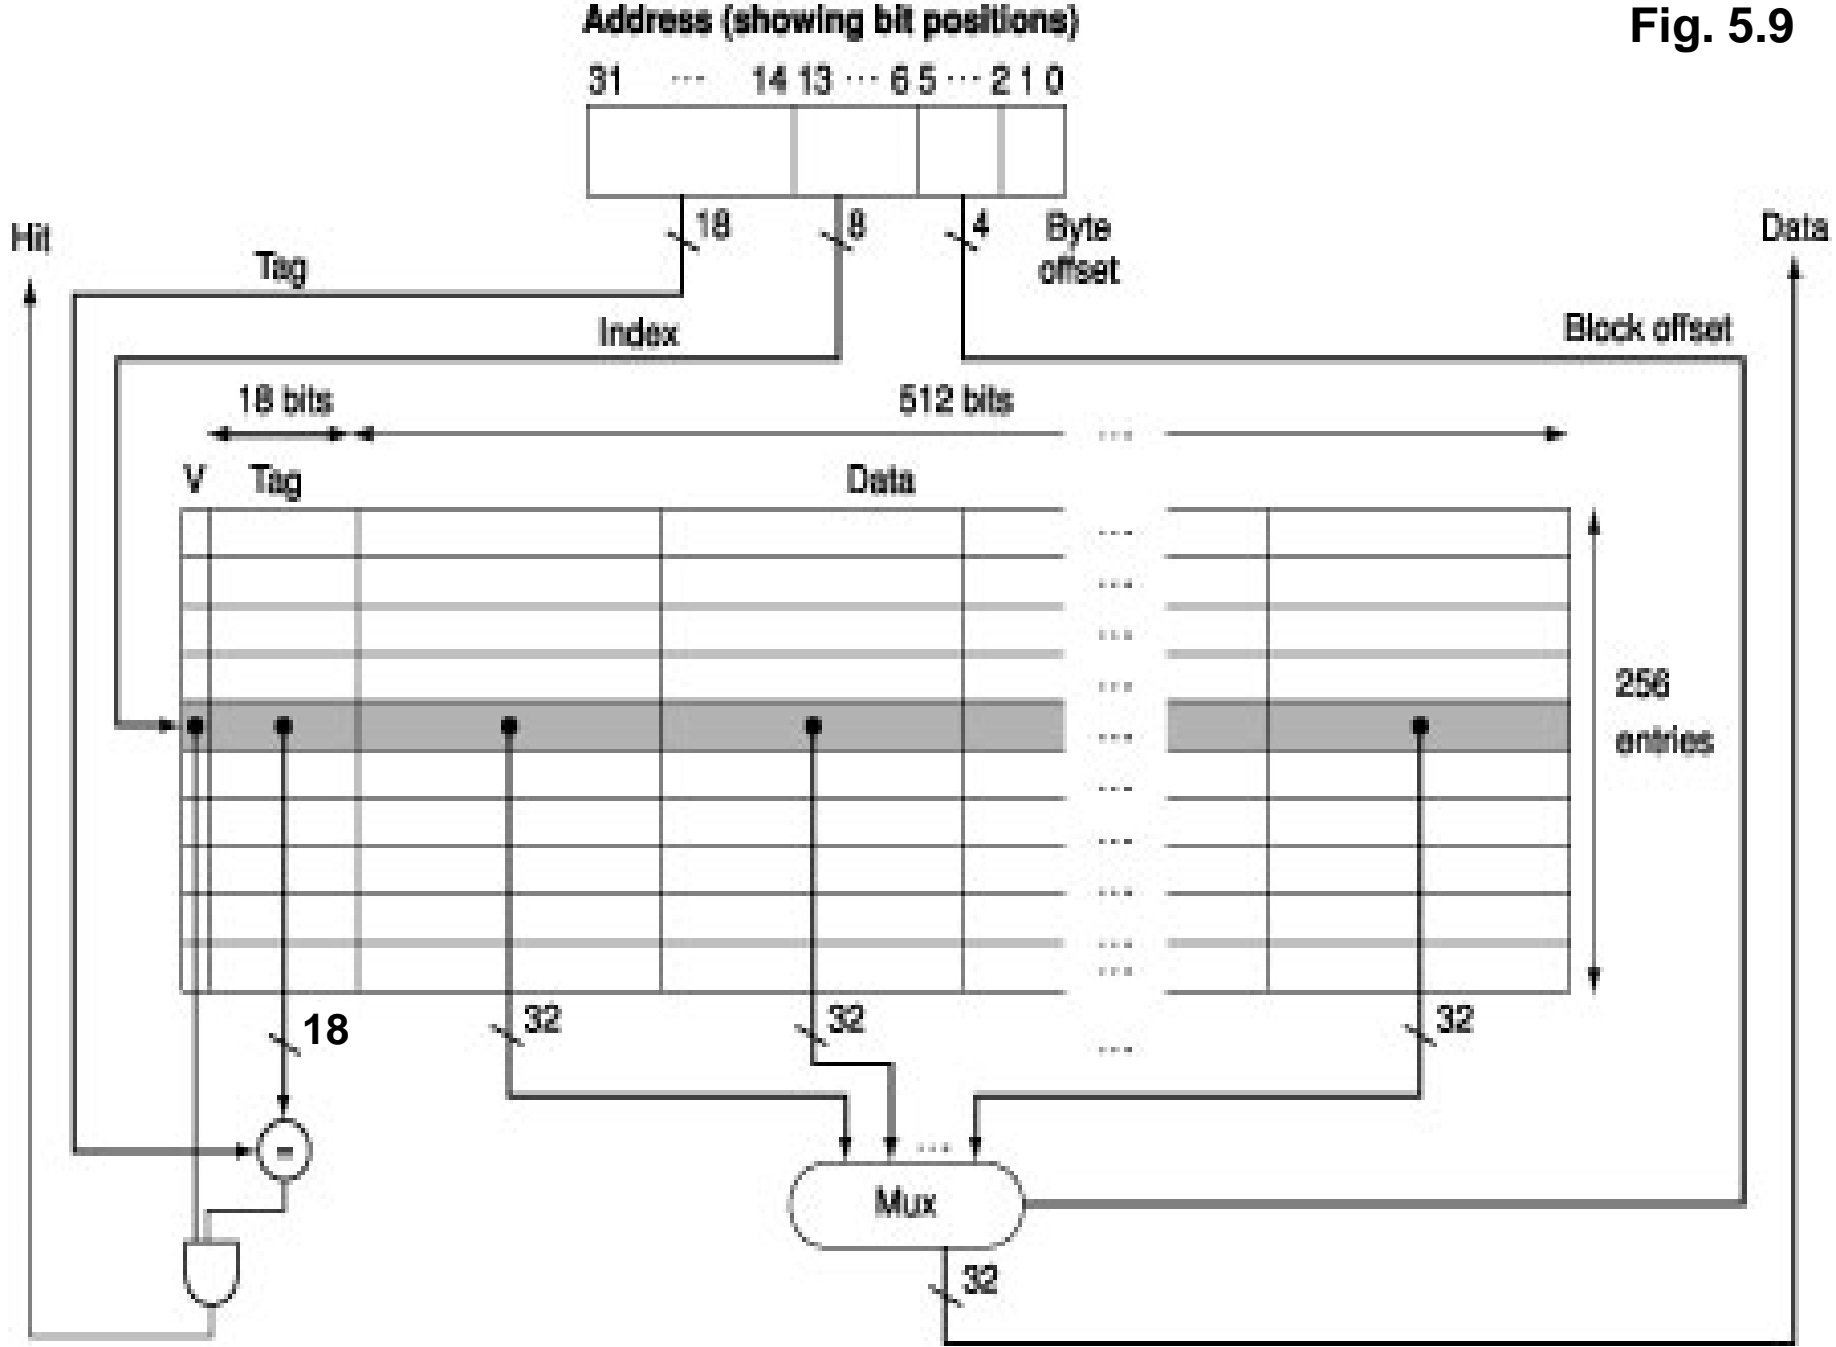
DecStation 3100 Cache - Old Fig. 7.10

## Variation: Larger Block Size

---

- ▶ Key advantage: take advantage of spatial locality
  - ▶ Move a number of locations to a cache at a time
- ▶ Disadvantages: competition, complicated writes


Fig. 5.9



# Example: Larger Block Size

---

- ▶ **64 blocks, 16 bytes/block**
- ▶ **Q: To what block number does address 1200 map?**
  - ▶ **Block address** =  $\lfloor 1200/16 \rfloor = 75$  (**Tag + Index**)
  - ▶ **Cache Index** (Address) =  $75 \text{ modulo } 64 = 11$  (**index**)



# Block Size Considerations

*larger block size vs. smaller block size with a fixed size of cache*

---

- ▶ Larger block size can result in **reduced miss rate**
  - ▶ Due to spatial locality
- ▶ But in a fixed-sized cache
  - ▶ Larger block size  $\Rightarrow$  fewer number of blocks (블락사이즈가 커지면 블락의 수는 줄어듦)
  - ▶ More competition  $\Rightarrow$  increased miss rate
  - ▶ Larger block size  $\Rightarrow$  pollution
- ▶ Larger **miss penalty** (블락사이즈가 커지면 Miss 패널티는 커짐)
  - ▶ Can override benefit of **reduced miss rate**
  - ▶ Early restart and critical-word-first can help

# Miss Rates vs. Block Size

- Note miss rate increases for larger block sizes




Fig. 5.8

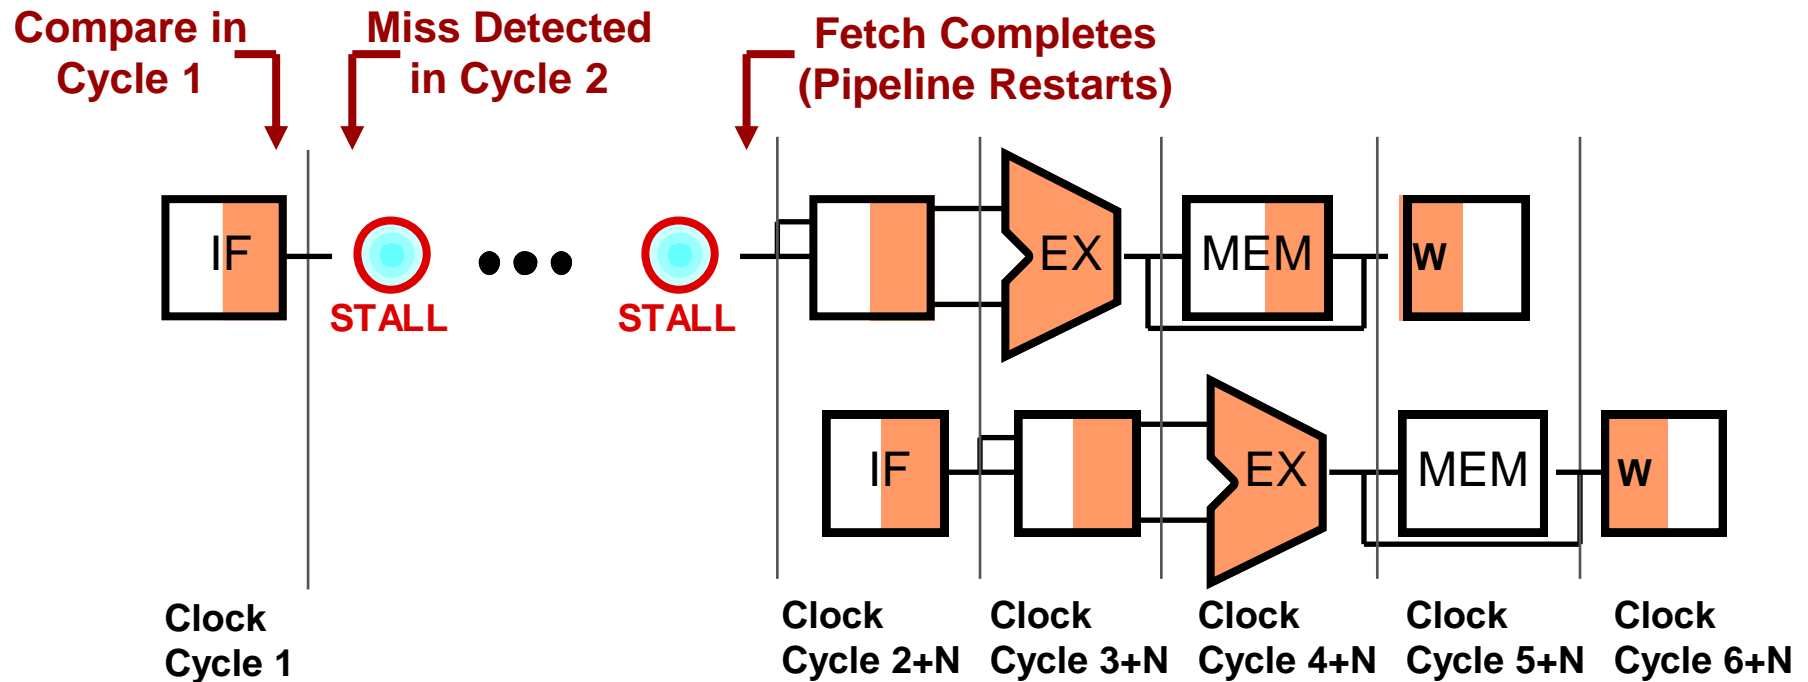
# Cache Misses

---

- ▶ On cache hit, CPU proceeds normally
- ▶ On cache miss
  - (1) Stall the CPU pipeline
  - (2) **Fetch block** from next level of memory hierarchy
  - (3) Restart instruction fetch OR  
Complete data access

# Cache Miss and the MIPS Pipeline

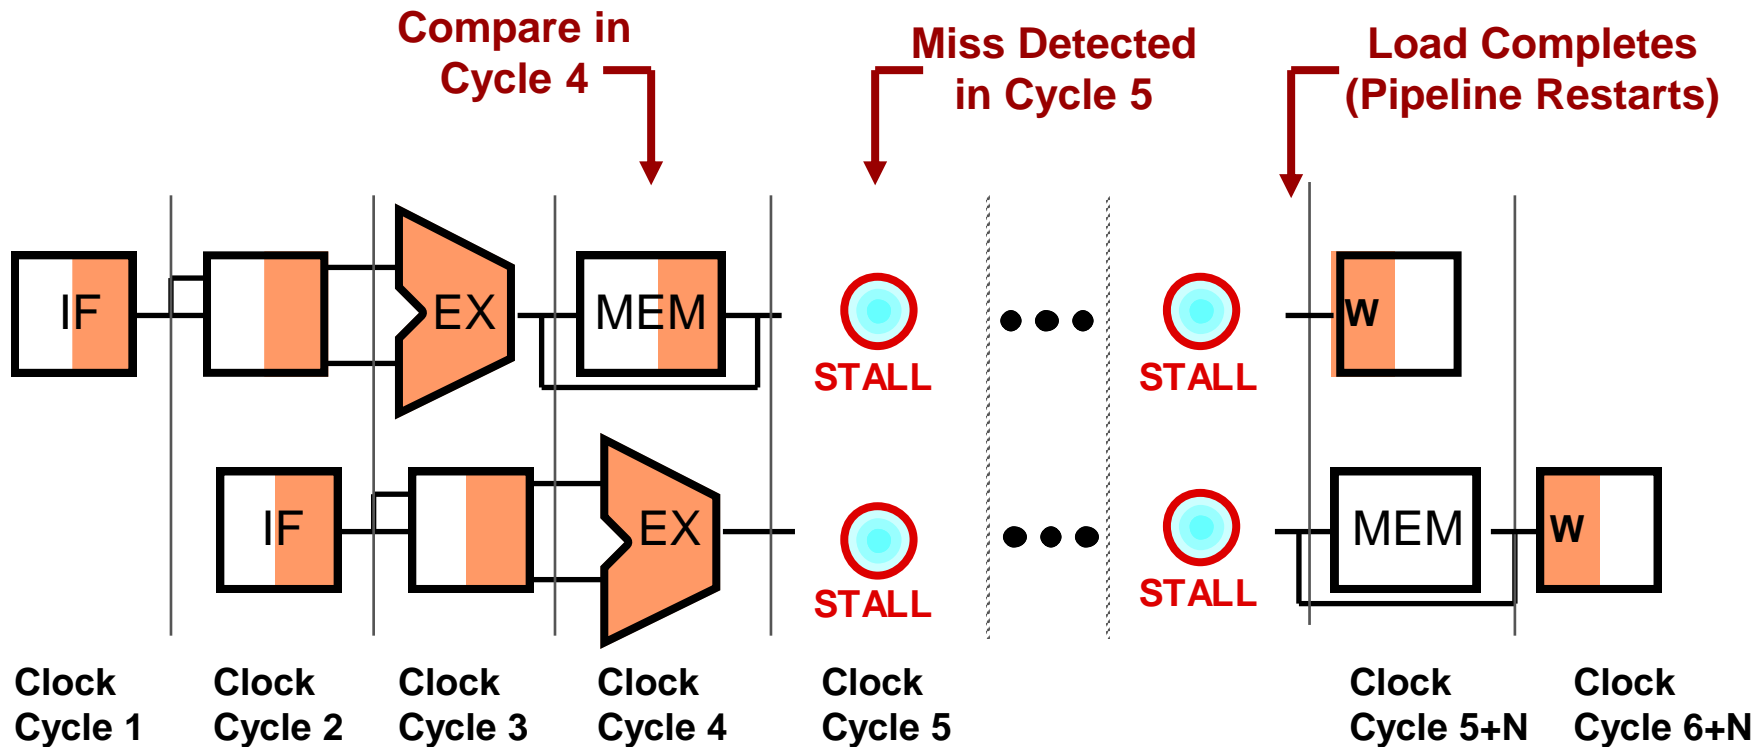
## ► Instruction Fetch





# Cache Miss and the MIPS Pipeline

## ► Load Instruction



**Today's lecture ends...**

# Four Key Cache Questions:

---

1. Where can block be placed in cache?  
(block placement)
2. How can block be found in cache? ...using a tag  
(block identification)
3. Which block should be replaced on a miss?  
(block replacement)
4. What happens on a write?  
(write strategy)

# Q4: Write Strategy

---

## ▶ What happens on a write?

- ① **Write through** - write to memory, stall processor until done
- ② **Write buffer** - place in buffer (allows pipeline to continue\*)
- ③ **Write back** - delay write to memory until block is replaced in cache

## ▶ Special considerations required when using DMA and multiprocessors

- ▶ DMA => coherence between different memory units
- ▶ Multiprocessors => **coherence** between caches

# Write-Through


---

- ▶ On **data-write hit**, just update the block in cache
  - ▶ But then cache and memory would be inconsistent
- ▶ ① **Write through**: also update memory at the same time
- ▶ But makes writes take longer
  - ▶ e.g., if base CPI = 1, 10% of instructions are stores, and write to memory takes 100 cycles,  
$$\text{Effective CPI} = 1 + 0.1 \times 100 = 11$$
- ▶ Solution: ② **Write buffer**
  - ▶ Holds data waiting to be written to memory
  - ▶ CPU continues immediately
    - Only stalls on write if write buffer is already full

### ③ Write-Back

---

- ▶ Alternative: On **data-write hit**, just update the block in cache
  - ▶ Keep track of whether each block is **dirty**
- ▶ When a **dirty** block is replaced
  - ▶ Write **it** back to memory
  - ▶ Can use a write buffer to allow **replacing block** to be read first



\***Dirty block** : a block with modified data

# Write Allocation

---

- ▶ What should happen on a **write miss**?
- ▶ Alternatives for **write-through**
  - ▶ Allocate on miss: fetch the block
  - ▶ Write around: don't fetch the block
    - Since programs often write a whole block before reading it (e.g., initialization)
- ▶ **For write-back**
  - ▶ Usually fetch the block

# Example: Intrinsity FastMATH

---

- ▶ **Embedded MIPS processor**
  - ▶ 12-stage pipeline
  - ▶ Instruction and data access on **each cycle**
- ▶ **Split cache: separate I-cache and D-cache**
  - ▶ Each 16KB: 256 blocks × 16 words/block
  - ▶ D-cache: write-through or write-back
- ▶ **SPEC2000 miss rates**
  - ▶ I-cache: 0.4%
  - ▶ D-cache: 11.4%
  - ▶ Weighted average: 3.2%



# Example: Intrinsity FastMATH

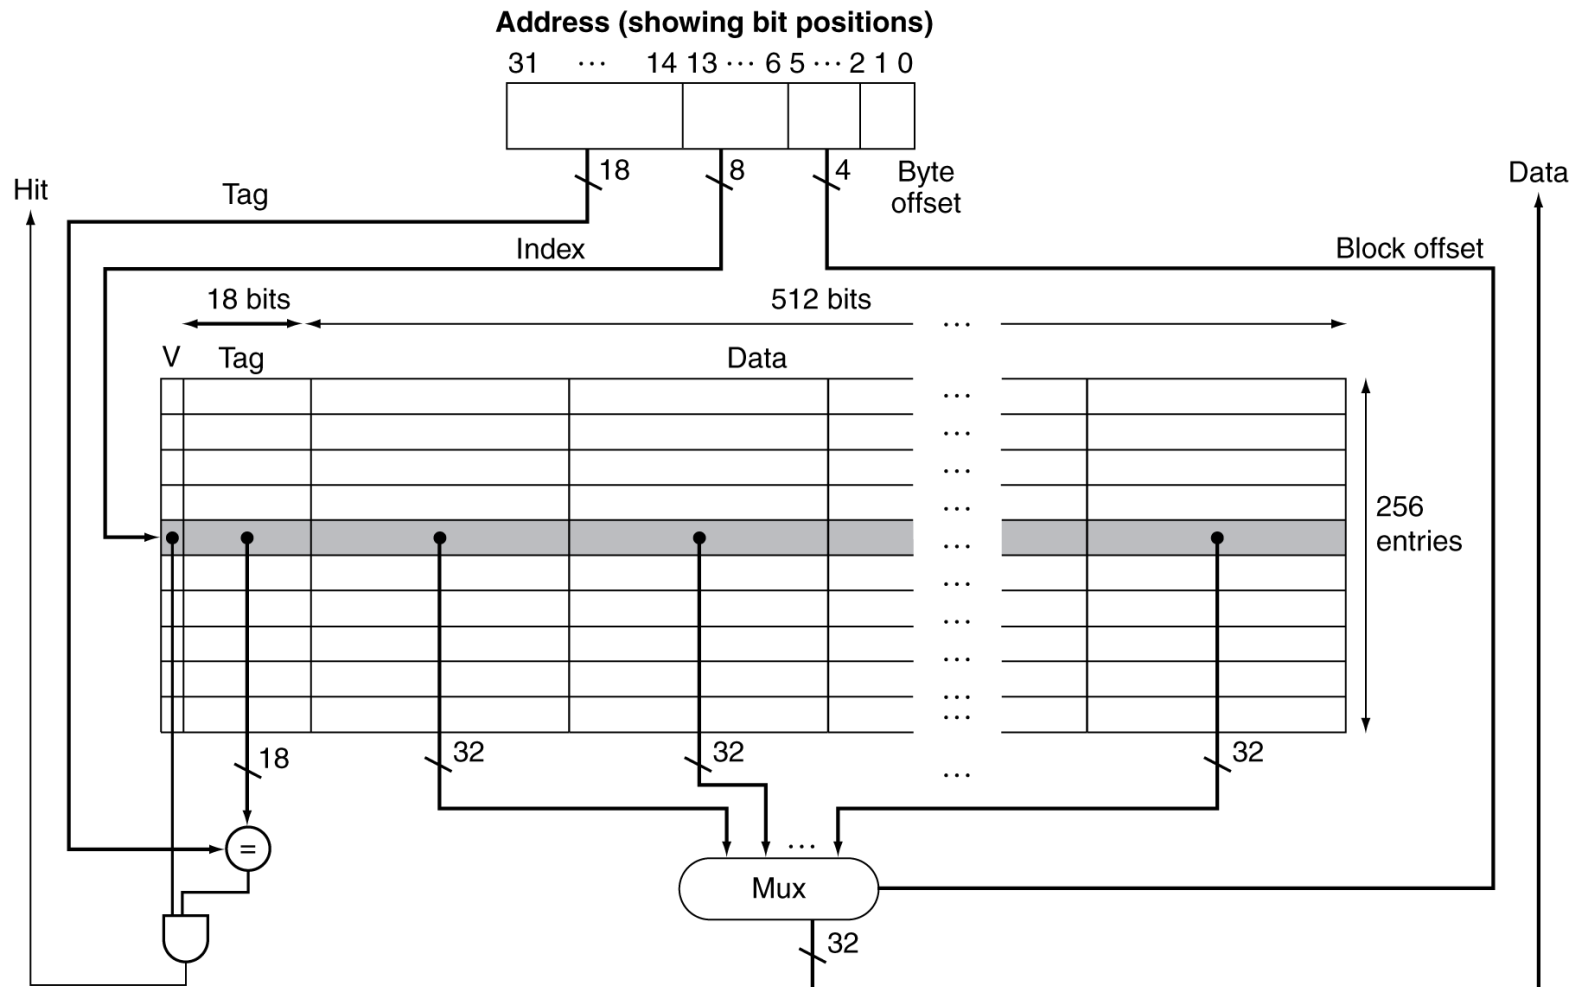


Fig. 5.9


# Main Memory Supporting Caches

---


- ▶ **Use DRAMs for main memory**
  - ▶ Fixed width (e.g., 1 word)
  - ▶ Connected by fixed-width clocked bus
    - Bus clock is typically slower than CPU clock
- ▶ **Example cache block read**
  - ▶ 1 bus cycle for address transfer
  - ▶ 15 bus cycles per DRAM access
  - ▶ 1 bus cycle per data transfer

# Increasing Memory Bandwidth


Fig. 5.11



a. One-word-wide memory organization



b. Wider memory organization



c. Interleaved memory organization

## ► 4-bank interleaved memory

- Miss penalty =  $1 + 15 + 4 \times 1 = 20$  bus cycles
- Bandwidth =  $16 \text{ bytes} / 20 \text{ cycles} = 0.8 \text{ B/cycle}$

## ► 4-word wide memory

- Miss penalty =  $1 + 15 + 1 = 17$  bus cycles
- Bandwidth =  $16 \text{ bytes} / 17 \text{ cycles} = 0.94 \text{ B/cycle}$

## ► For 4-word cache block, 1-word-wide DRAM

- Miss penalty =  $1 + 4 \times 15 + 4 \times 1 = 65$  bus cycles
- Bandwidth = 16 bytes / 65 cycles =  $0.25 \text{ B/cycle}$

# Measuring Cache Performance

---

## ▶ Components of CPU time

- ▶ Program execution cycles
  - Includes cache hit time
- ▶ Memory stall cycles
  - Mainly from cache misses

## ▶ With simplifying assumptions:

Memory stall cycles

$$= \frac{\text{Memory accesses}}{\text{Program}} \times \text{Miss rate} \times \text{Miss penalty}$$


$$= \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Misses}}{\text{Instruction}} \times \text{Miss penalty}$$

# Cache Performance Example (p. 477)

---

## ▶ Given

- ▶ I-cache miss rate = 2%
- ▶ D-cache miss rate = 4%
- ▶ Miss penalty = 100 cycles
- ▶ Base CPI (ideal cache) = 2
- ▶ Load & stores are 36% of instructions



## ▶ Miss cycles per instruction

- ▶ I-cache:  $0.02 \times 100 = 2$  (명령어는 cache miss 하면 모두 메모리 액세스)
- ▶ D-cache:  $0.36 \times 0.04 \times 100 = 1.44$  (데이터는 36%중에서 cache miss하면 메모리 액세스)

## ▶ Actual CPI = $2 + 2 + 1.44 = 5.44$

- ▶ Ideal CPU (with miss rate = 0) is  $5.44/2 = 2.72$  times faster than actual system

# Cache Performance Example (cont'd)

---

- ▶ What happens if we decrease CPI from 2 to 1?
- ▶ Miss cycles per instruction - no change!
  - ▶ I-cache:  $0.02 \times 100 = 2$
  - ▶ D-cache:  $0.36 \times 0.04 \times 100 = 1.44$
- ▶ Actual CPI =  $1 + 2 + 1.44 = 4.44$ 
  - ▶ Ideal CPU is  $4.44/1 = 4.44$  times faster than actual system
  - ▶ Speedup over original one is  $5.44/4.44 = 1.23X$
- ▶ As CPI drops, **fraction of time stalled** increases:
  - ▶ CPI = 2:  $3.44 / 5.44 = 63\%$
  - ▶ CPI = 1:  $3.44 / 4.44 = 77\%$


# Average Memory Access Time

---

- ▶ Hit time is also important for performance
- ▶ Average memory access time (AMAT)
  - ▶  $\text{AMAT} = \text{Hit time} + \text{Miss rate} \times \text{Miss penalty}$

- ▶ **Example: CPU with**

- ▶ 1ns clock
- ▶ **Hit time = 1 cycle**
- ▶ Miss penalty = 20 cycles
- ▶ I-cache miss rate = 5%
- ▶  **$\text{AMAT} = 1 + 0.05 \times 20 = 2\text{ns}$  (2 cycles/instruction)**



# Performance Summary

---

- ▶ When CPU performance increased
  - ▶ Miss penalty becomes more significant
- ▶ Decreasing base CPI (example in PPT 37)
  - ▶ Greater proportion of time spent on memory stalls
- ▶ Increasing clock rate
  - ▶ Memory stalls account for more CPU cycles
- ▶ Can't neglect cache behavior when evaluating system performance



# Block Replacement Policy


---

- ▶ Direct mapped: no choice
- ▶ Set associative
  - ▶ Prefer non-valid entry, if there is one
  - ▶ Otherwise, choose among entries in the set
- ▶ Least-recently used (LRU)
  - ▶ Choose the one unused for the longest time
    - Simple for 2-way, manageable for 4-way, too hard beyond that
- ▶ Random
  - ▶ Gives approximately the same performance as LRU for high associativity

# Multilevel Caches

---

- ▶ L1 cache attached to CPU
  - ▶ Small, but fast
- ▶ L2 cache services **primary (L1) cache misses**
  - ▶ Larger, slower, but still faster than main memory
- ▶ Main memory services **L2 cache misses**



# Multilevel Cache Example

---

## ▶ Given

- ▶ CPU base CPI = 1
- ▶ Clock rate = 4GHz => 0.25ns
- ▶ Miss rate/instruction = 2%
- ▶ Main memory access time = 100ns

## ▶ What is effective CPI with just primary cache?

- ▶ Miss penalty =  $100\text{ns} / 0.25\text{ns} = 400 \text{ cycles}$
- ▶ Effective CPI =  $1 + 0.02 \times 400 = 9$

Memory access penalty



# Example (cont.)

- ▶ Now add L2 cache
  - ▶ Access time = 5ns
  - ▶ Global miss rate to main memory = 0.5%
- ▶ Primary miss with L2 hit
  - ▶ **Penalty =  $5\text{ns}/0.25\text{ns} = 20$  cycles (L2 access cycle)**
- ▶ Primary miss with L2 miss
  - ▶ Extra penalty = 400 cycles (**memory access cycle**)
- ▶  $\text{CPI} = 1 + 0.02 \times 20 + 0.005 \times 400 = 3.4$
- ▶ Performance ratio =  $9/3.4 = 2.6$

Miss rate/instruction = 2%

